



BROWN

# DATA WAREHOUSING


## INTRODUCTION TO DATA SCIENCE

**CARSTEN BINNIG**  
**BROWN UNIVERSITY**



# DATA WAREHOUSES

**Definition:** A data warehouse is a **database that is optimized for analytical workloads** which **integrates data from independent and heterogeneous data sources**



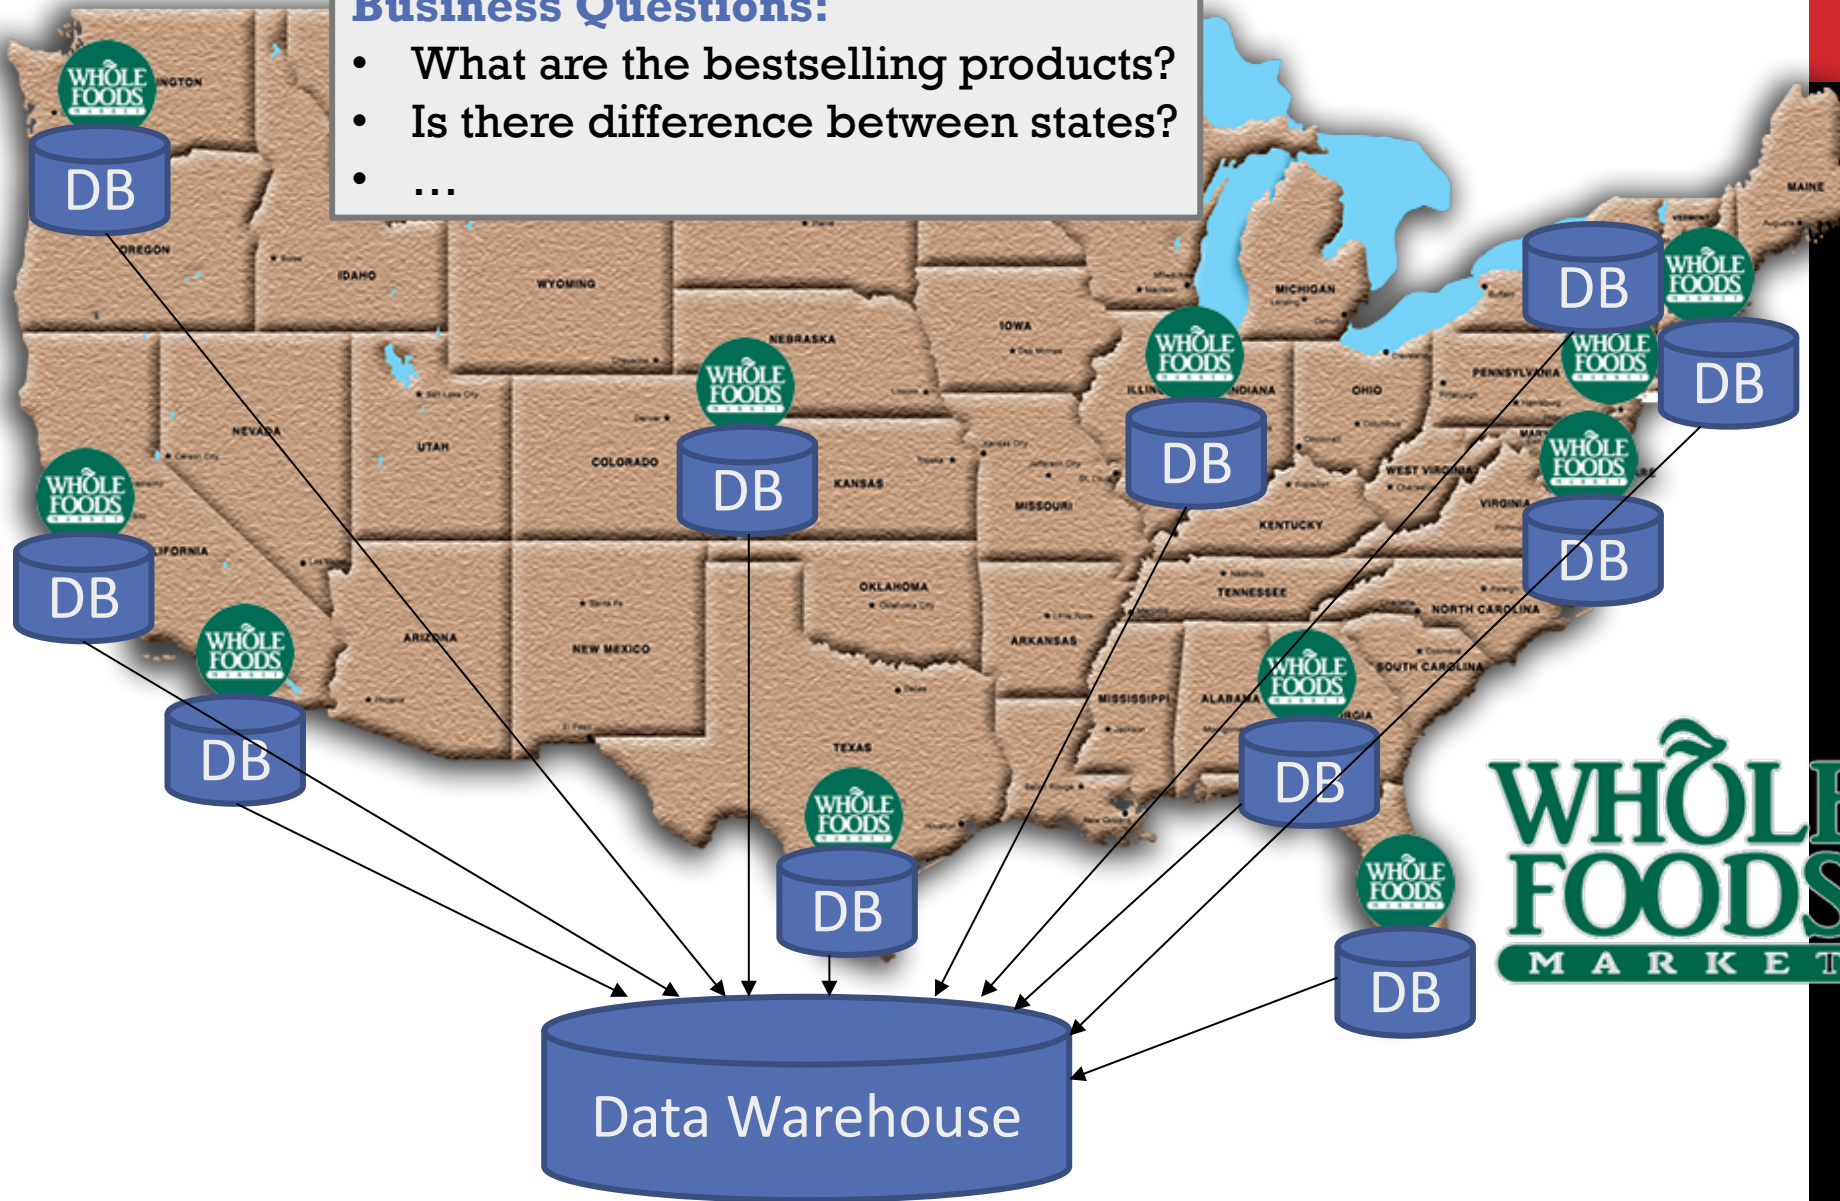
**Heterogeneous Data Sources**

**Decision Support / Data Mining**

# ENTERPRISE SCENARIO: WHOLEFOODS

## Business Questions:

- What are the bestselling products?
- Is there difference between states?
- ...



# OTHER APPLICATION DOMAINS

**Restaurant Chains (McDonalds, etc.)**

**Retailers (Nike, ...)**

**Insurance Companies**

**Banks**

**...**

# HISTORY OF DATABASES

## Age of Online Transaction Processing - OLTP (> 1970)

- **Goal: have access to up-to-date business transactions**
- 60s: IMS (hierarchical data model) => financial domain
- 80s: Oracle (relational data model) => most other domains (ERP, CRM)

## Age of Online Analytical Processing - OLAP (> 1990)

- **Goal: make business decisions**
- 90s: Data-warehousing extensions to relational databases
- Recently: New systems like in-memory column stores

# OLTP VS. OLAP

## **Online Transaction Processing (OLTP)**

- Current state of data is important
- Queries read / update only few records; AKA point queries or CRUD workloads (Create, Read, Update, Delete)
- Data Modeling: Avoid redundancy, normalize schemas

**Goal:** High throughput of transactions (Oracle 1995)

# EXAMPLE: CUSTOMER DATA (OLTP)

**CREATE**  
**READ**  
**UPDATE**  
**DELETE**

Standardadresse

Druckvorschau

Straßenansicht

Straße/Hausnummer: Adalperostraße 80

Postleitzahl/Ort: 85737 Ismaning

Land: DE Deutschland Region:

Zeitzone: CET

Geo-Daten

Längengrad: 701,07958984375000

Breitengrad: 2.893,77172851562500

Ortszeit:

Bundesland: Bayern

Dauer (hh:mm:ss): 04:40:48 Distanz (km): 404,680 Geodaten holen

# OLTP VS. OLAP

## Online Analytical Processing

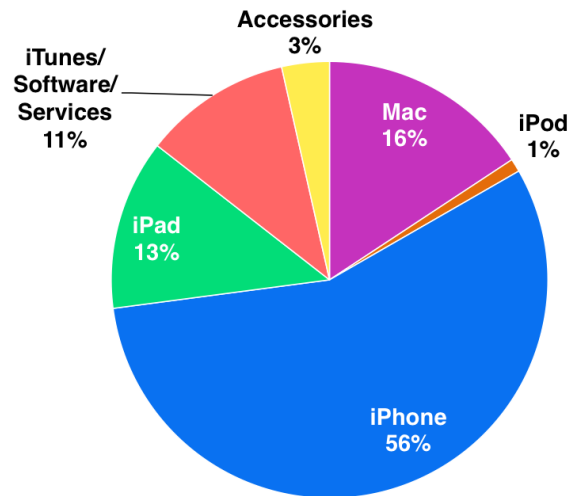
- History of data is important (not only the current state)
- Big queries (aggregate data, joins);
  - No Updates, only bulk loads
  - Data freshness is not that important!
- Modeling: Redundancy is a feature (i.e., de-normalized schemas are preferred)

**Goal:** Low latency of “big” queries ( $\leq 500\text{ms}$ )




# EXAMPLE: REVENUE (OLAP)

**Revenue By Product - Q4, 2014**



**Revenue by region – Q4, 2014**



# THE BIG PICTURE


## Data Sources



## Load



## Business Intelligence and Analytics



# FOR THE BUSINESS PERSON

## **Data Sources**

- CSV files

## **ETL**

- Copy and paste to Excel
- References + functions

## **Data Warehouse**

- Excel Sheets

## **Business Intelligence and Analytics**

- Excel functions
- Excel charts

# FOR THE BUSINESS PERSON

## **Data Sources**

- Web scraping, web services API
- Databases

## **ETL**

- Visual transformation tools
- Informatica, IBM DataStage, Ab Initio, Talend

## **Data Warehouse**

- Teradata, Oracle, IBM DB2, Microsoft SQL Server

## **Business Intelligence and Analytics**

- Business Objects, Cognos, Microstrategy
- SAS, SPSS, R

# FOR THE “HIP” WEB ENTERPRISE

## **Data Sources**

- Logs from the services tier
- User clicks, user comments, web crawl data...

## **ETL**

- Flume, Sqoop, Pig,/Crunch, Oozie (Workflow Scheduler)
- Hadoop/Hive, Spark/SparkSQL

## **Business Intelligence and Analytics**

- Custom web-based dashboards
- R

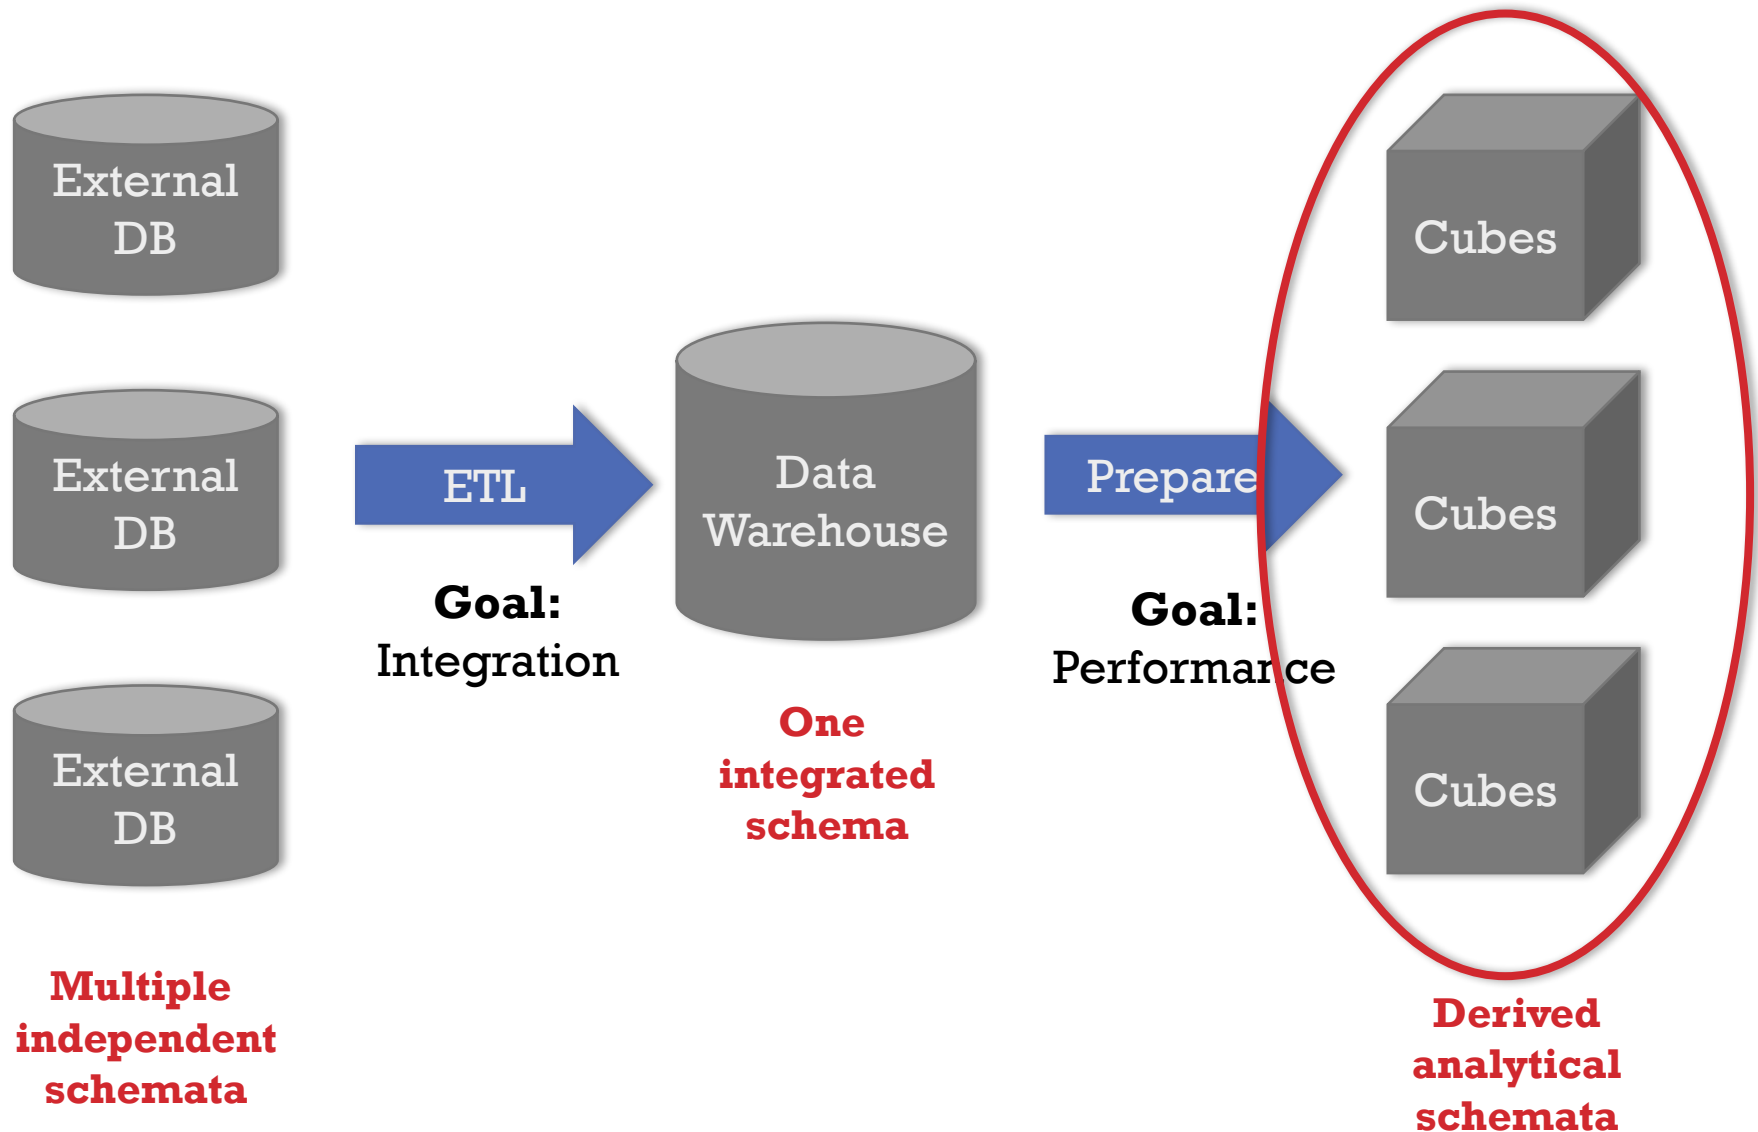
# DATA WAREHOUSING STEPS



# DATA WAREHOUSING STEPS




# A MORE TECHNICAL VIEW






# CUBE: MULTIDIMENSIONAL DATA MODEL



# MULTIDIMENSIONAL DATA (CUBE)



# DATA OPERATIONS

**Slice:** Cut a slice out of the cube (e.g. **product=„Beer“**)



**Dice:** Cut a smaller cube out of the data  
(e.g. **product=„Beer“** and **year=2013**)

**Drill-down:** Show details on the next level of detail  
(e.g., **zoom into from sales per month to sales per week**)

**Roll-up:** Aggregate data along a hierarchy  
(e.g., **zoom out from sales per month to sales per year**)

# RELATIONAL OLAP (ROLAP)

**ROLAP:** Store multidimensional cube data in a **relational database**




Relational Schema?

## Star-Schema:

**Fact table:** Store key figures (e.g., revenue, number of products sold, margin, ...)

**Dimension tables:** Store values on the axis!

# STAR-SCHEMA (ROLAP)



# DIMENSION TABLE

## Data in dimension tables:

- Distinct values of one axis of the cube (e.g. dates, product names, ...)
- Many different data types (texts, dates, ...)
- Often de-normalized (why?)

One dimension table is typically the **Time / Date table**

## Used to ...

- **Select data in fact table** (e.g. revenue in 2011): by joining dimension table with fact table
- **Group results** (e.g., revenue grouped by year)

# EXAMPLE: DIMENSION TABLE (CUSTOMER)

<u>custkey</u>	lastName	firstName	city	Country	region
1	Binnig	Carsten	Mannheim	Germany	Europe
2	Tellex	Stephanie	Providenc e	USA	North America
...	...	...	...	...	...

**PK** **Attributes for selection and grouping**

# FACT TABLE

## Data in fact tables

- Numeric key figures for aggregation e.g. revenue
- Foreign keys to dimensions (tables: customer, Product, date, ...)
- Mostly numeric data

**Key figures** are used for **aggregations** (e.g., total of orders, quantity of sales)

Data in fact table is **constantly growing**!

**Primary key of fact table:** Composed of all foreign keys



# EXAMPLE: FACT TABLE (SALES)

<u>customer key</u>	<u>product key</u>	<u>date key</u>	...	revenue	quantity
1	1	1	...	1000	10
1	2	1	...	100	1
2	1	3	...	800	9
...	...	...	...	...	...

Foreign Keys to dimension tables

Key figures

# DIMENSIONS: HIERARCHIES

Dimensions often **describe a hierarchy** (i.e., 1:N relationships between entities)

**Static hierarchies:** Levels in hierarchy is fixed  
(e.g. Year->Month->Day or Region->Country->City)

**Flexible hierarchies:** Dynamic number of levels  
(e.g. management hierarchies, bill of materials – BOM)

# STATIC HIERARCHIES

**Levels of a hierarchy** are represented by **different columns**

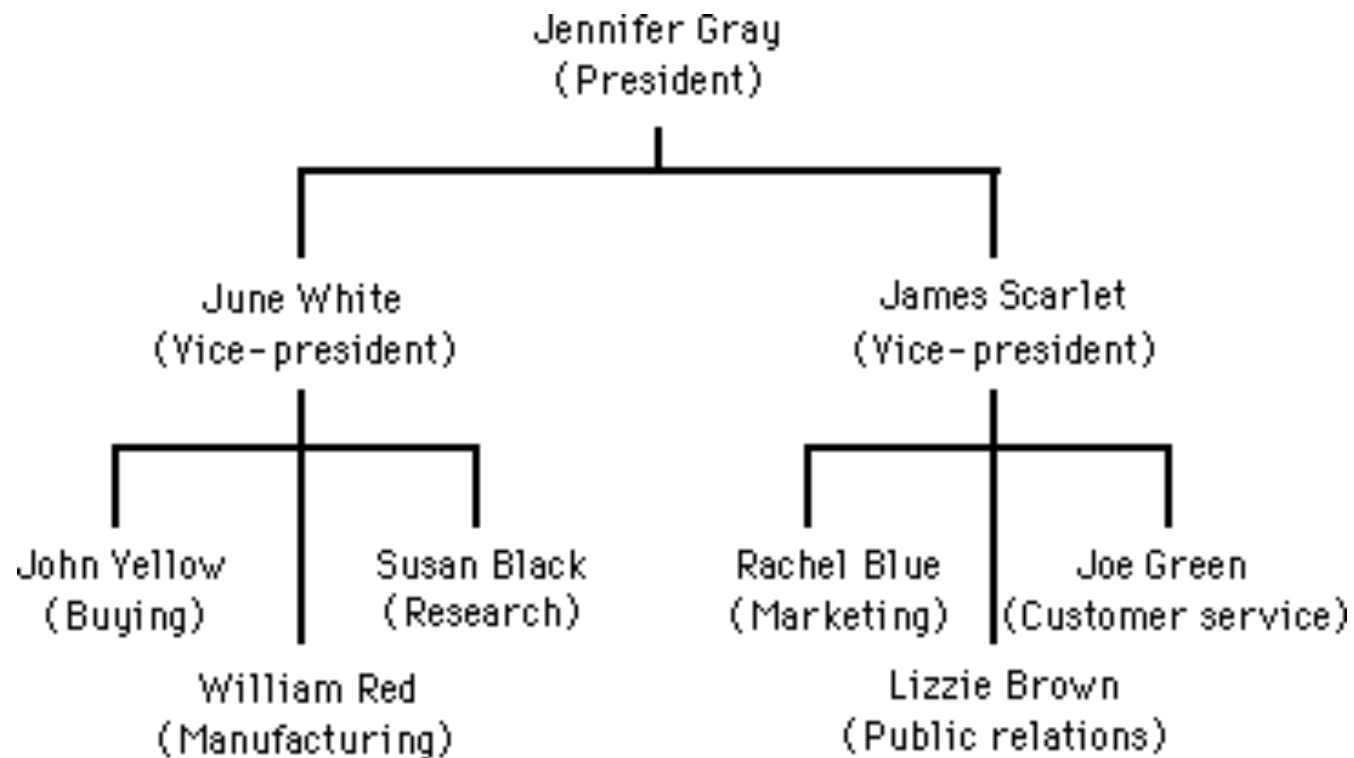
## Region->Country->City

City	Country	Region
Mannheim	Germany	Europe
Mosbach	Germany	Europe
...	...	...

## Year->Month->Day

Year	Month	Day
2012	01	1
2012	01	2
...	...	...
2012	01	31
...	...	...


# FLEXIBEL HIERARCHIES



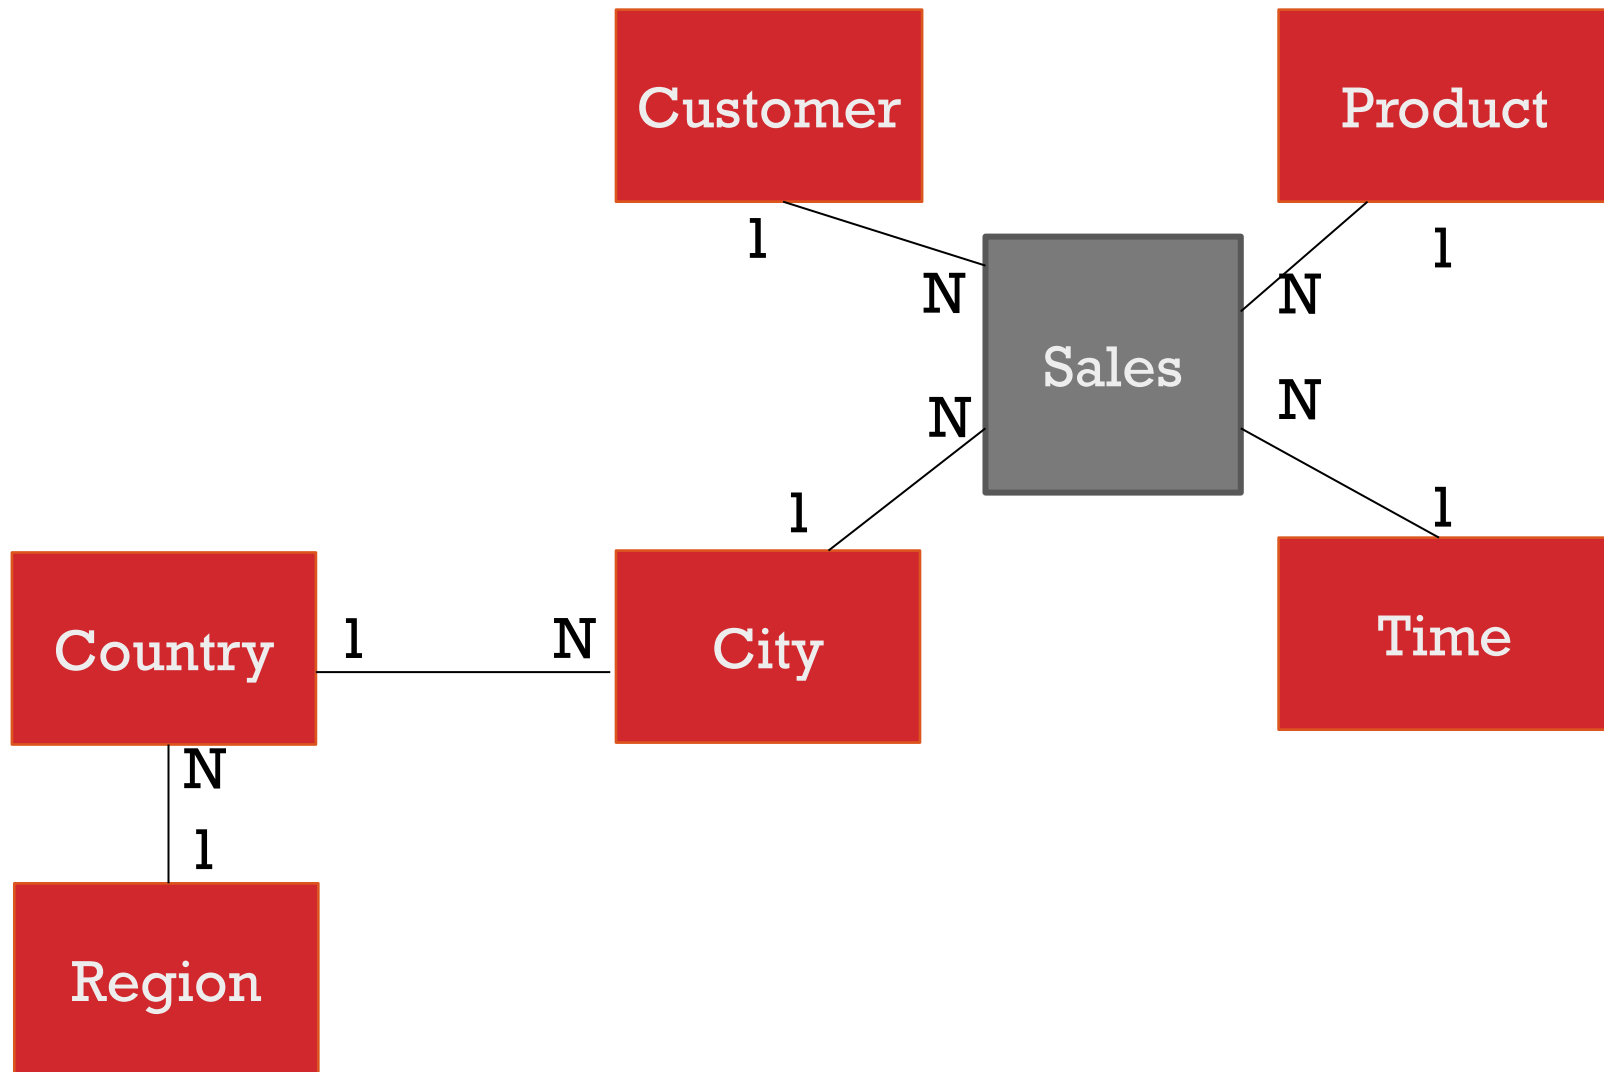
# FLEXIBEL HIERARCHIES

**Levels of hierarchy** are represented as **recursive relationships** (e.g., management hierarchy)

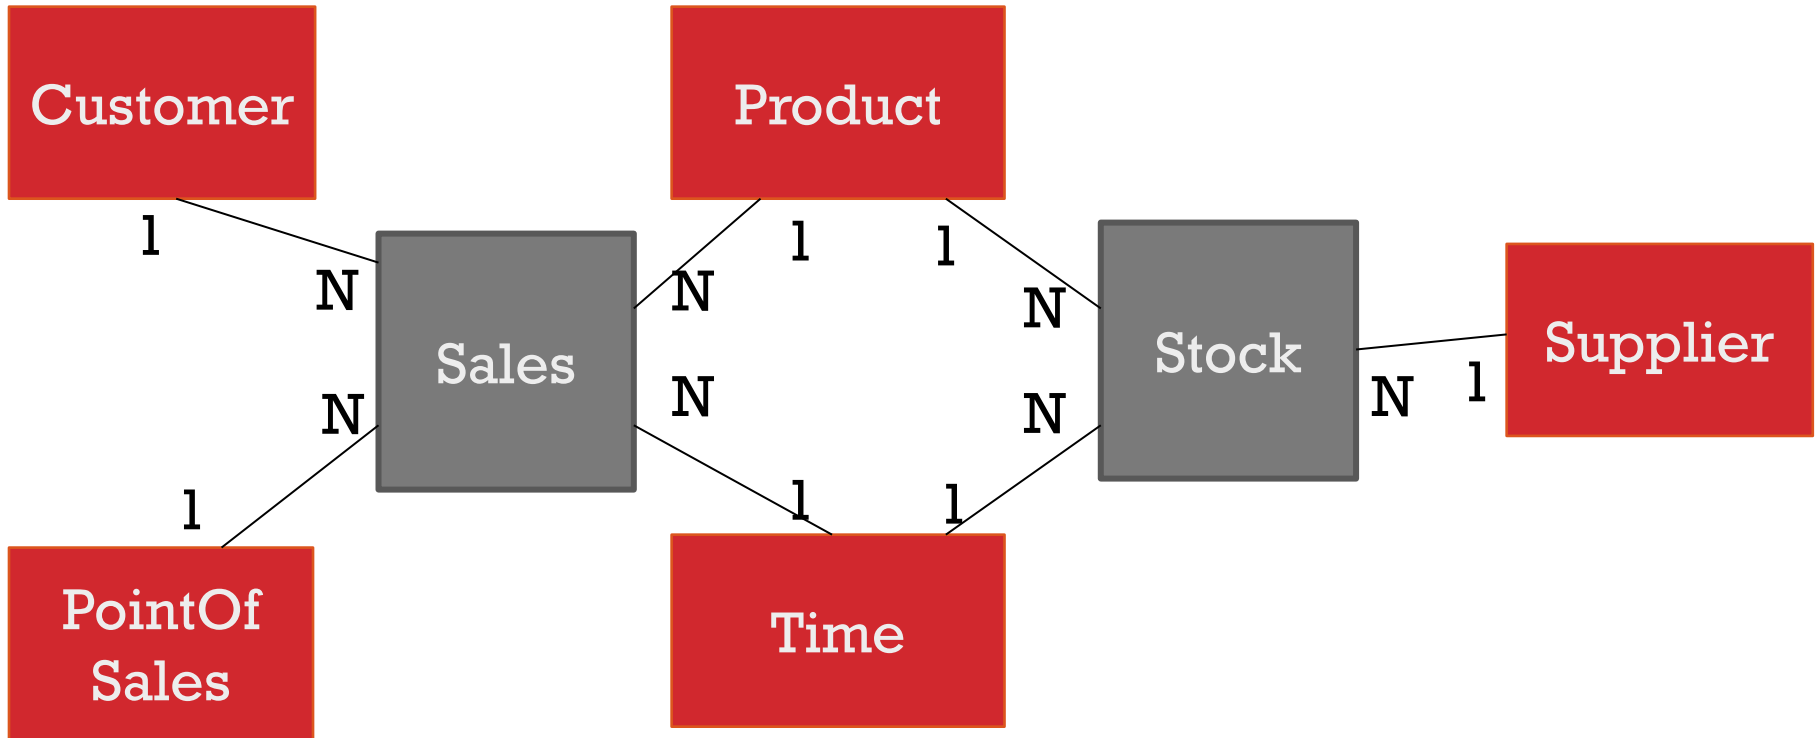
empKey	lastName	bossKey
1	...	NULL
2	...	1
3	...	1
4	...	2
5	...	2



# OTHER SCHEMATA: SNOWFLAKE




# OTHER SCHEMATA: GALAXY



# CLICKER QUESTION

**An OLTP database tracks which user has borrowed which books for how long. We want to be able to answer questions like ‘who are the users with the longest lending (per book, per genre)?’**




**How should the fact table look like?**

- A) Lendings(bookId, genreId, userId, days)
- B) Lendings(bookId, genreID, days)
- C) Lendings(bookId, userId, days)



# CLICKER QUESTION

**An OLTP database tracks which user has borrowed which books for how long. We want to be able to answer questions like ‘who are the users with the longest lending (per book, per genre)?’**



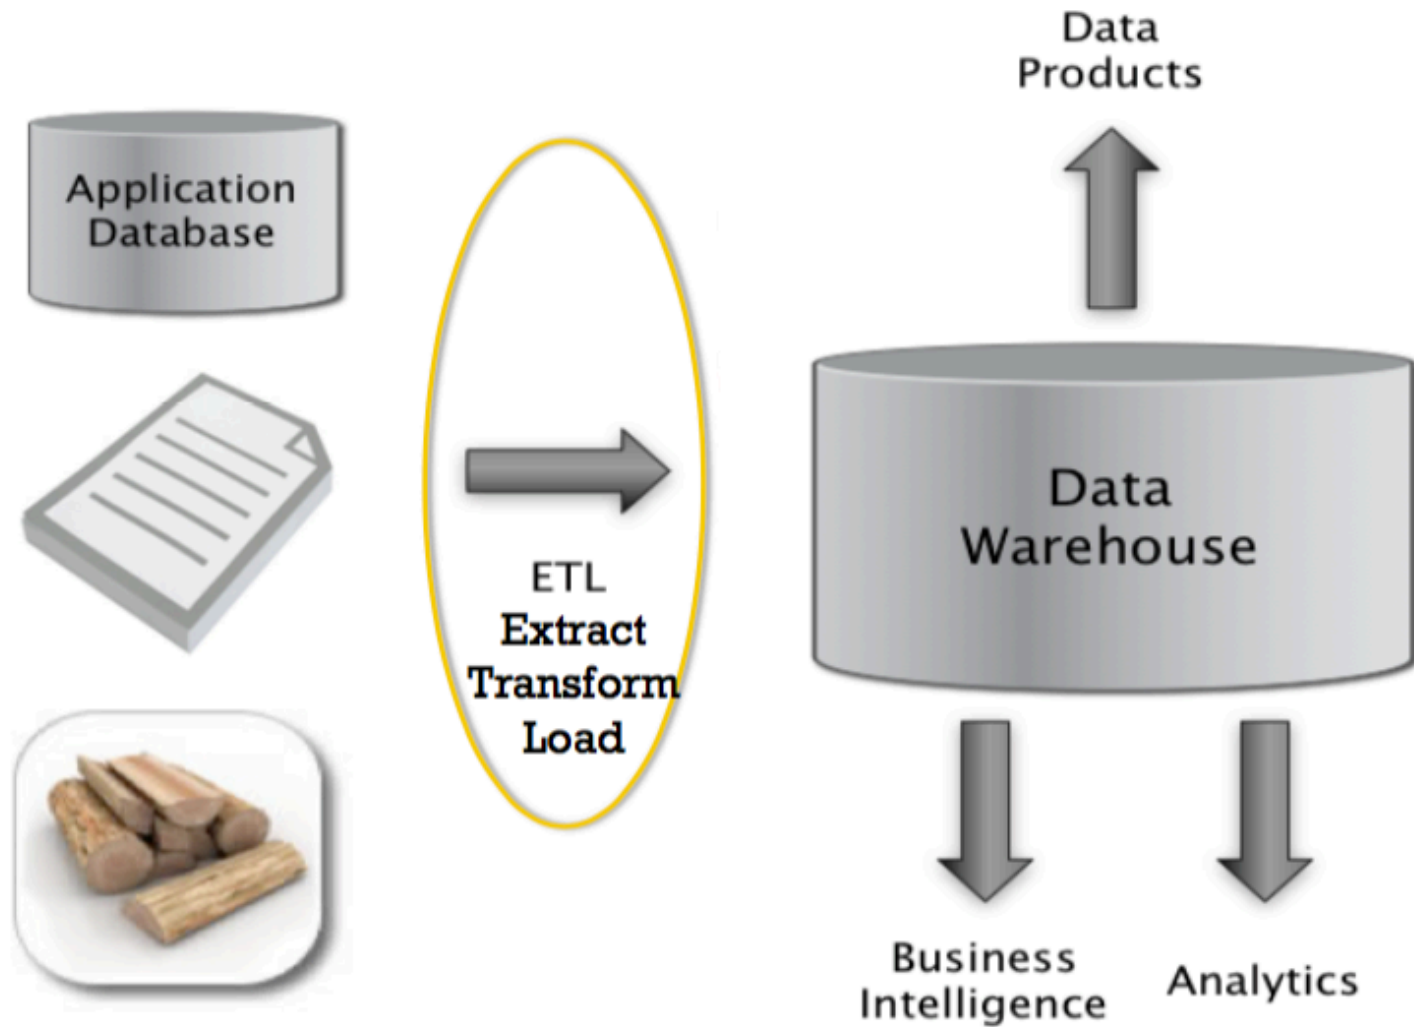
**How should the fact table look like?**

- A) **Lendings(bookId, genreId, userId, days)**
- B) Lendings(bookId, genreID, days)
- C) Lendings(bookId, userId, days)

# DATA WAREHOUSING STEPS



# THE BIG PICTURE



# DATA INTEGRATION

Data integration is done by ETL Processes

- **Extract:** extract data out of an operational data source
- **Transform:** cleanse it and transform it into the target schema (e.g., split first and last names)
- **Load:** append it to the tables of a data warehouse

**Operational Sources:** files, databases, event logs, ...

**Sink (Data Warehouse):** RDBMS, specialized OLAP engines, ...

# ETL WORKFLOWS

**The ETL pipeline or workflow often consists of many sequential steps**


- Often a mix of tools involved (Web-Service APIs, tools to reformat data, ... )
- Analogy: Unix pipes and filters -> `$ cat data_science.txt | wc | mail -s "word count" hammer@example.com`

**If the workflow is to be run more than once, it can be scheduled**

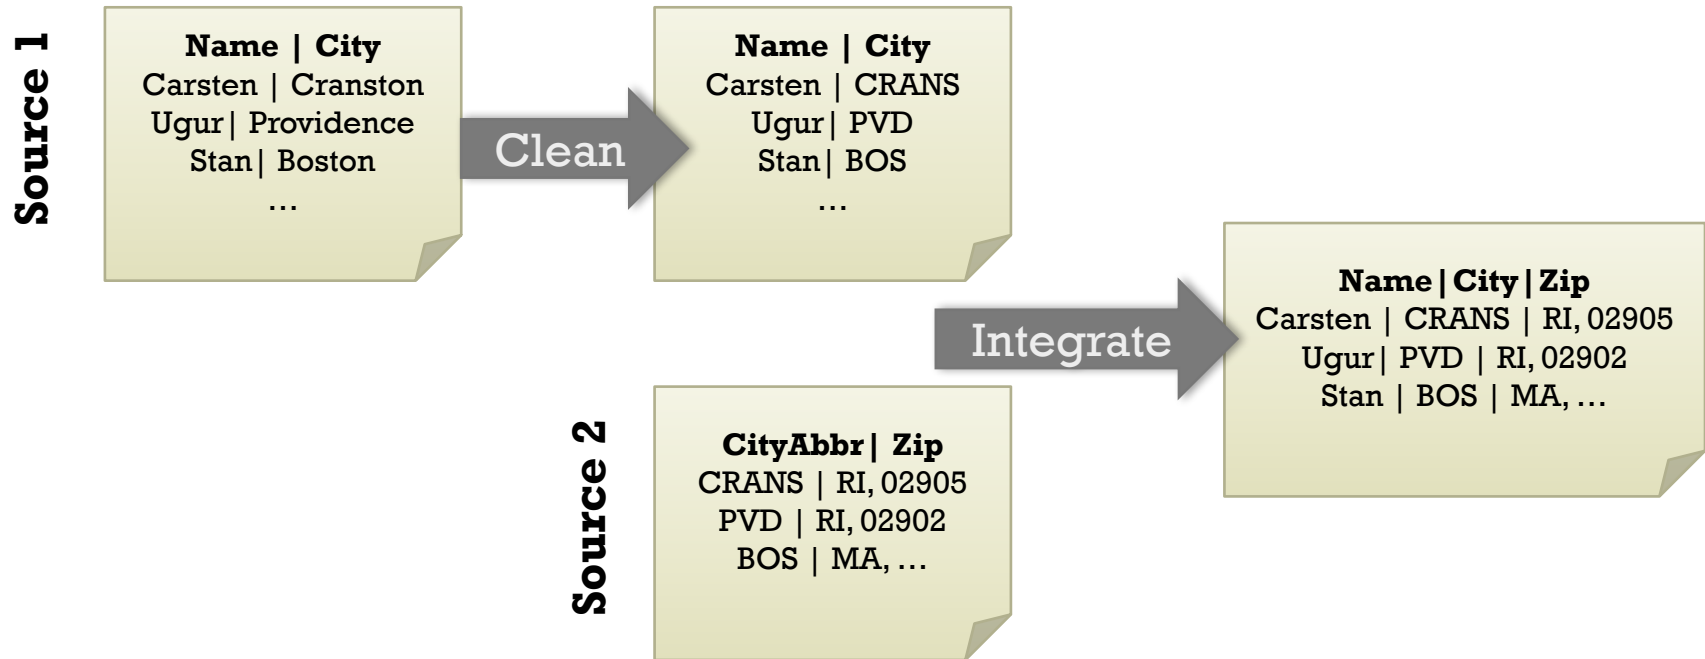
- Scheduling can be time-based or event-based

**Transformations are most complex Product (Separate slides on Data Integration!)**

# EXTRACT DATA




# TRANSFORM DATA




## Typical Tasks

- Clean Data (e.g., add missing values, correct mis-spellings, ...)
- Integrate Data when using multiple sources (e.g., schema matching )
- Execute relational transformations (e.g., joins)

# LOAD DATA



Load into Warehouse  
(e.g., generate keys)






# DATA WAREHOUSING STEPS



# RECAP: STAR-SCHEMA (ROLAP)



# STAR-QUERY

Star query = typical query pattern for star schema

**Example:** Total revenue in a given year (e.g. 2013) per product

Join of multiple dimension tables with fact table +

- **Selection (WHERE):** on attributes in dimension tables
- **Grouping (GROUP BY):** on attributes in dimension tables
- **Aggregation (SUM, AVG, COUNT, ... and HAVING-clause):** on attributes in fact table

# EXAMPLE: STAR-QUERY

## Total revenue in 2013 per product

```
select sum(revenue) as total, by p.ProductKey,  
p.name  
from Linitem l, Customer c, Product p, Date d  
where l.custKey = c.custKey  
and l.ProductKey = p.ProductKey  
and l.dateKey = d.dateKey  
And d.year = 2013  
group by p.ProductKey, p.name
```

# CLICKER QUESTION

**The following star schema is used to track user who borrowed which books over time**

***Dimensions:***

Book(bookId, title)

User(userId, name, DOB)

Genre(genreId, title)

***Fact Table:***

Lendings(bookId, userId, genreId, days)

# CLICKER QUESTION (CNT)

Book(bookId, title)

User(userId, name, DOB)

Genre(genreId, title)

Lendings(bookId, userId, genreId, days)

**Which SQL query returns the total number of books from the genre “Fantasy” for more than 90 days on average?**

A) SELECT g.genre, COUNT(\*)  
FROM BorrowedBooks bb, Books b , Genre g  
WHERE bb.bookID=b.bookID AND  
bb.genreID=g.genreID AND  
g.genre='Fantasy' AND  
bb.days > 90  
GROUP BY b.genre

B) SELECT genre, COUNT(\*)  
FROM BorrowedBooks bb, Genre g  
WHERE bb.genreID=g.genreID AND  
g.genre='Fantasy'  
HAVING AVG(bb.days) > 90

# CLICKER QUESTION (CNT)

Book(bookId, title)

User(userId, name, DOB)

Genre(genreId, title)

Lendings(bookId, userId, genreId, days)

**Which SQL query returns the total number of books from the genre “Fantasy” for more than 90 days on average?**

A) `SELECT g.genre, COUNT(*)  
FROM BorrowedBooks bb, Books b , Genre g  
WHERE bb.bookID=b.bookID AND  
bb.genreID=g.genreID AND  
g.genre='Fantasy' AND  
bb.days > 90  
GROUP BY b.genre`

B) `SELECT genre, COUNT(*)  
FROM BorrowedBooks bb, Genre g  
WHERE bb.genreID=g.genreID AND  
g.genre='Fantasy'  
HAVING AVG(bb.days) > 90`

# SQL-EXTENSIONS

SQL has **different extensions** to support **analytical queries**

**Rollup (Grouping Sets)/ Cube:** special grouping by different sets of dimensional attributes

**Top(k)/Limit:** Top-k results ordered by a given key figure (e.g., top-10 customer which produced maximal total revenue)

**Skyline:** Finding optimal along multiple dimensions (e.g., hotels that are cheap and are close to the beach)



# ROLLUP

**Rollup:** special grouping by different sets along a hierarchy

## **Example (IBM DB2):**

```
select sum(revenue) as total, region, country, city
from Linitem l, Customer c
where l.custKey = c.custKey
group by rollup(region, country, city)
```

Query groups result by the following attribute sets:  
(region), (region, country) and (region, country, city)

# EXAMPLE: ROLLUP

total	region	country	city	
1.435.789	Europe	-	-	(region)
232.199	Europe	France	-	(region, Country)
634.124	Europe	Germany	-	
119.566	Europe	Germany	Munich	(region, Country, city)
35.234	Europe	Germany	Mannheim	
...	...	...	...	
210.199	Europe	France	Paris	
...	...	...	...	

# GROUPING SETS

**Alternative for Rollup:** Grouping sets define the set of group-by attributes explicitly

## Example (Oracle):

```
select sum(total) as total, region, Country, city
from Linitem l, Customer c
where l.custKey = c.custKey
group by grouping sets((region, country),
                       (region, country, city))
```

Query groups result by the following attribute combinations:  
(region, Country) und (region, Country, city)

# EXAMPLE: GROUPING SETS

total	region	Country	city	
232.199	Europe	France	-	(region, country)
634.124	Europe	Germany	-	
119.566	Europe	Germany	Munich	(region, country, city)
35.234	Europe	Germany	Mannheim	
...	...	...	...	
210.199	Europe	France	Paris	
...	...	...	...	

# TOP(K) OR LIMIT

## Top-k/ LIMIT functionality:

- Sort aggregated result
- Limit result size by given k

## Example (PostgreSQL):

```
select sum(total) as total, region
from Litem l, Customer c
where l.custKey = c.custKey
group by region
order by total
limit 5;
```

# SKYLINE

**Skyline** is a **multi-dimensional top(k)**

- Skyline returns all **tuples that are not dominated by any other point** in the given set of dimensions
- Qualifying tuples also known as „**Pareto-Optimum**“


**Example: Hotels low distance to beach + low price**

```
select *
```

```
from hotels h
```

```
skyline of h.distance min, p.price min
```

# EXAMPLE: SKYLINE (HOTELS)



# SUMMARY

## **Data Modeling**

- Multi-dimensional Model / Cube
- Star Schema / Snowflake Schema
- Hierarchies

## **ETL-Processes**

## **SQL Extensions**

- ROLLUP / GROUPING SETS
- TOP(k)
- SKYLINE



# WHAT IS A GOOD DATA WAREHOUSE?

“A Data Warehouse is a

- **subject-oriented,**
- **integrated,**
- **non-volatile and time-variant**

collection of data in **support of managements decisions”**

**(W. H. Inmon, Building the Data Warehouse, 1996)**


# SUBJECT ORIENTED DATABASE

## Operational Databases:

- Are **application oriented** (e.g., bank accounts, loans, ...)
- Each DB manages only a **subset-of the overall data**

## Data Warehouses:

- **Global view** on all data about a given subject / entity (e.g., customer)
- **Not targeted** towards **one application**



# INTEGRATED DATABASE


A data warehouse **integrates (inconsistent) data** coming from different sources in a **consistent way**

DB 1 – m, f


DB 2 – male,  
female

DB 3 – 1,0

**inconsistent**



**consistent**




# NON-VOLATILE AND TIME-VARIANT

**Operational Databases** represents the **most up-to-date snapshot**

**Data Warehouses** represents the **history of all changes**:

- New Data is only appended / never updated
- All entries have a timestamp
- Comparison over time are possible



*updated constantly*



*snapshotted data*