

MapReduce

February 21, 2019

Data Science CSCI 1951A

Brown University

Instructor: Ellie Pavlick

HTAs: Wennie Zhang, Maulik Dang, Gurnaaz Kaur

Announcements

Announcements

- Next assignment released today—due March 7

Announcements

- Next assignment released today—due March 7
- First project checkin will be assigned next week, so this will be happening...in parallel (<- that was a hilarious pun, FYI)

Announcements

- Next assignment released today—due March 7
- First project checkin will be assigned next week, so this will be happening...in parallel (<- that was a hilarious pun, FYI)
- The HTAs are complaintless 👍

Announcements

- Next assignment released today—due March 7
- First project checkin will be assigned next week, so this will be happening...in parallel (<- that was a hilarious pun, FYI)
- The HTAs are complaintless 👍
- Questions? Concerns? Anything?

Today



MapReduce

MapReduce

- Functional-programming paradigm (inspired by LISP and friends)

MapReduce

- Functional-programming paradigm (inspired by LISP and friends)
- Two functions:

MapReduce

- Functional-programming paradigm (inspired by LISP and friends)
- Two functions:
 - Map: (in_key, in_value) -> list_of(out_key, intermediate_value)

MapReduce

- Functional-programming paradigm (inspired by LISP and friends)
- Two functions:
 - Map: (in_key, in_value) -> list_of(out_key, intermediate_value)
 - Reduce: (out_key, list_of(intermediate_value)) -> list_of(out_value)

MapReduce

- Functional-programming paradigm (inspired by LISP and friends)

Extremely Vague

- Two functions:

General

- Map: (in_key, in_value) -> list_of(out_key, intermediate_value)
- Reduce: (out_key, list_of(intermediate_value)) -> list_of(out_value)

MapReduce

distributed grep

distributed sort

web link-graph reversal

web access log stats

inverted index construction

document clustering

machine learning

statistical machine translation

...

Map Reduce

- One “master” scheduler which assigns tasks (mapping or reducing) to machines

Map Reduce

- One “master” scheduler which assigns tasks (mapping or reducing) to machines
- No shared state between machines—massively parallelizable

Map Reduce

- One “master” scheduler which assigns tasks (mapping or reducing) to machines
- No shared state between machines—massively parallelizable
- Assume very high failure rates on workers

Counting Words

Documents

hello world

oh hi there
world

why hello
there ,
world

world ! how
the hell are
ya ?

Counting Words

Documents

hello world

oh hi there
world

why hello
there ,
world

world ! how
the hell are
ya ?

hello 2
world 4
oh 1
hi 1
there 2
why 1
! 1
how 1
...

19

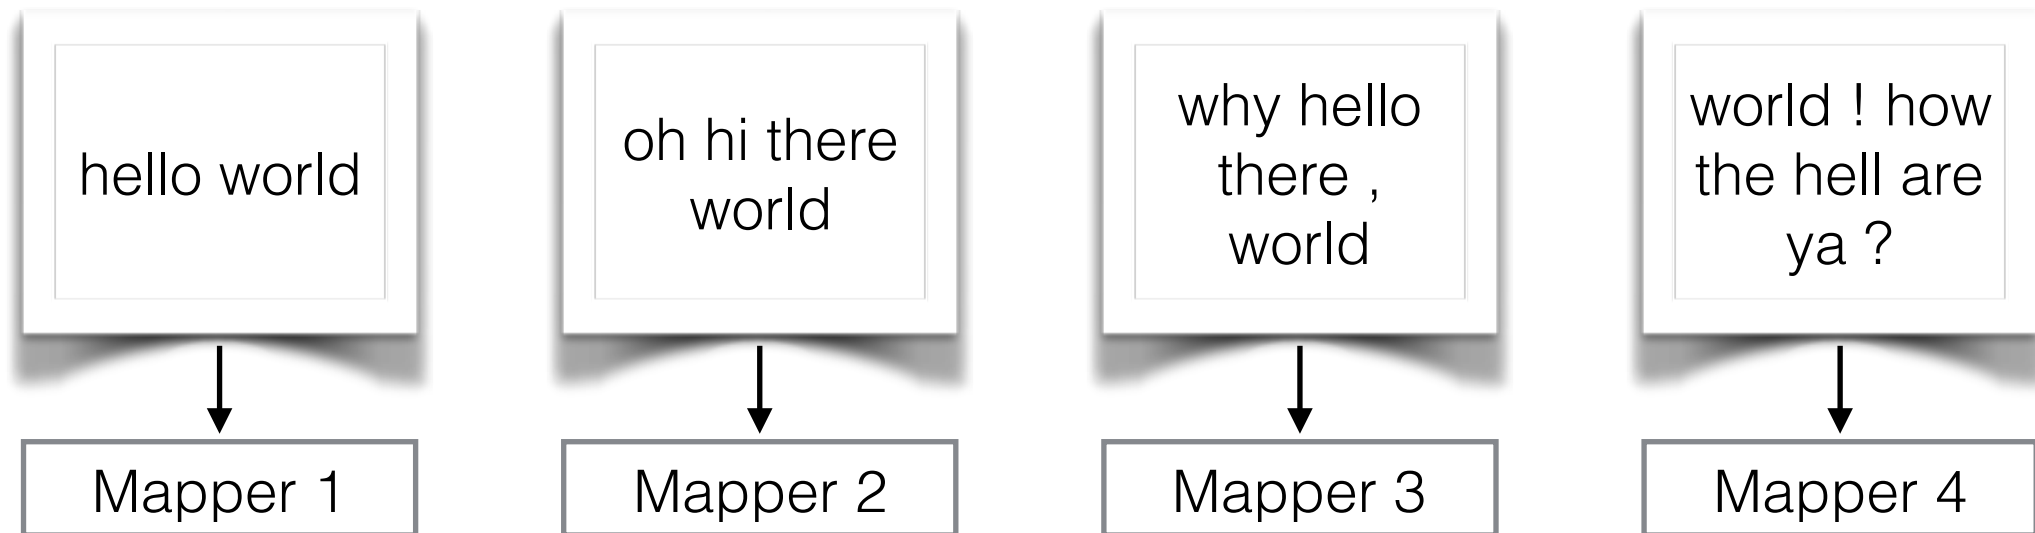
Counts for
each word

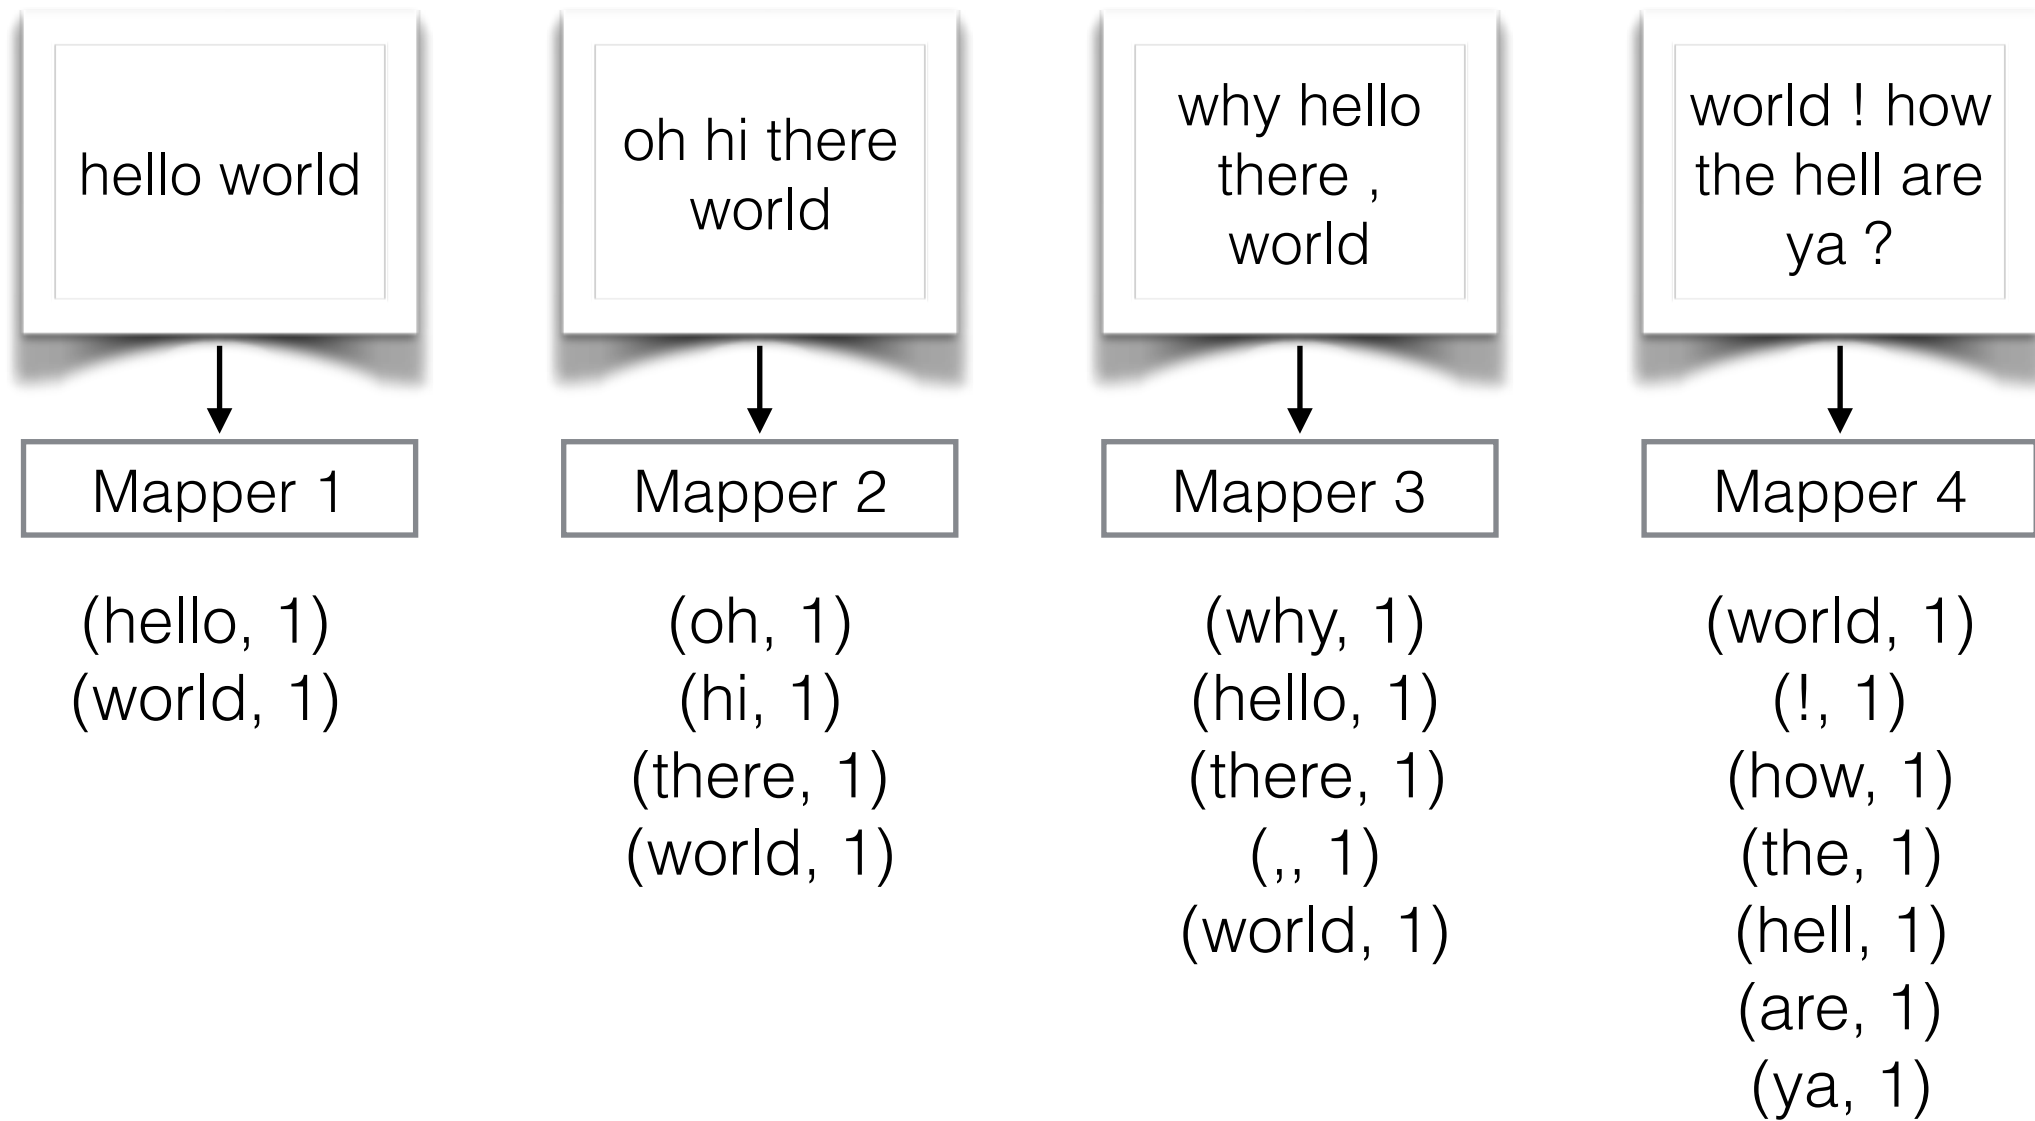
hello world

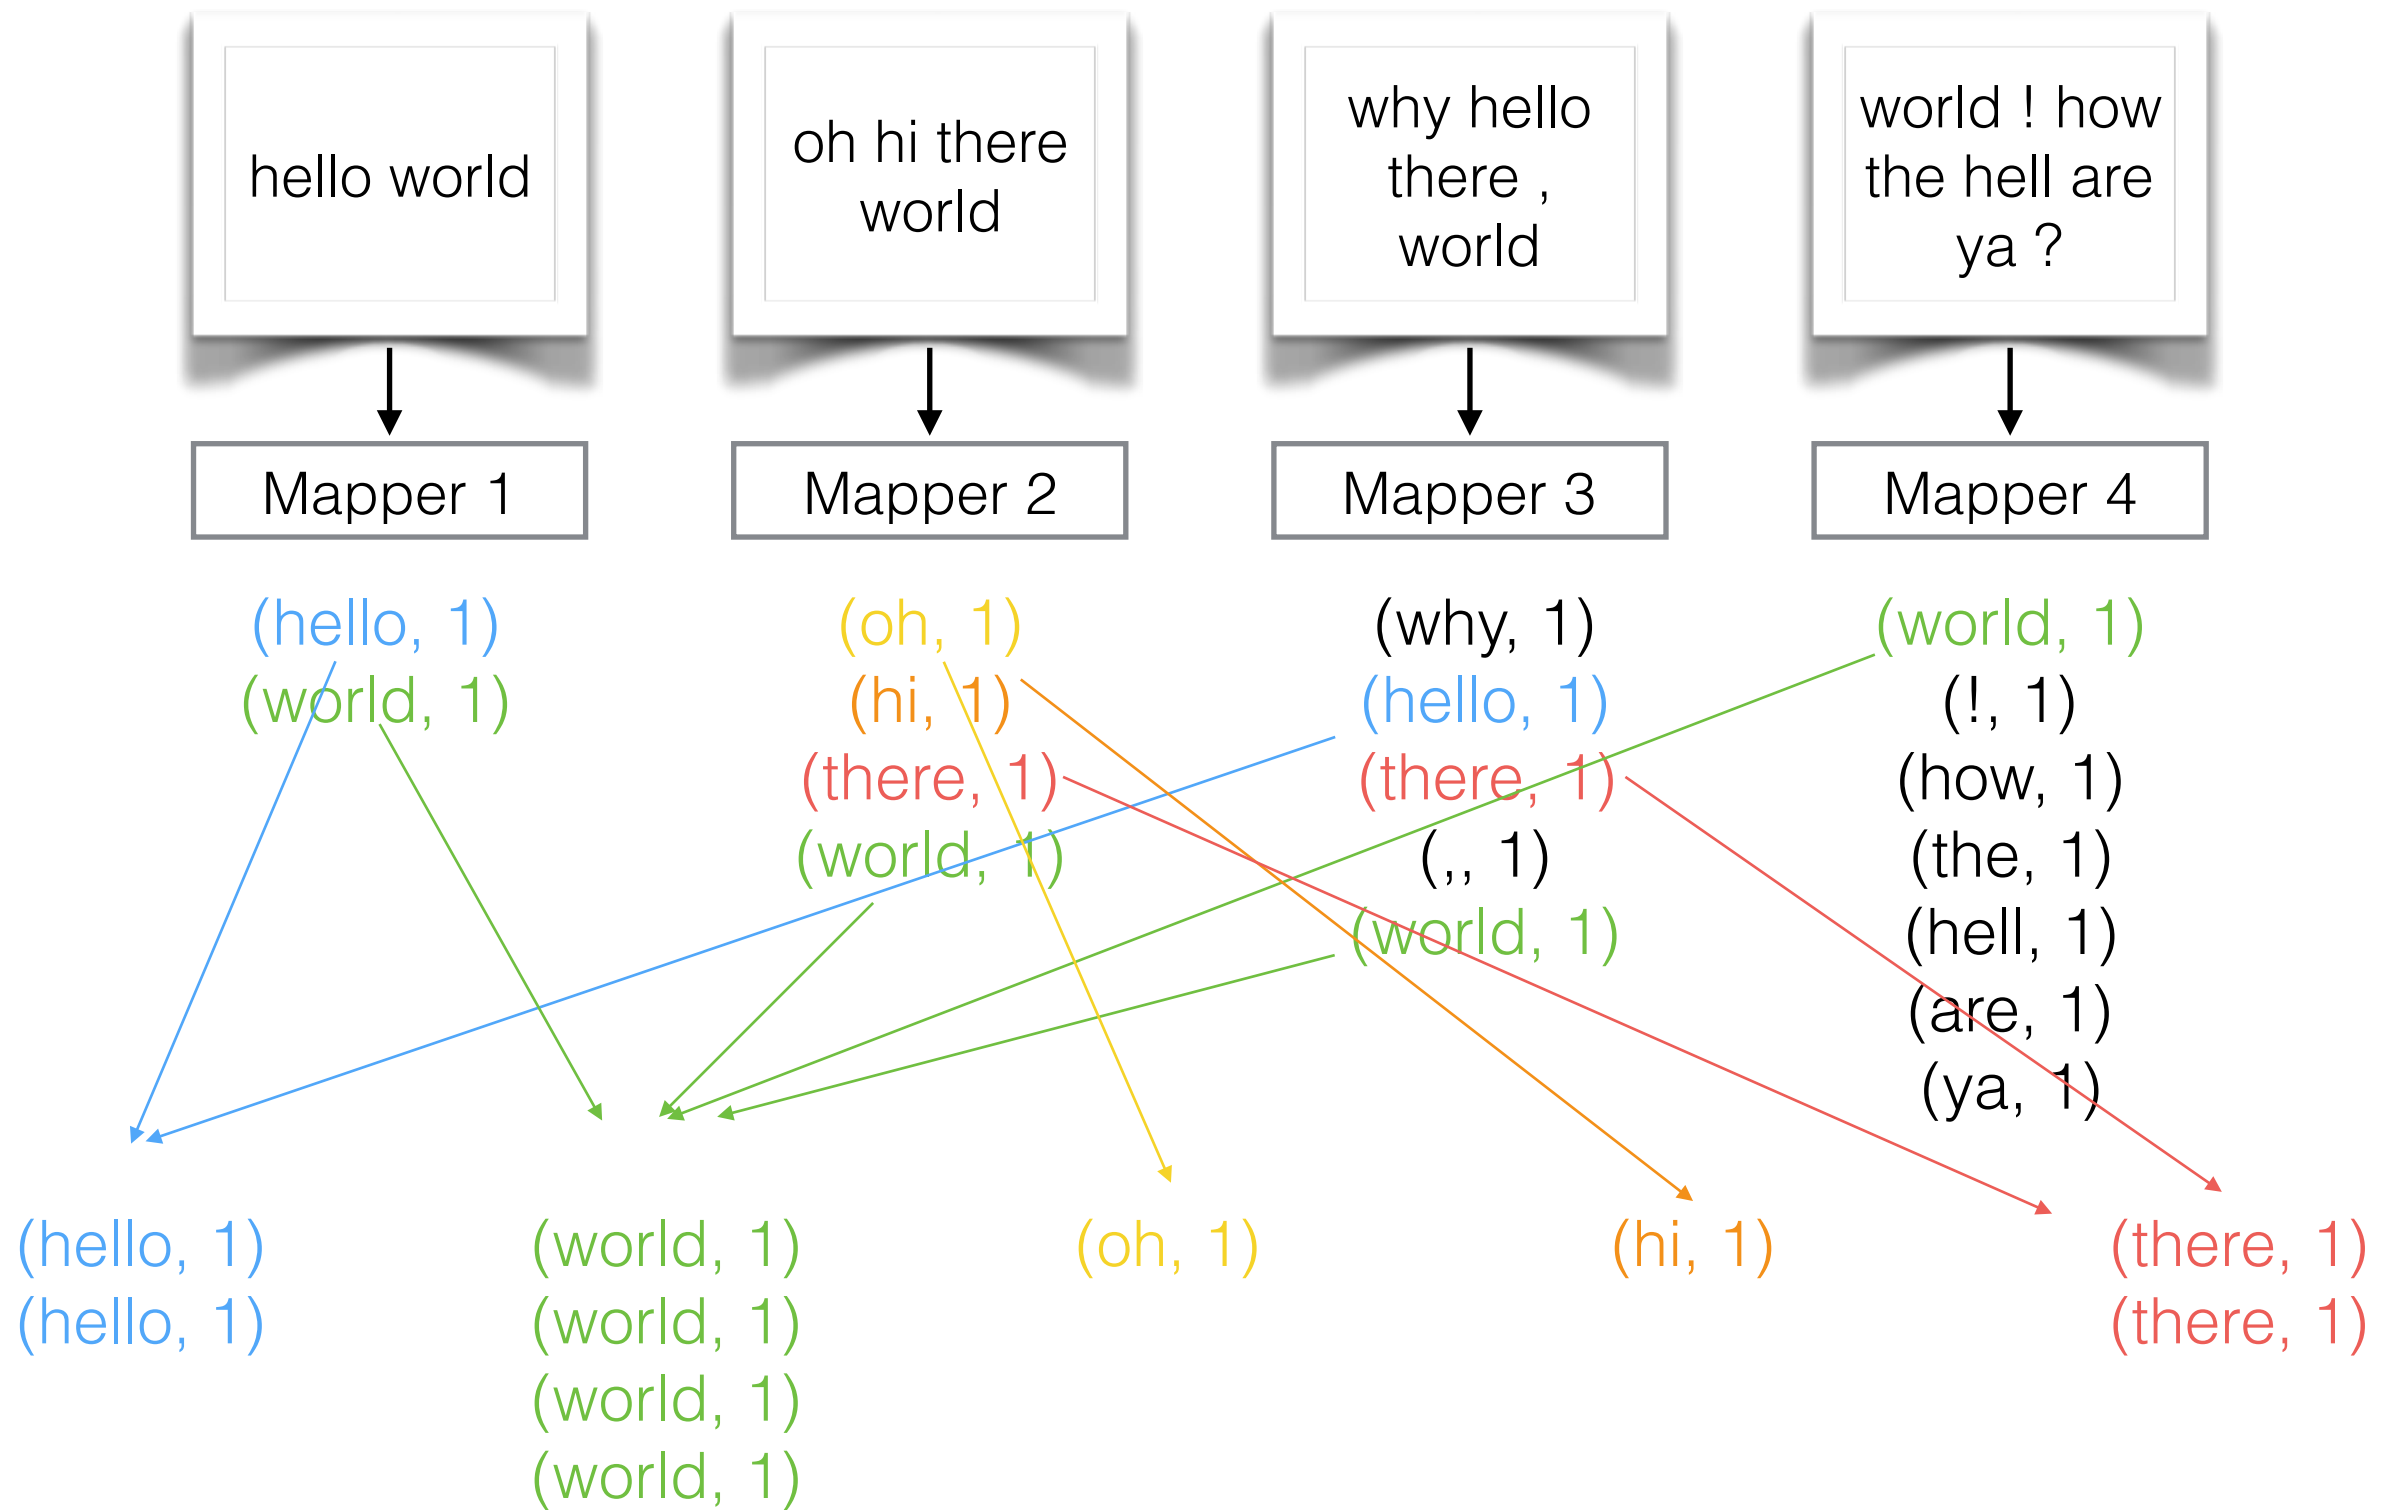
oh hi there
world

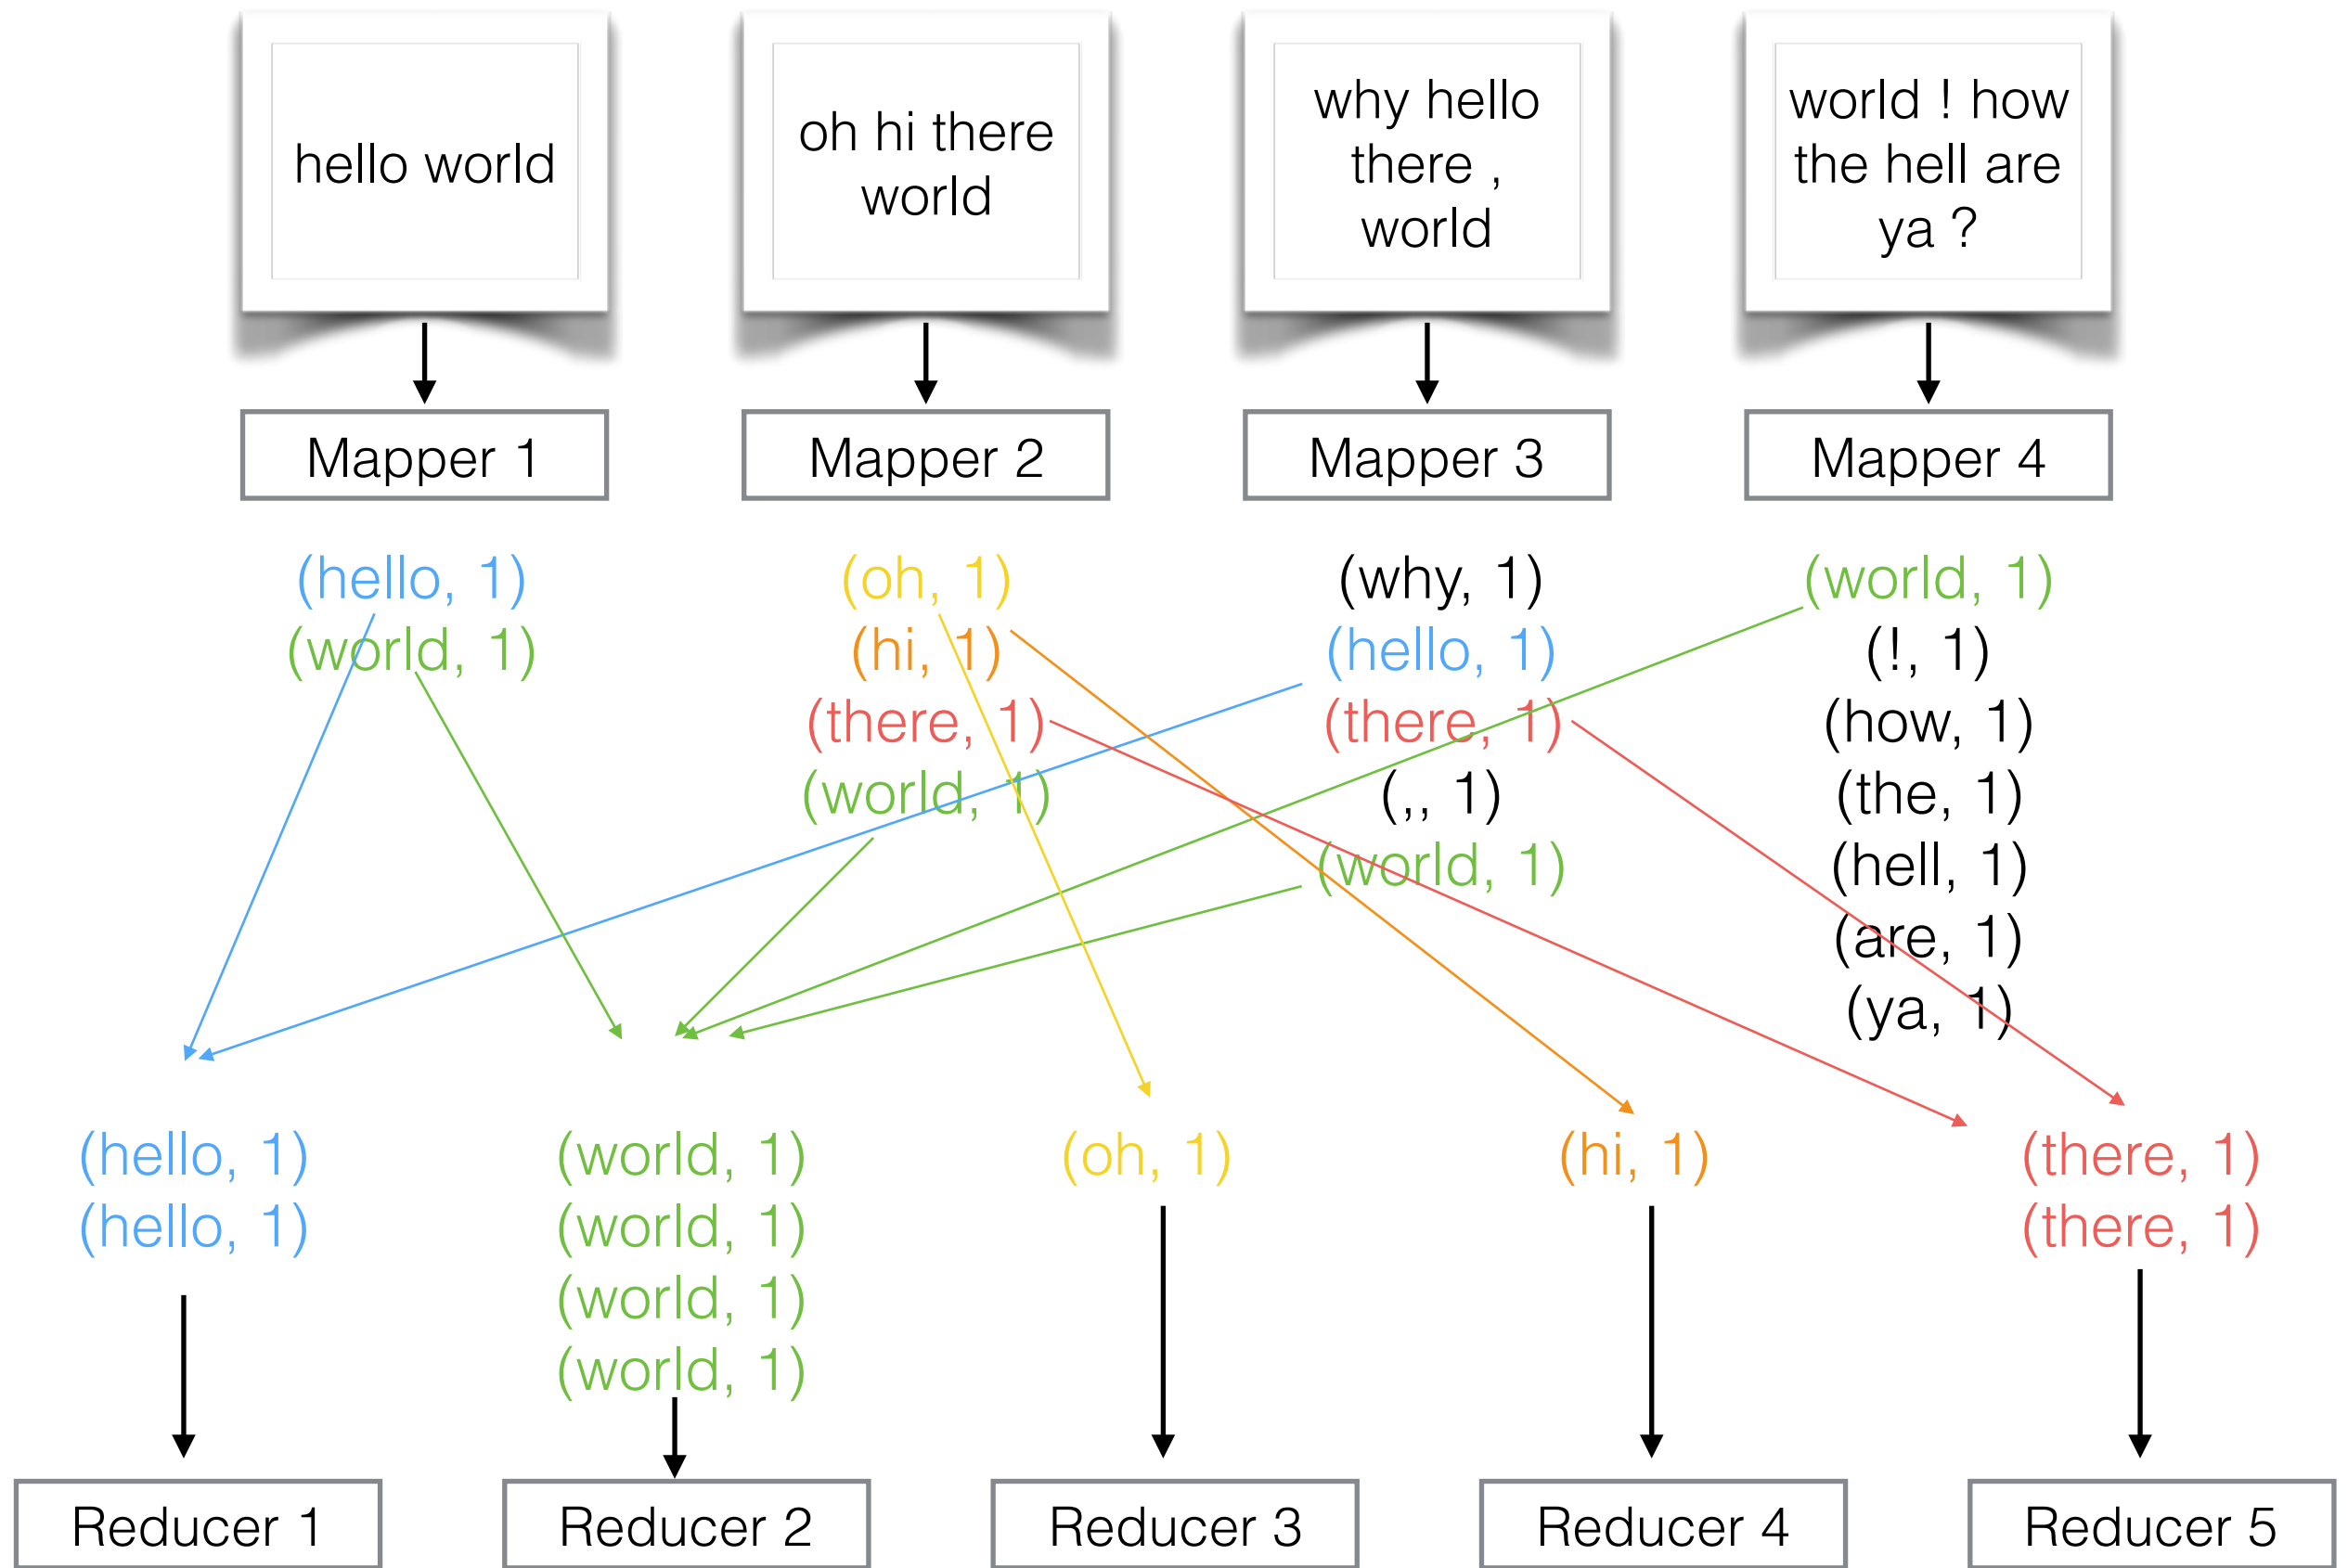
why hello
there ,
world

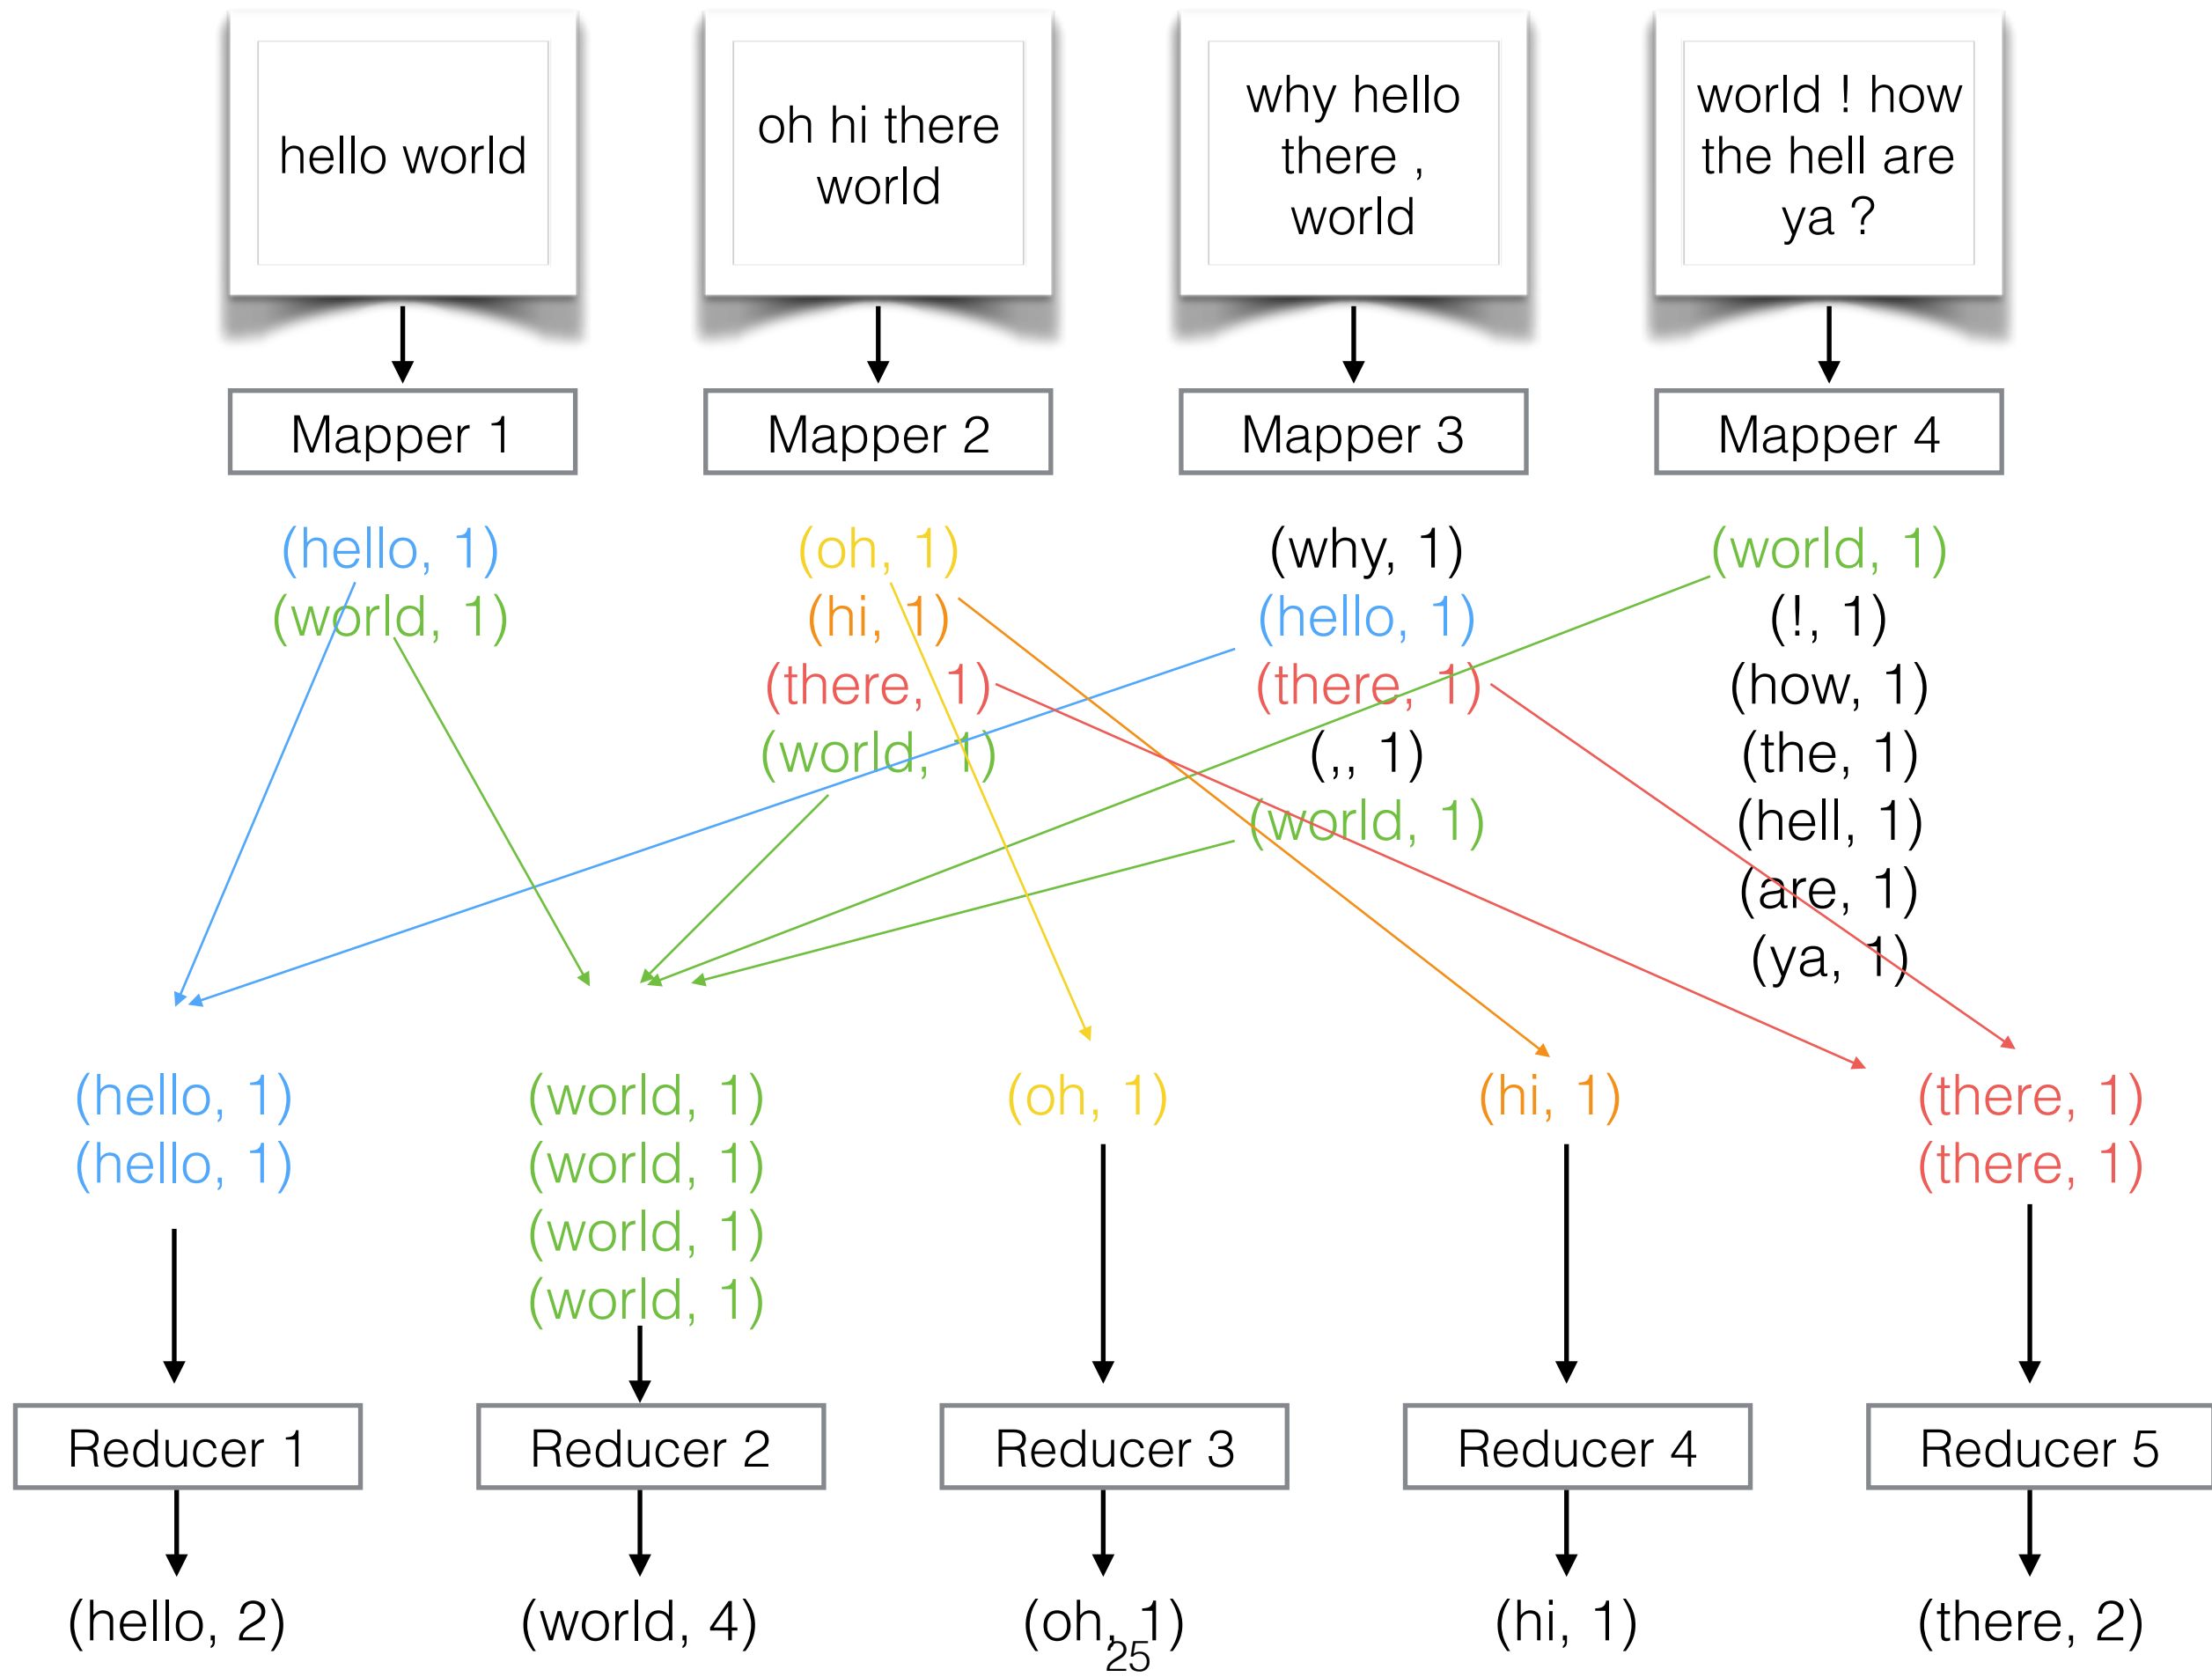
world ! how
the hell are
ya ?














Mapper 1

Mapper 2

Mapper 3

Mapper 4

(hello, 1)
(world, 1)

(oh, 1)
(hi, 1)
(there, 1)
(world, 1)

(why, 1)
(hello, 1)
(there, 1)
(,, 1)
(world, 1)

(world, 1)
(!, 1)
(how, 1)
(the, 1)
(hell, 1)
(are, 1)
(ya, 1)

(hello, 1)
(hello, 1)

(world, 1)
(world, 1)
(world, 1)
(world, 1)

(oh, 1)

(hi, 1)

(there, 1)
(there, 1)

Reducer 1

Reducer 2

Reducer 3

Reducer 4

Reducer 5

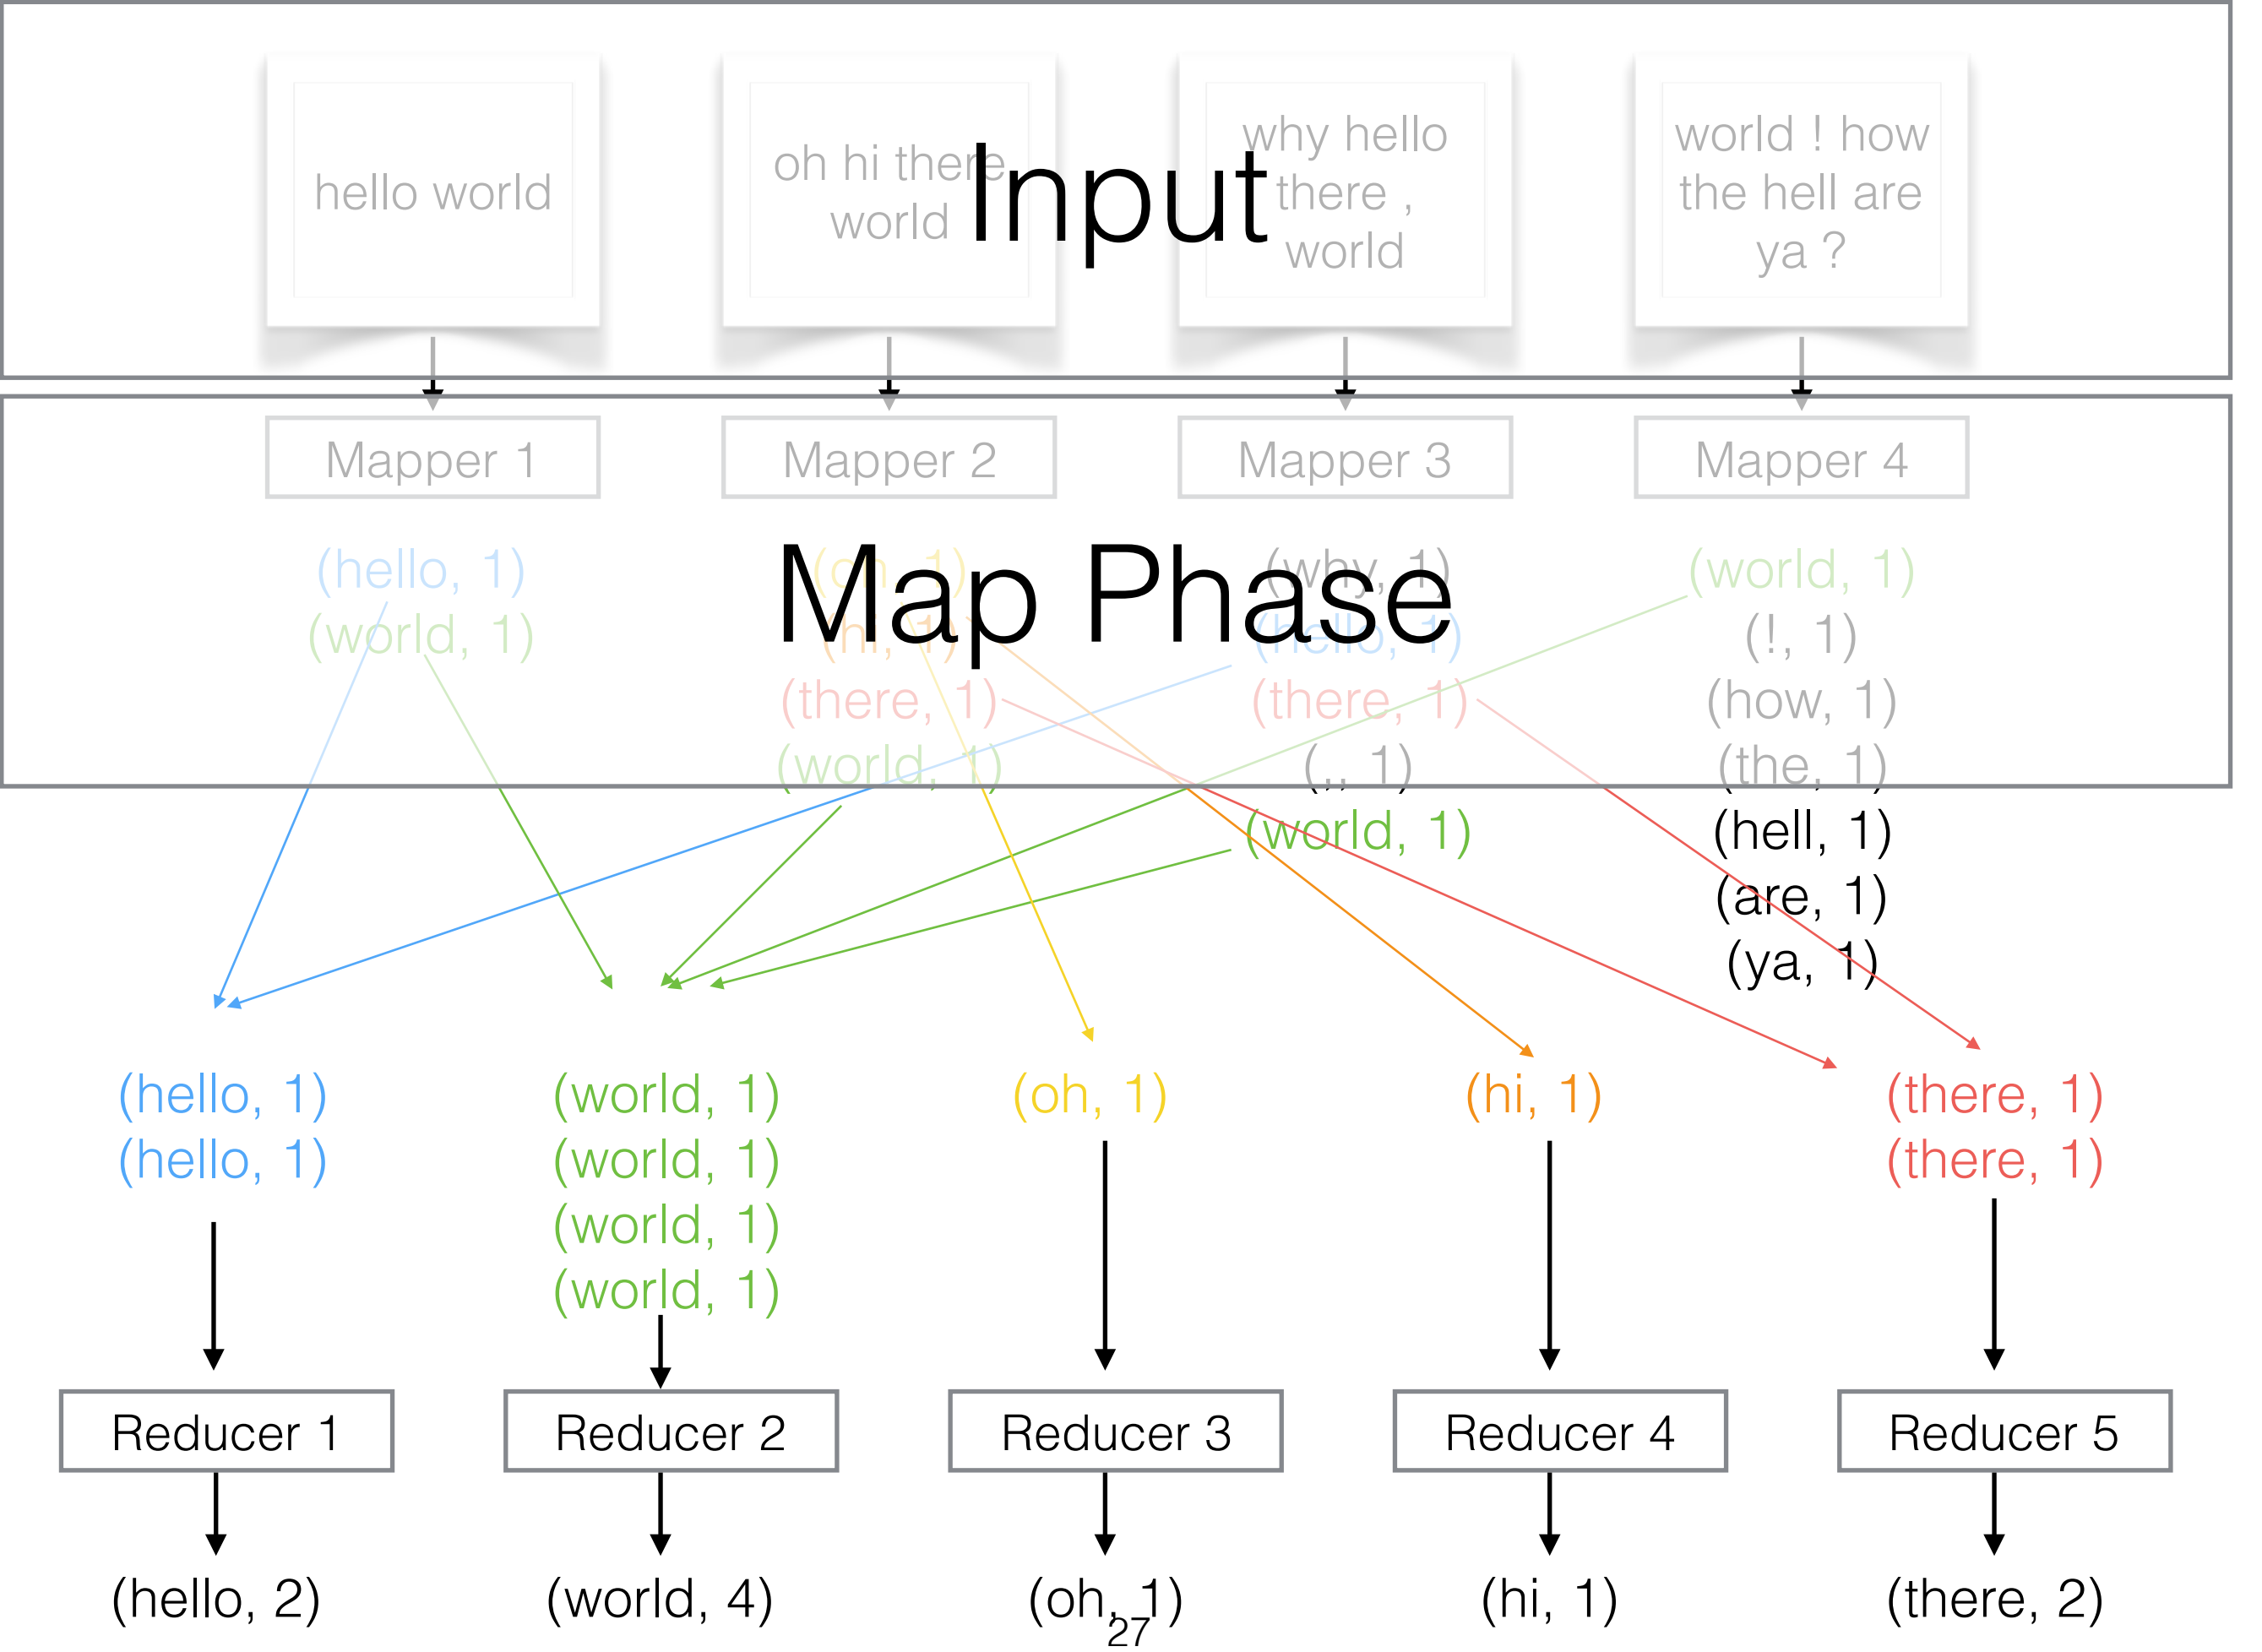
(hello, 2)

(world, 4)

(oh, 1)
26

(hi, 1)

(there, 2)



hello world

oh hi there
world

why hello
there ,
world

world ! how
the hell are
ya ?

Input

Mapper 1

Mapper 2

Mapper 3

Mapper 4

Map Phase

(hello, 1)
(world, 1)

(oh, 1)
(hi, 1)
(there, 1)
(world, 1)

(why, 1)
(hello, 1)
(there, 1)
(, , 1)

(world, 1)
(!, 1)
(how, 1)
(the, 1)

(hell, 1)
(are, 1)
(ya, 1)

Shuffle Phase ("Group By")

(hello, 1)
(hello, 1)

(world, 1)
(world, 1)
(world, 1)
(world, 1)

(oh, 1)

(hi, 1)

(there, 1)
(there, 1)

Reducer 1

Reducer 2

Reducer 3

Reducer 4

Reducer 5

(hello, 2)

(world, 4)

(oh, 1)
28

(hi, 1)

(there, 2)

hello world

oh hi there
world

why hello
there ,
world

world ! how
the hell are
ya ?

Input

Mapper 1

Mapper 2

Mapper 3

Mapper 4

Map Phase

(hello, 1)
(world, 1)

(oh, 1)
(hi, 1)
(there, 1)
(world, 1)

(why, 1)
(hello, 1)
(there, 1)
(, , 1)

(world, 1)
(!, 1)
(how, 1)
(the, 1)
(hell, 1)
(are, 1)
(ya, 1)

NOT! Sort (No guarantee about order of values...)

Shuffle Phase ("Group By")

(hello, 1)
(hello, 1)

(world, 1)
(world, 1)
(world, 1)
(world, 1)

(oh, 1)

(hi, 1)

(there, 1)
(there, 1)

Reducer 1

Reducer 2

Reducer 3

Reducer 4

Reducer 5

(hello, 2)

(world, 4)

(oh, 1)
29

(hi, 1)

(there, 2)

hello world

oh hi there
world

why hello
there ,
world

world ! how
the hell are
ya ?

Input

Mapper 1

Mapper 2

Mapper 3

Mapper 4

Map Phase

(hello, 1)
(world, 1)

(oh, 1)
(hi, 1)
(there, 1)
(world, 1)

(why, 1)
(hello, 1)
(there, 1)
(,, 1)
(world, 1)

(world, 1)
(!, 1)
(how, 1)
(the, 1)
(hell, 1)
(are, 1)
(ya, 1)

Shuffle Phase ("Group By")

(hello, 1)
(hello, 1)

(world, 1)
(world, 1)
(world, 1)
(world, 1)

(oh, 1)

(hi, 1)

(there, 1)
(there, 1)

Reducer 1

Reducer 2

Reducer 3

Reducer 4

Reducer 5

Reduce Phase

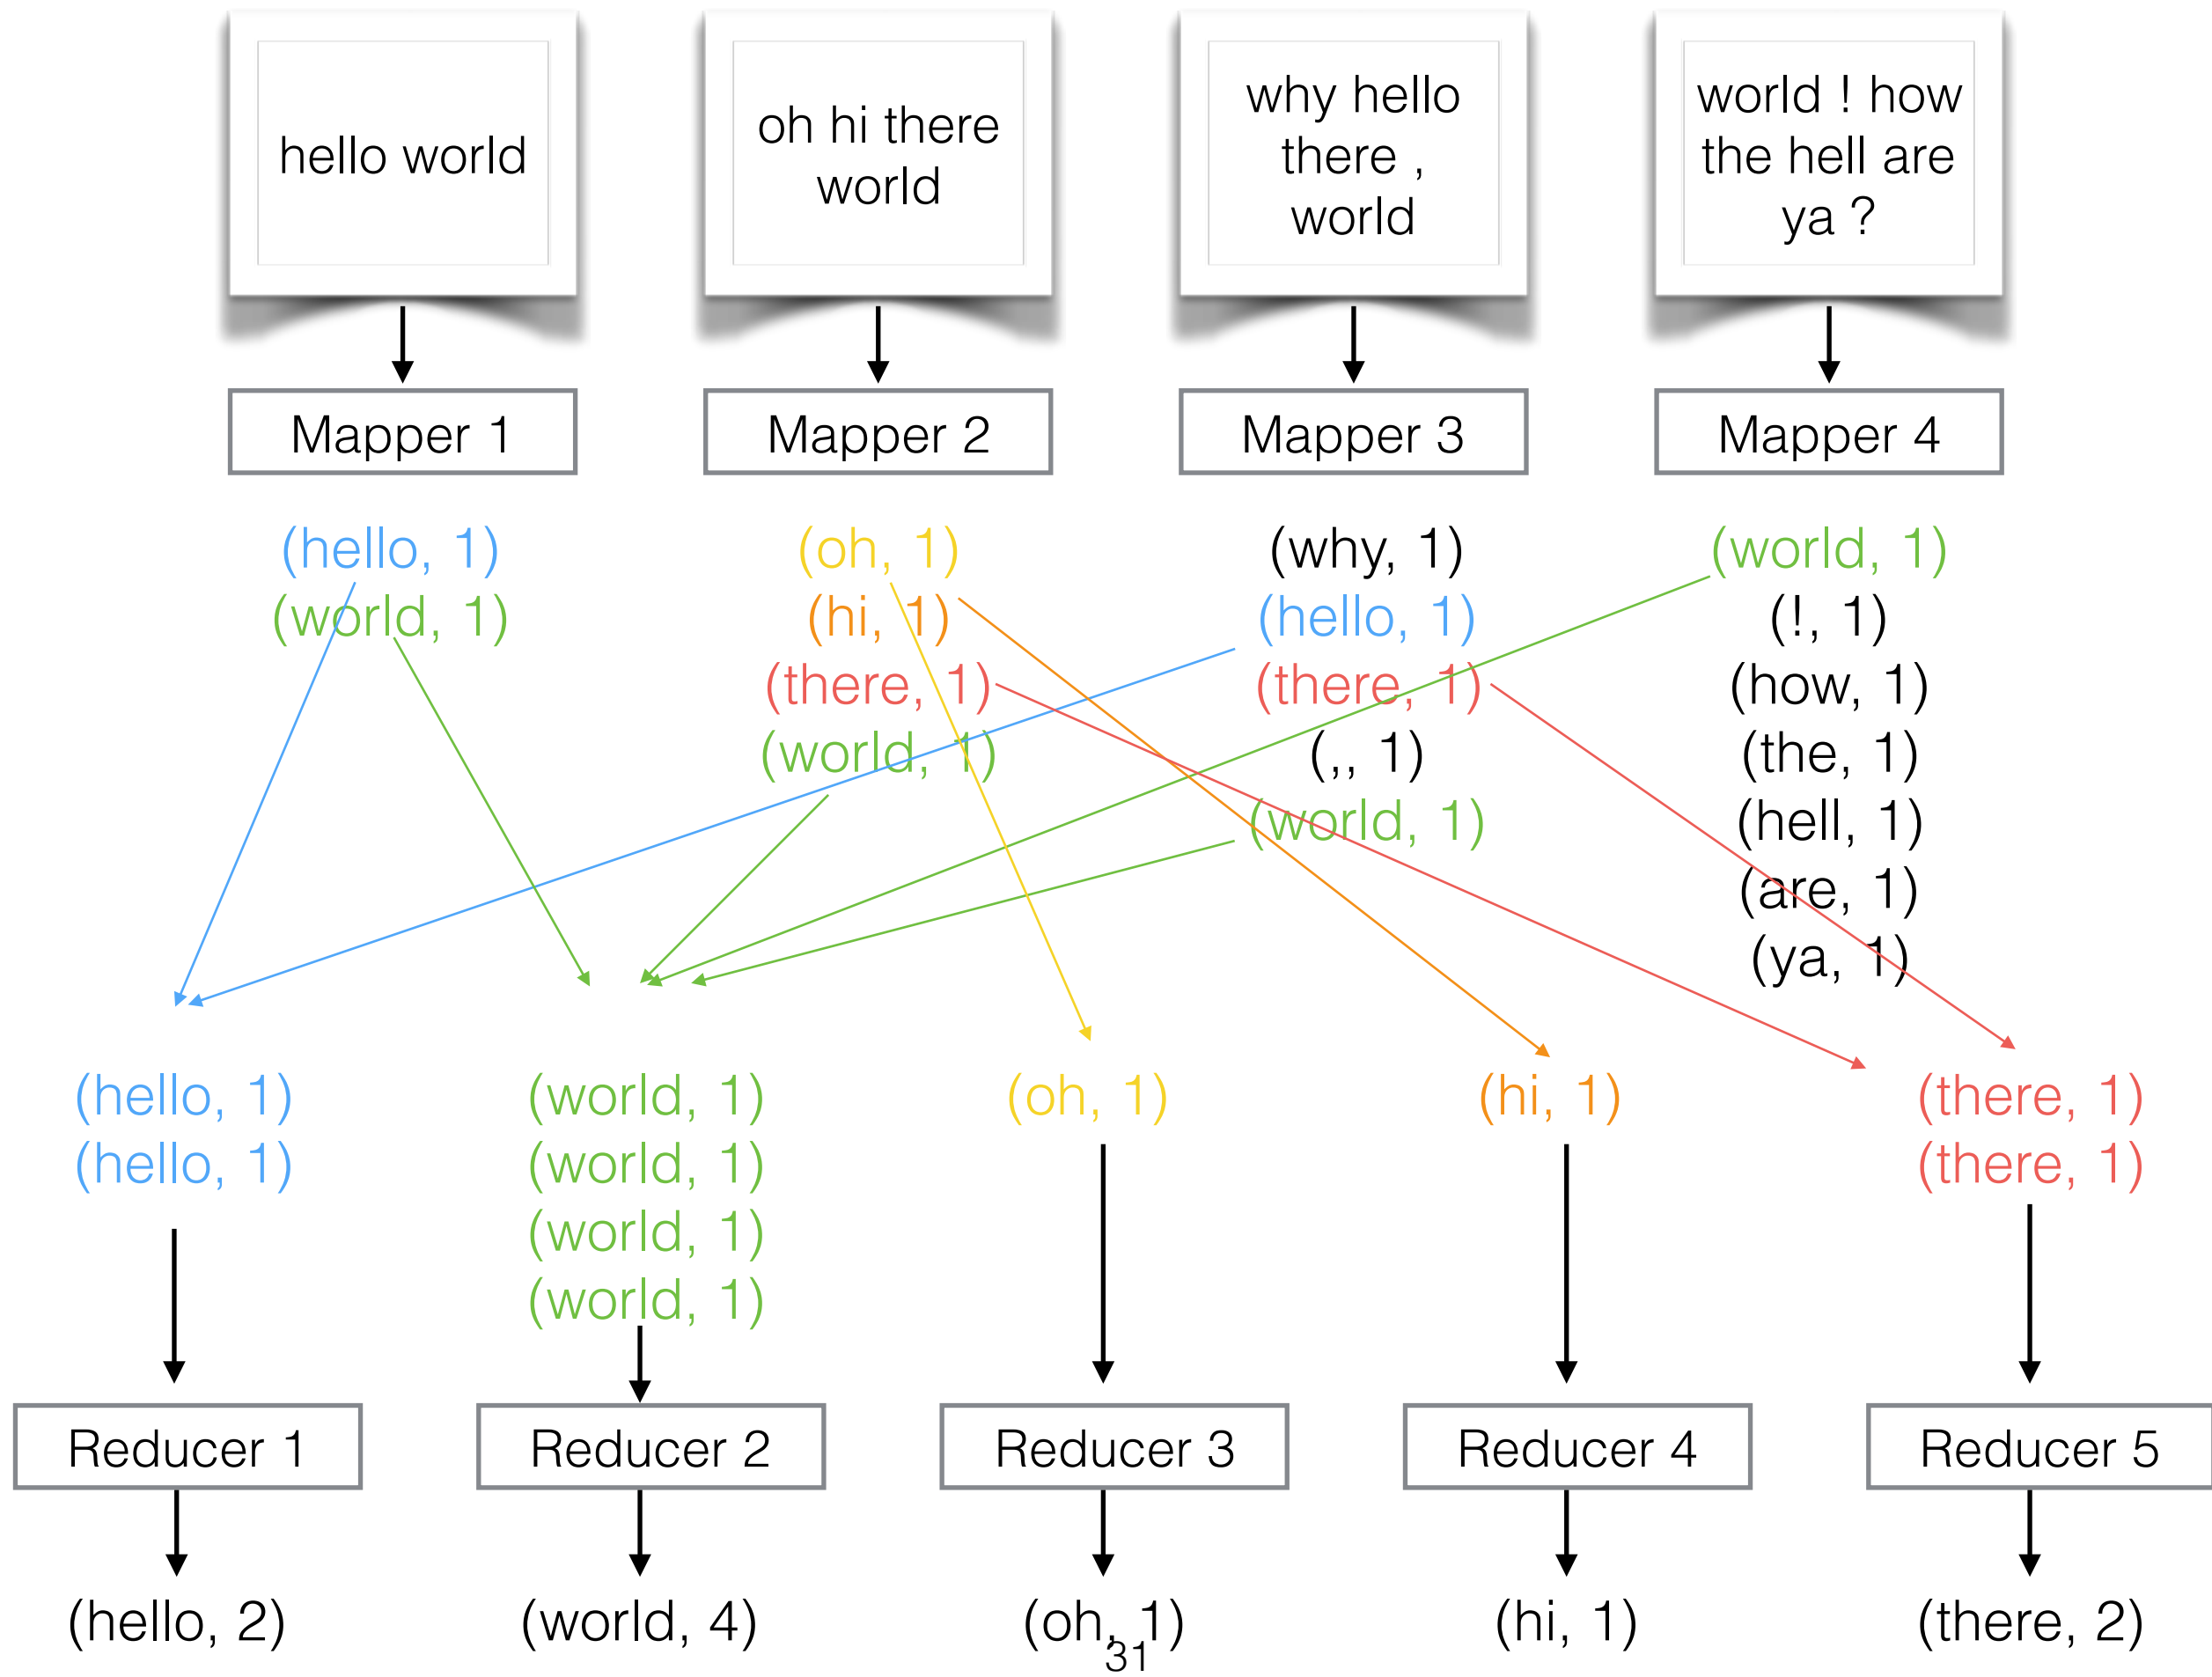
(hello, 2)

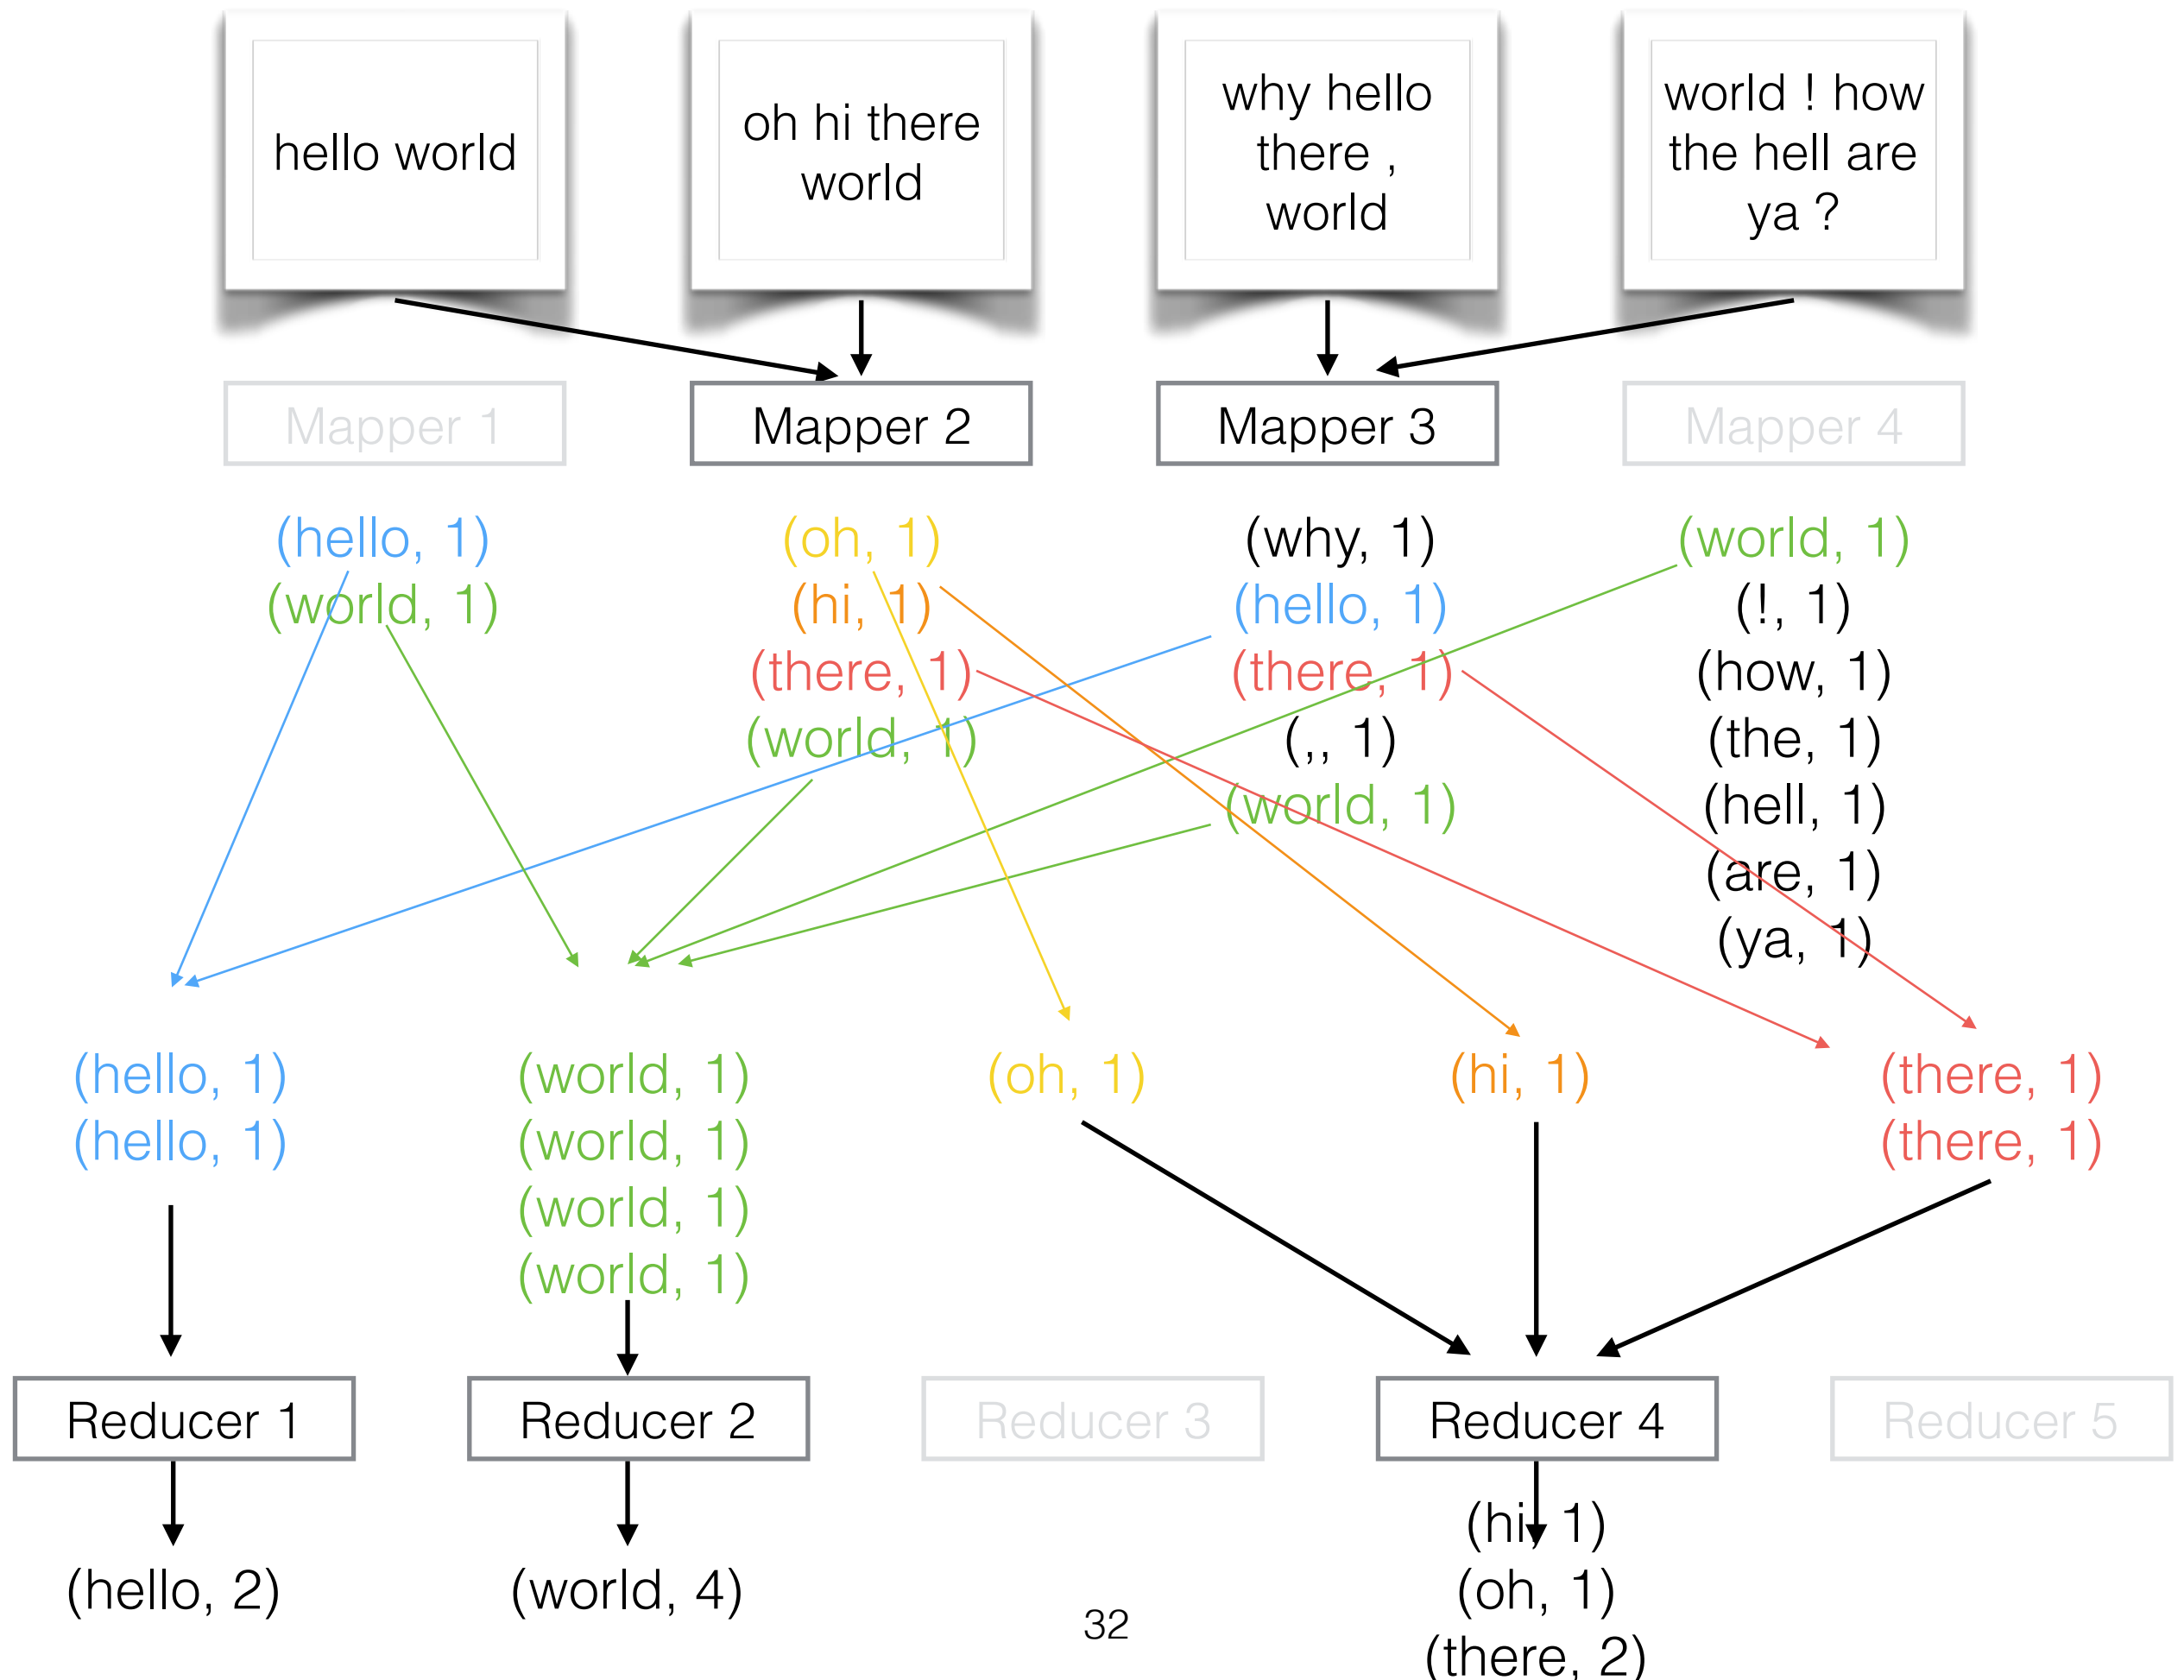
(world, 4)

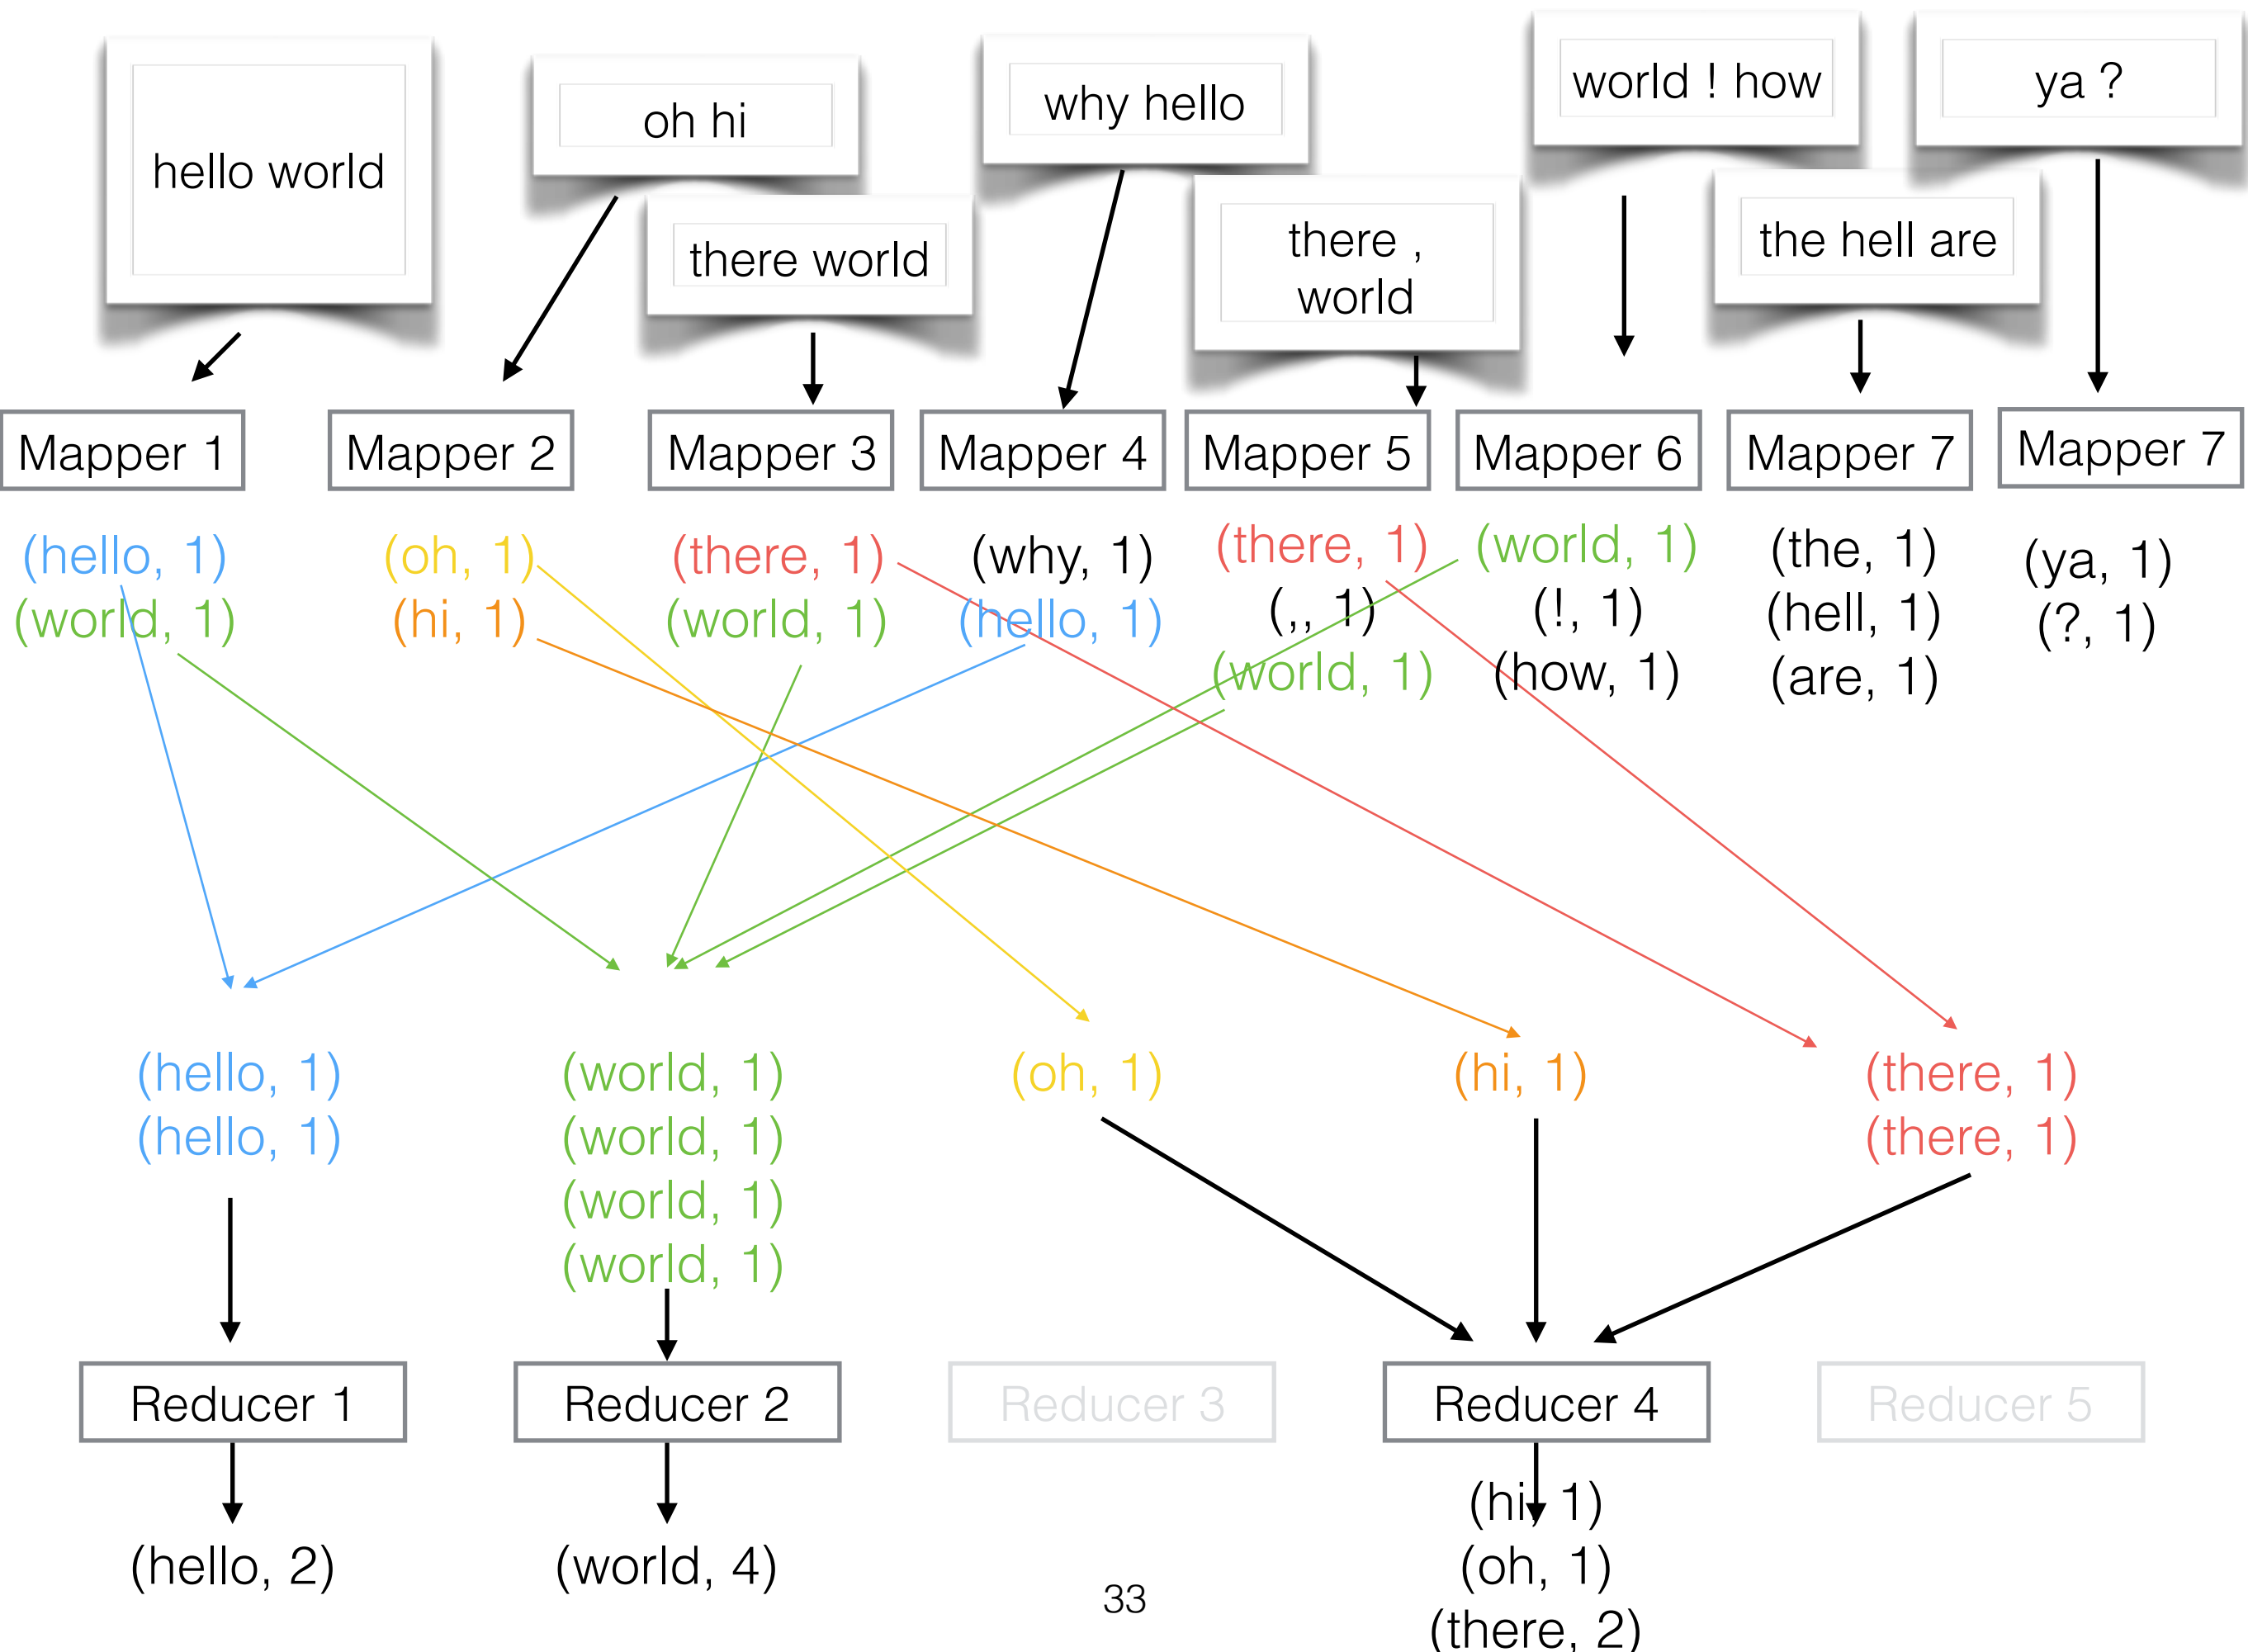
(oh, 1)

(hi, 1)

(there, 2)







Map Reduce

```
//define your mapper function(s)
def MapFn: (String, String) -> (String, Int) {
  TODO;
}

//define your reduce function(s)
def ReduceFn: (String, Int) -> (String, Int) {
  TODO;
}

//define your pipeline
Table<String, String> table = read(table_path)
Table<String, Int> output =
    table.MapFn().ReduceFn();
write(output)
```

Map Reduce

```
//define your mapper function(s)
def MapFn: (String, String) -> (String, Int) {
  TODO;
}
```

```
//define your reduce function
def ReduceFn: (String, Int) -> (String, Int) {
  TODO;
}
```

table

DocID	Text
1	hello world
2	oh hi there world
3	why hello there , world
4	world ! how the hell are ya ?

```
//define your pipeline
Table<String, String> table = read(table_path)
Table<String, Int> output =
  table.MapFn().ReduceFn();
write(output)
```

Map Reduce

```
//define your mapper function(s)
def MapFn: (String, String) -> (String, Int) {
  TODO;
}
```

output

Word	Count
hello	2
world	4
oh	1
hi	1
there	2

table

DocID	Text
1	hello world
2	oh hi there world
3	why hello there , world
4	world ! how the hell are ya ?

```
//define your pipeline
Table<String, String> table = read(table_path)
Table<String, Int> output =
  table.MapFn().ReduceFn();
write(output)
```

Map Reduce

```
//define your mapper function(s)
def MapFn: (String, String) -> (String, Int) {
  TODO;
}
```

```
//define your reduce function(s)
def ReduceFn: (String, Int) -> (String, Int) {
  TODO;
}
```

Lots of data types:

String, Int, Float, Tuples thereof

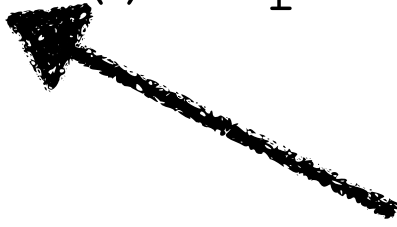
```
//define your pipeline
Table<String, String> table = read(table_path)
Table<String, Int> output =
  table.MapFn().ReduceFn();
write(output)
```

Map Reduce

```
// enumerate occurrences of each word, with
// count of 1
def MapFn: (String, String) -> (String, Int) {
    for w in input.value().split() {
        emit(w, 1);
    }
}
```

Map Reduce

```
// enumerate occurrences of each word, with  
// count of 1  
def MapFn: (String, String) -> (String, Int) {  
    for w in input.value().split() {  
        emit(w, 1);  
    }  
}
```



String



Map Reduce

```
// sum the total counts of each word
def ReduceFn:(String, Int) -> (String, Int) {
    sum = 0;
    for c in input.value() {
        sum += c;
    }
    emit(input.key(), sum);
}
```


Map Reduce

```
// sum the total counts of each word
def ReduceFn:(String, Int) -> (String, Int){
    sum = 0;
    for c in input.value() ← List of ints (counts)
        sum += c;
    }
    emit(input.key(), sum);
}
```

Map Reduce

```
// sum the total counts of each word
def ReduceFn: (String, Int) -> (String, Int) {
  sum = 0;
  for c in input.value()  List of ints (counts)
    sum += c;
  }
  emit(input.key()  the word, sum)
}
```

Find the number of occurrences of each word?

```
// enumerate occurrences of each word
// with count of 1
def MapFn: (String, String) -> (String, Int) {
    for w in input.split() {
        emit(w, 1);
    }
}

// sum the total counts of each word
def ReduceFn: (String, Int) -> (String, Int) {
    sum = 0;
    for (w, c) in input { sum += c; }
    emit(w, sum);
}

// define your pipeline
def main() {
    Table<String, String> table = read(table_path)
    Table<String, Int> output =
        table.MapFn().ReduceFn();
    write(output)
}
```

Input: String



Map: output (word, 1)
for every word.



Reduce: Sum counts
for each word

(non)Clicker Question!

Find the number of unique documents that each word occurs in?

(non)Clicker Question!

Find the number of unique documents that each word occurs in?

```
// enumerate occurrences of each word
// with count of 1
def MapFn1: String -> (String, Int) {
    ???
}
def ReduceFn1: String -> (String, Int) {
    ???
}
// sum the total counts of each word
def ReduceFn2: (String, Int) -> (String, Int) {
    ???
}
// define your pipeline
def main() {
    Table<String, String> table = read(table_path)
    Table<String, Int> output =
        table.MapFn1().ReduceFn1().ReduceFn2();
    write(output)
}
```

(non)Clicker Question!

Find the number of unique documents that each word occurs in?

```
// enumerate occurrences of each word
// with count of 1
def MapFn1: String -> (String, Int) {
  ???
}
def ReduceFn1: String -> (String, Int) {
  ???
}
// sum the total counts of each word
def ReduceFn2: (String, Int) -> (String, Int) {
  ???
}
// define your pipeline
def main() {
  Table<String, String> table = read(table_path)
  Table<String, Int> output =
    table.MapFn1().ReduceFn1().ReduceFn2();
  write(output)
}
```

No using sets!

(non)Clicker Question!

Find the number of unique documents that each word occurs in?

```
// enumerate occurrences of each word
// with count of 1
def MapFn1: String -> (String, Int) {
  ???
}
def ReduceFn1: String -> (String, Int) {
  ???
}
// sum the total counts of each word
def ReduceFn2: (String, Int) -> (String, Int) {
  ???
}
// define your pipeline
def main() {
  Table<String, String> table = read(table_path)
  Table<String, Int> output =
    table.MapFn1().ReduceFn1().ReduceFn2()
  write(output)
}
```

No using sets!



D1

hello world,
just saying
hello

D2

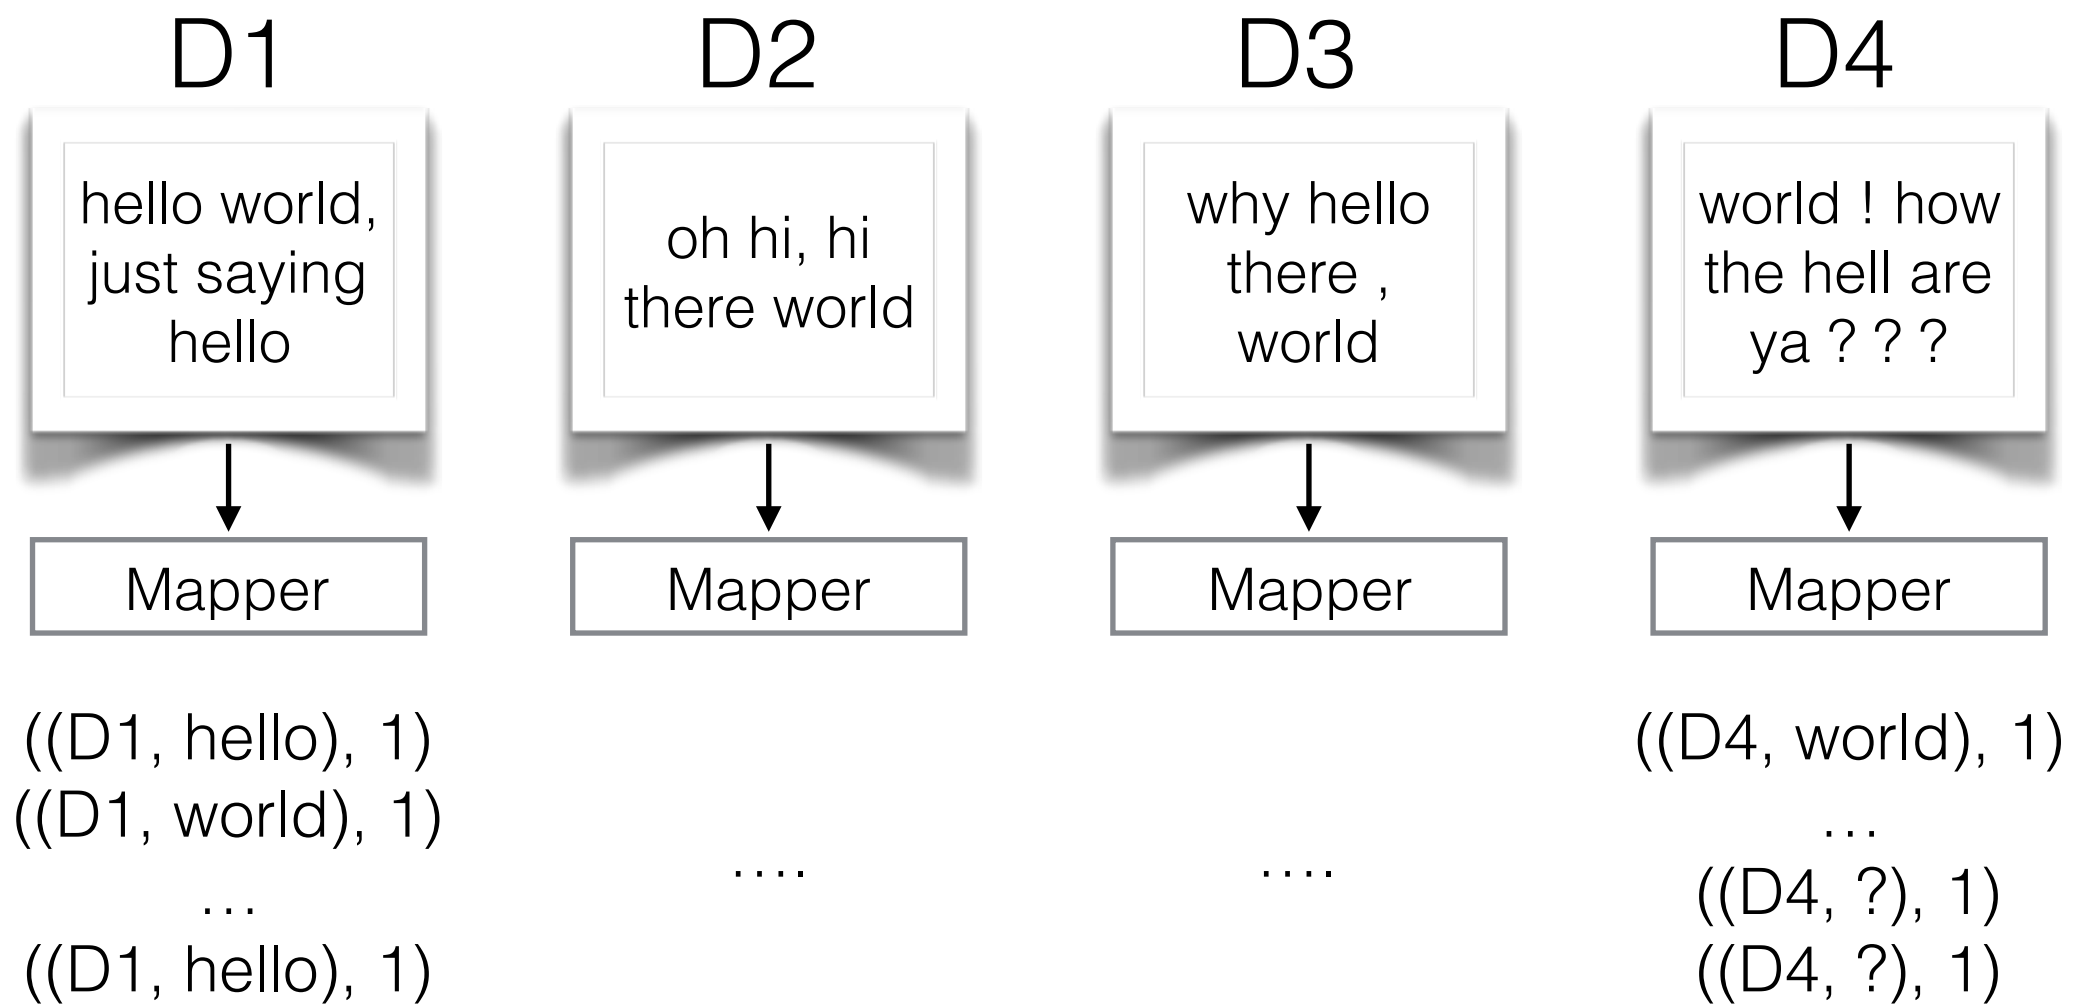
oh hi, hi
there world

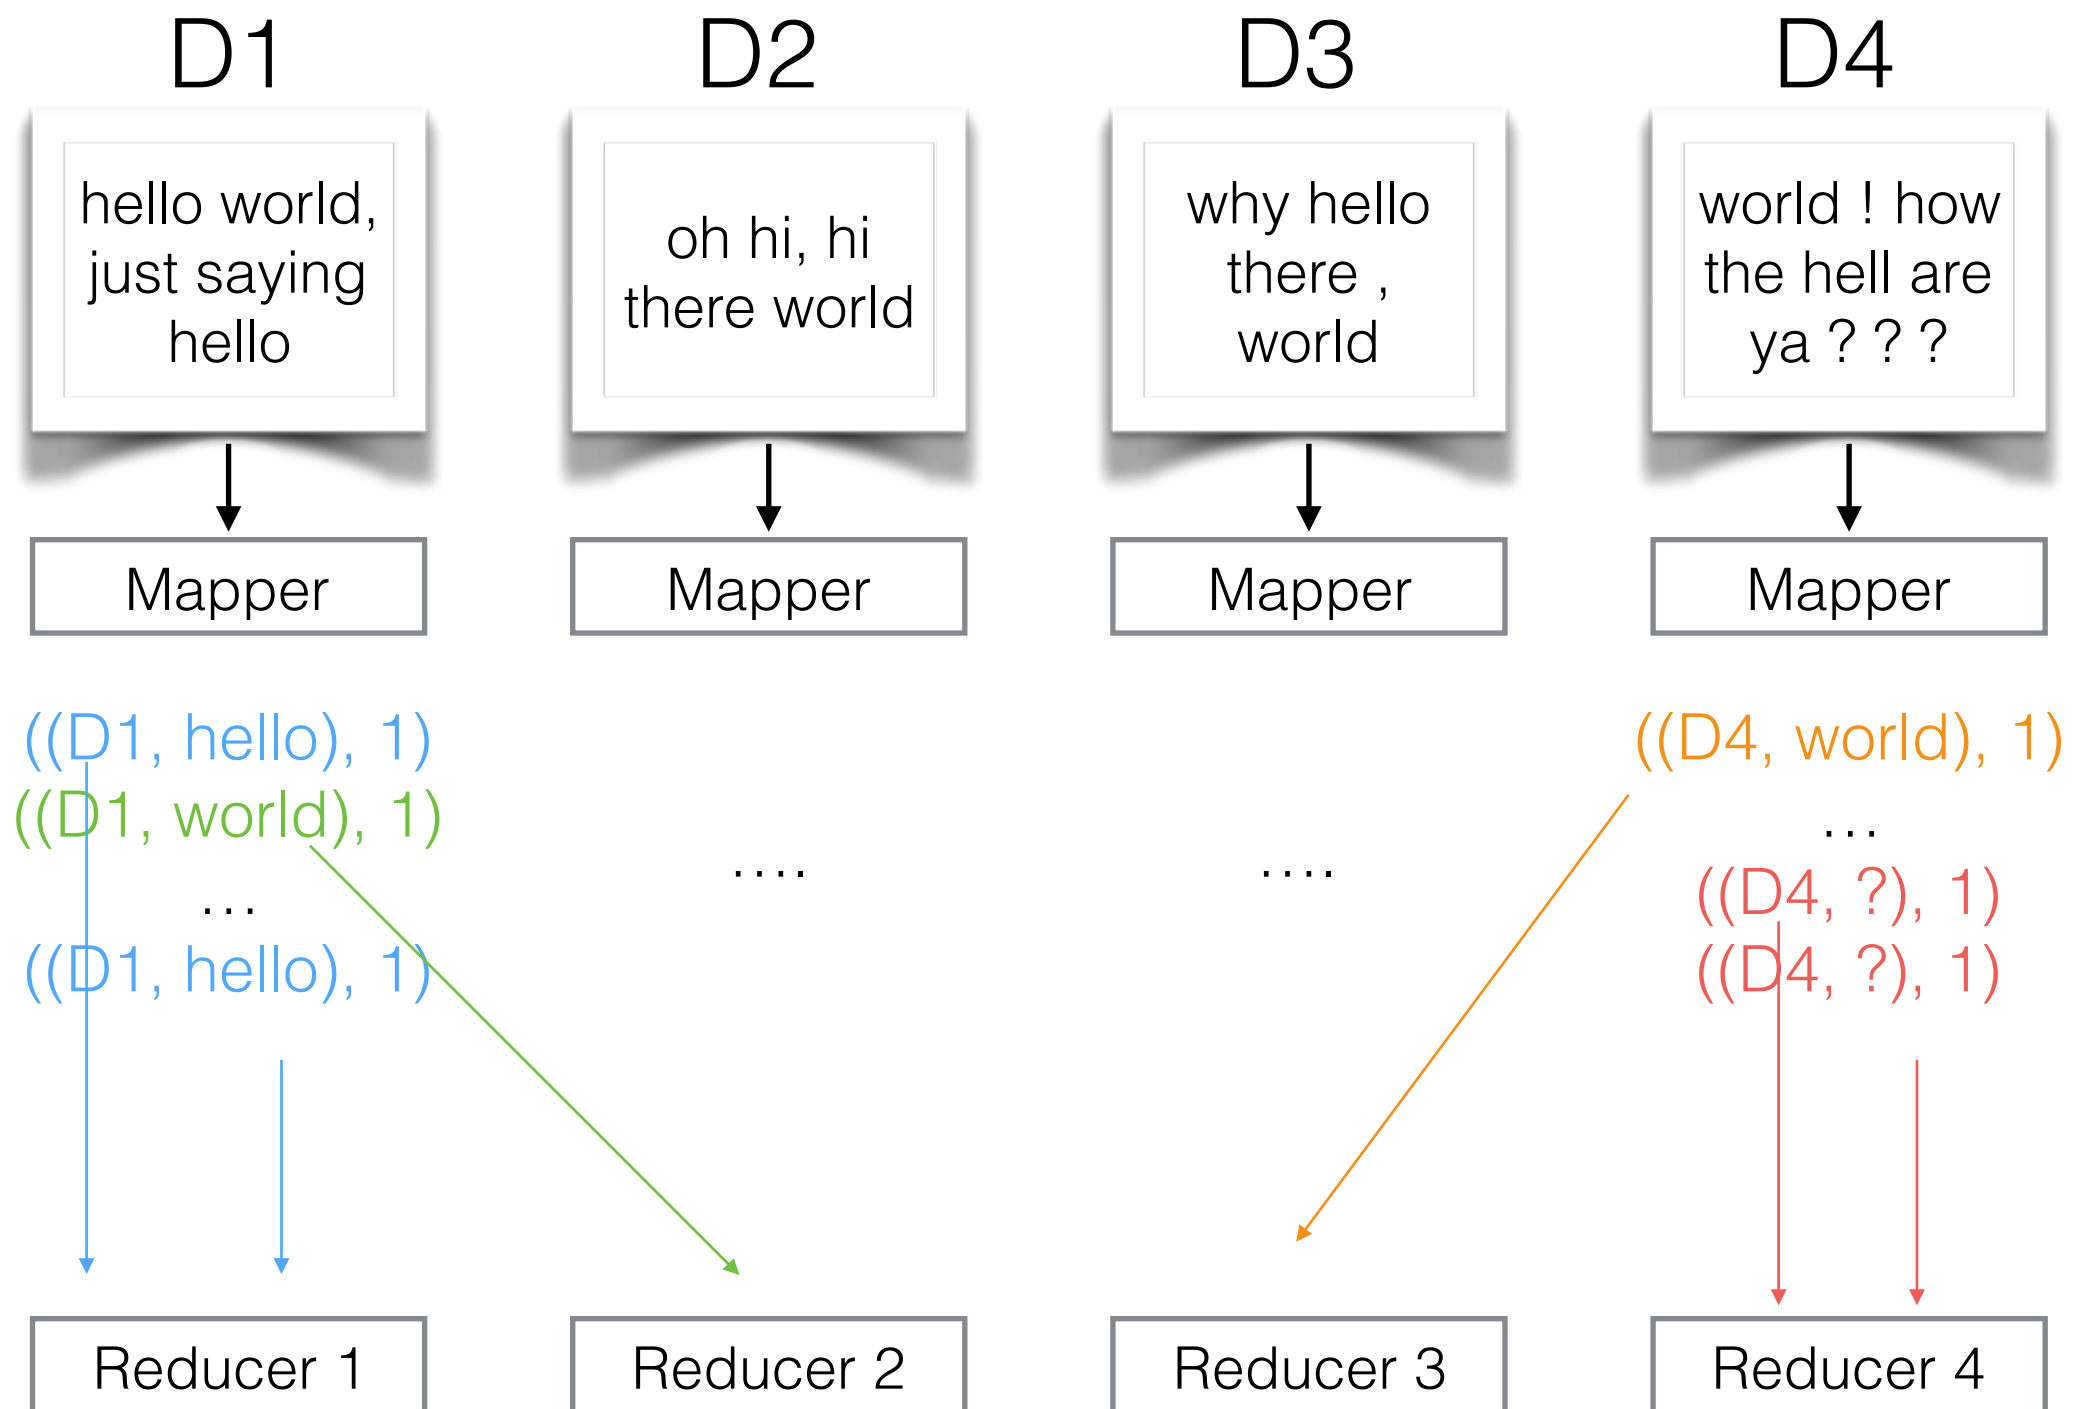
D3

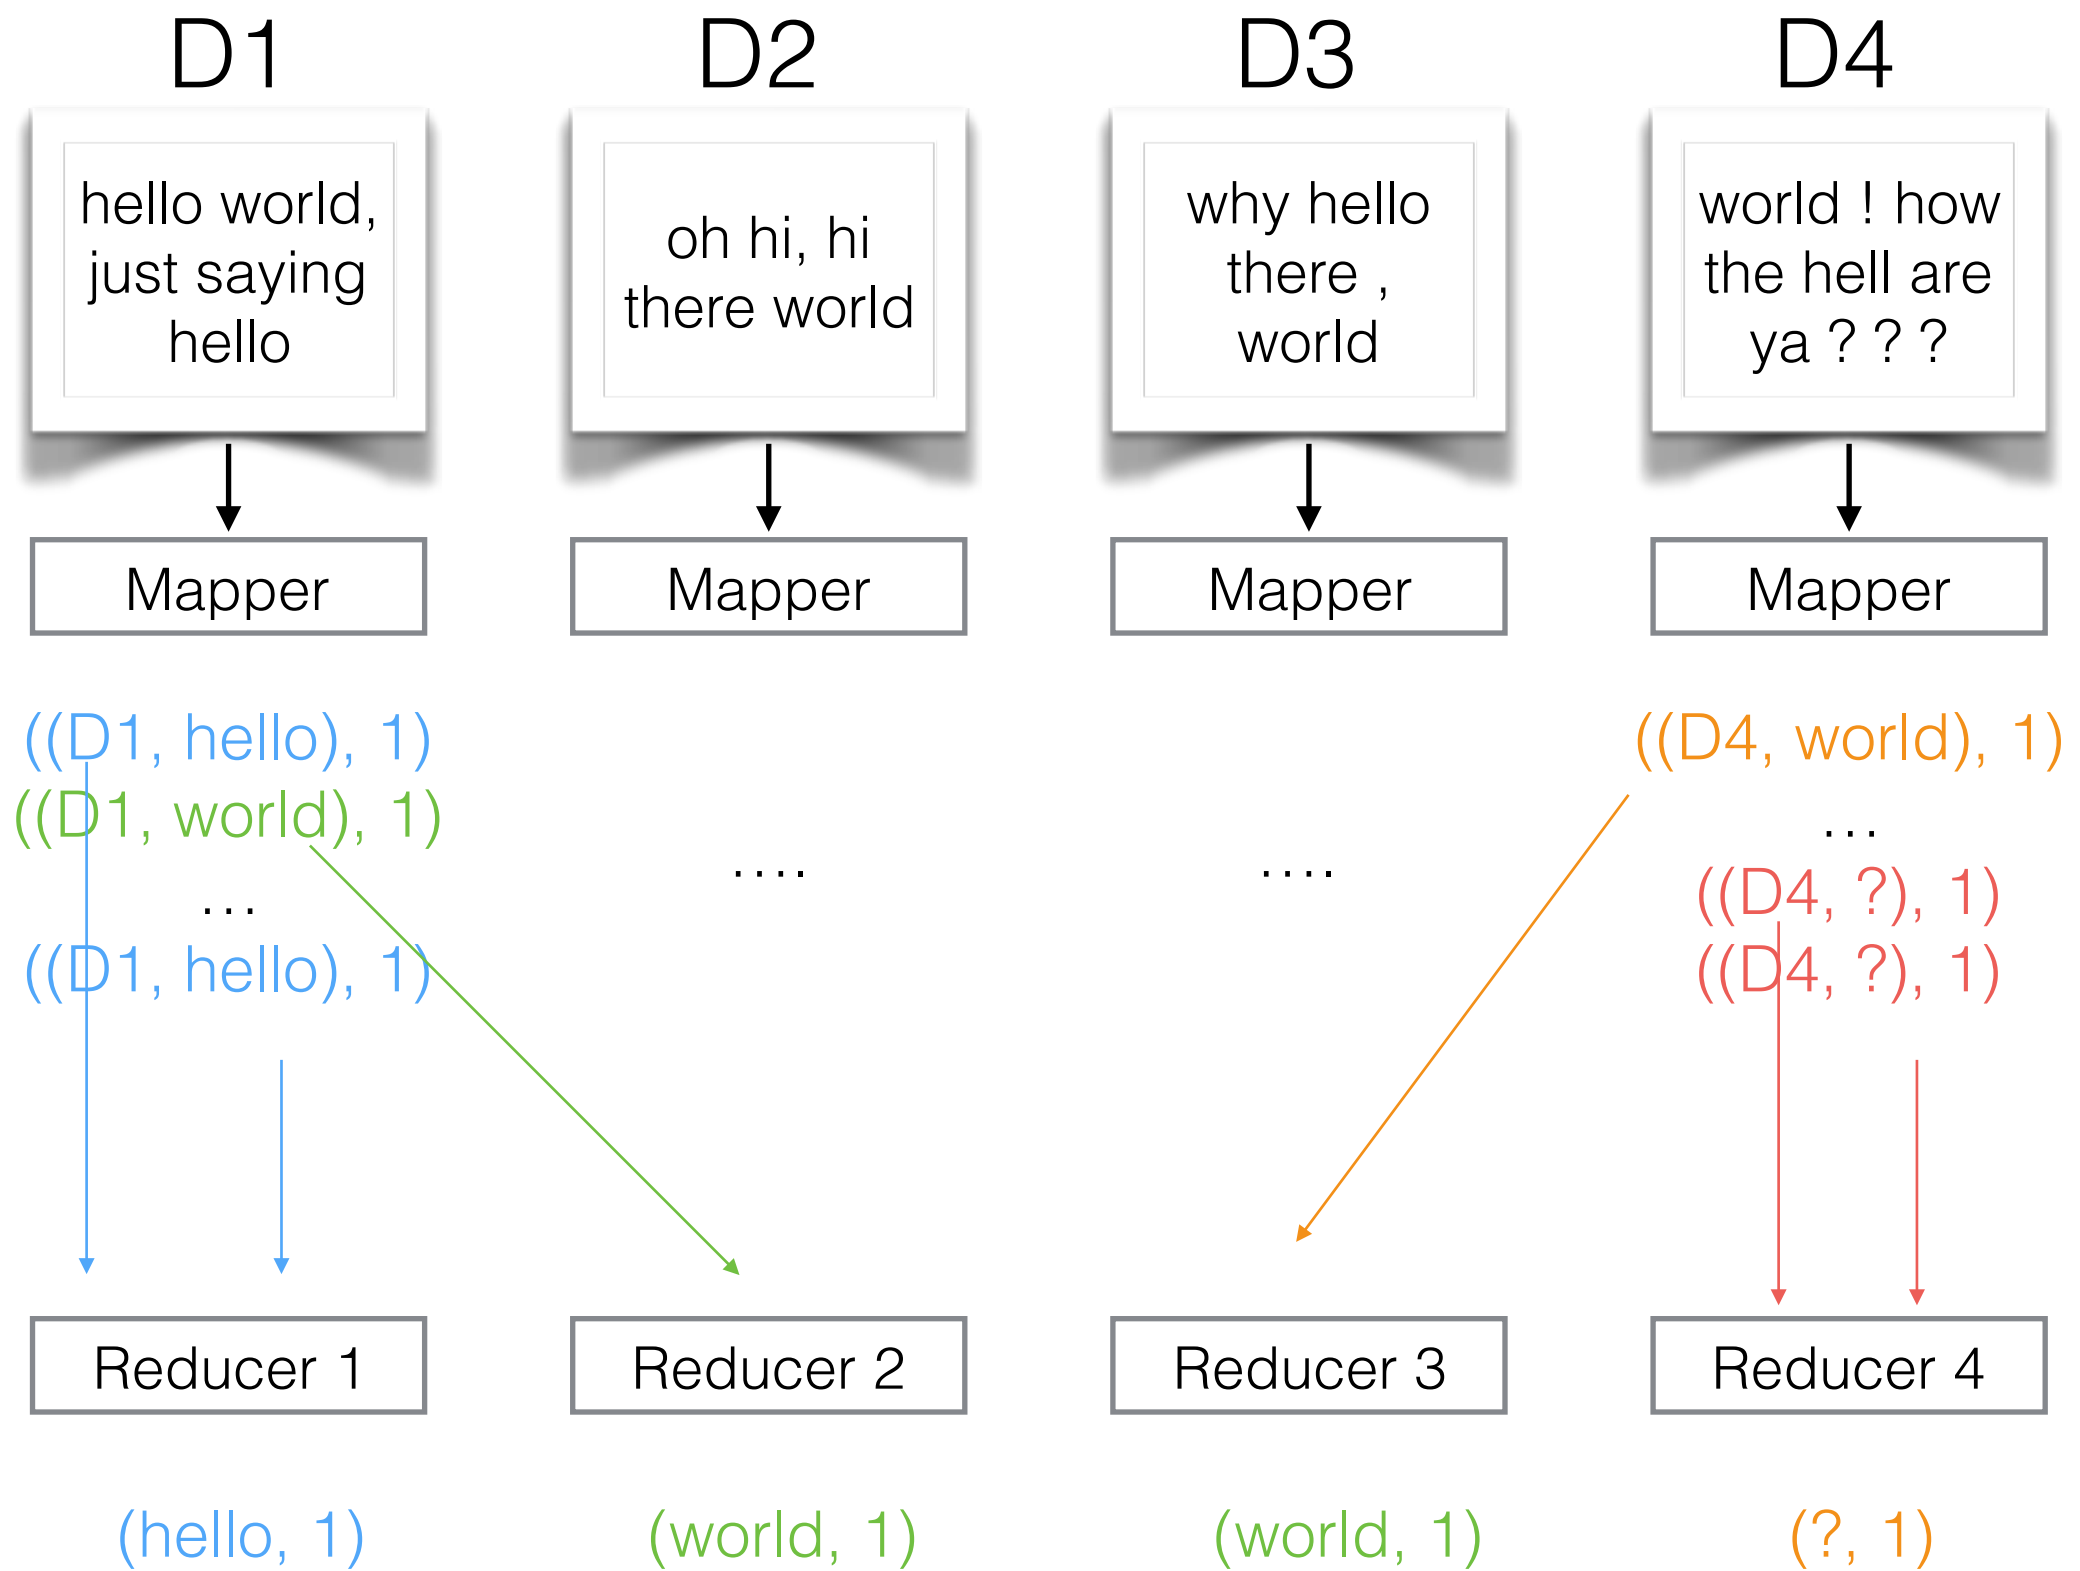
why hello
there ,
world

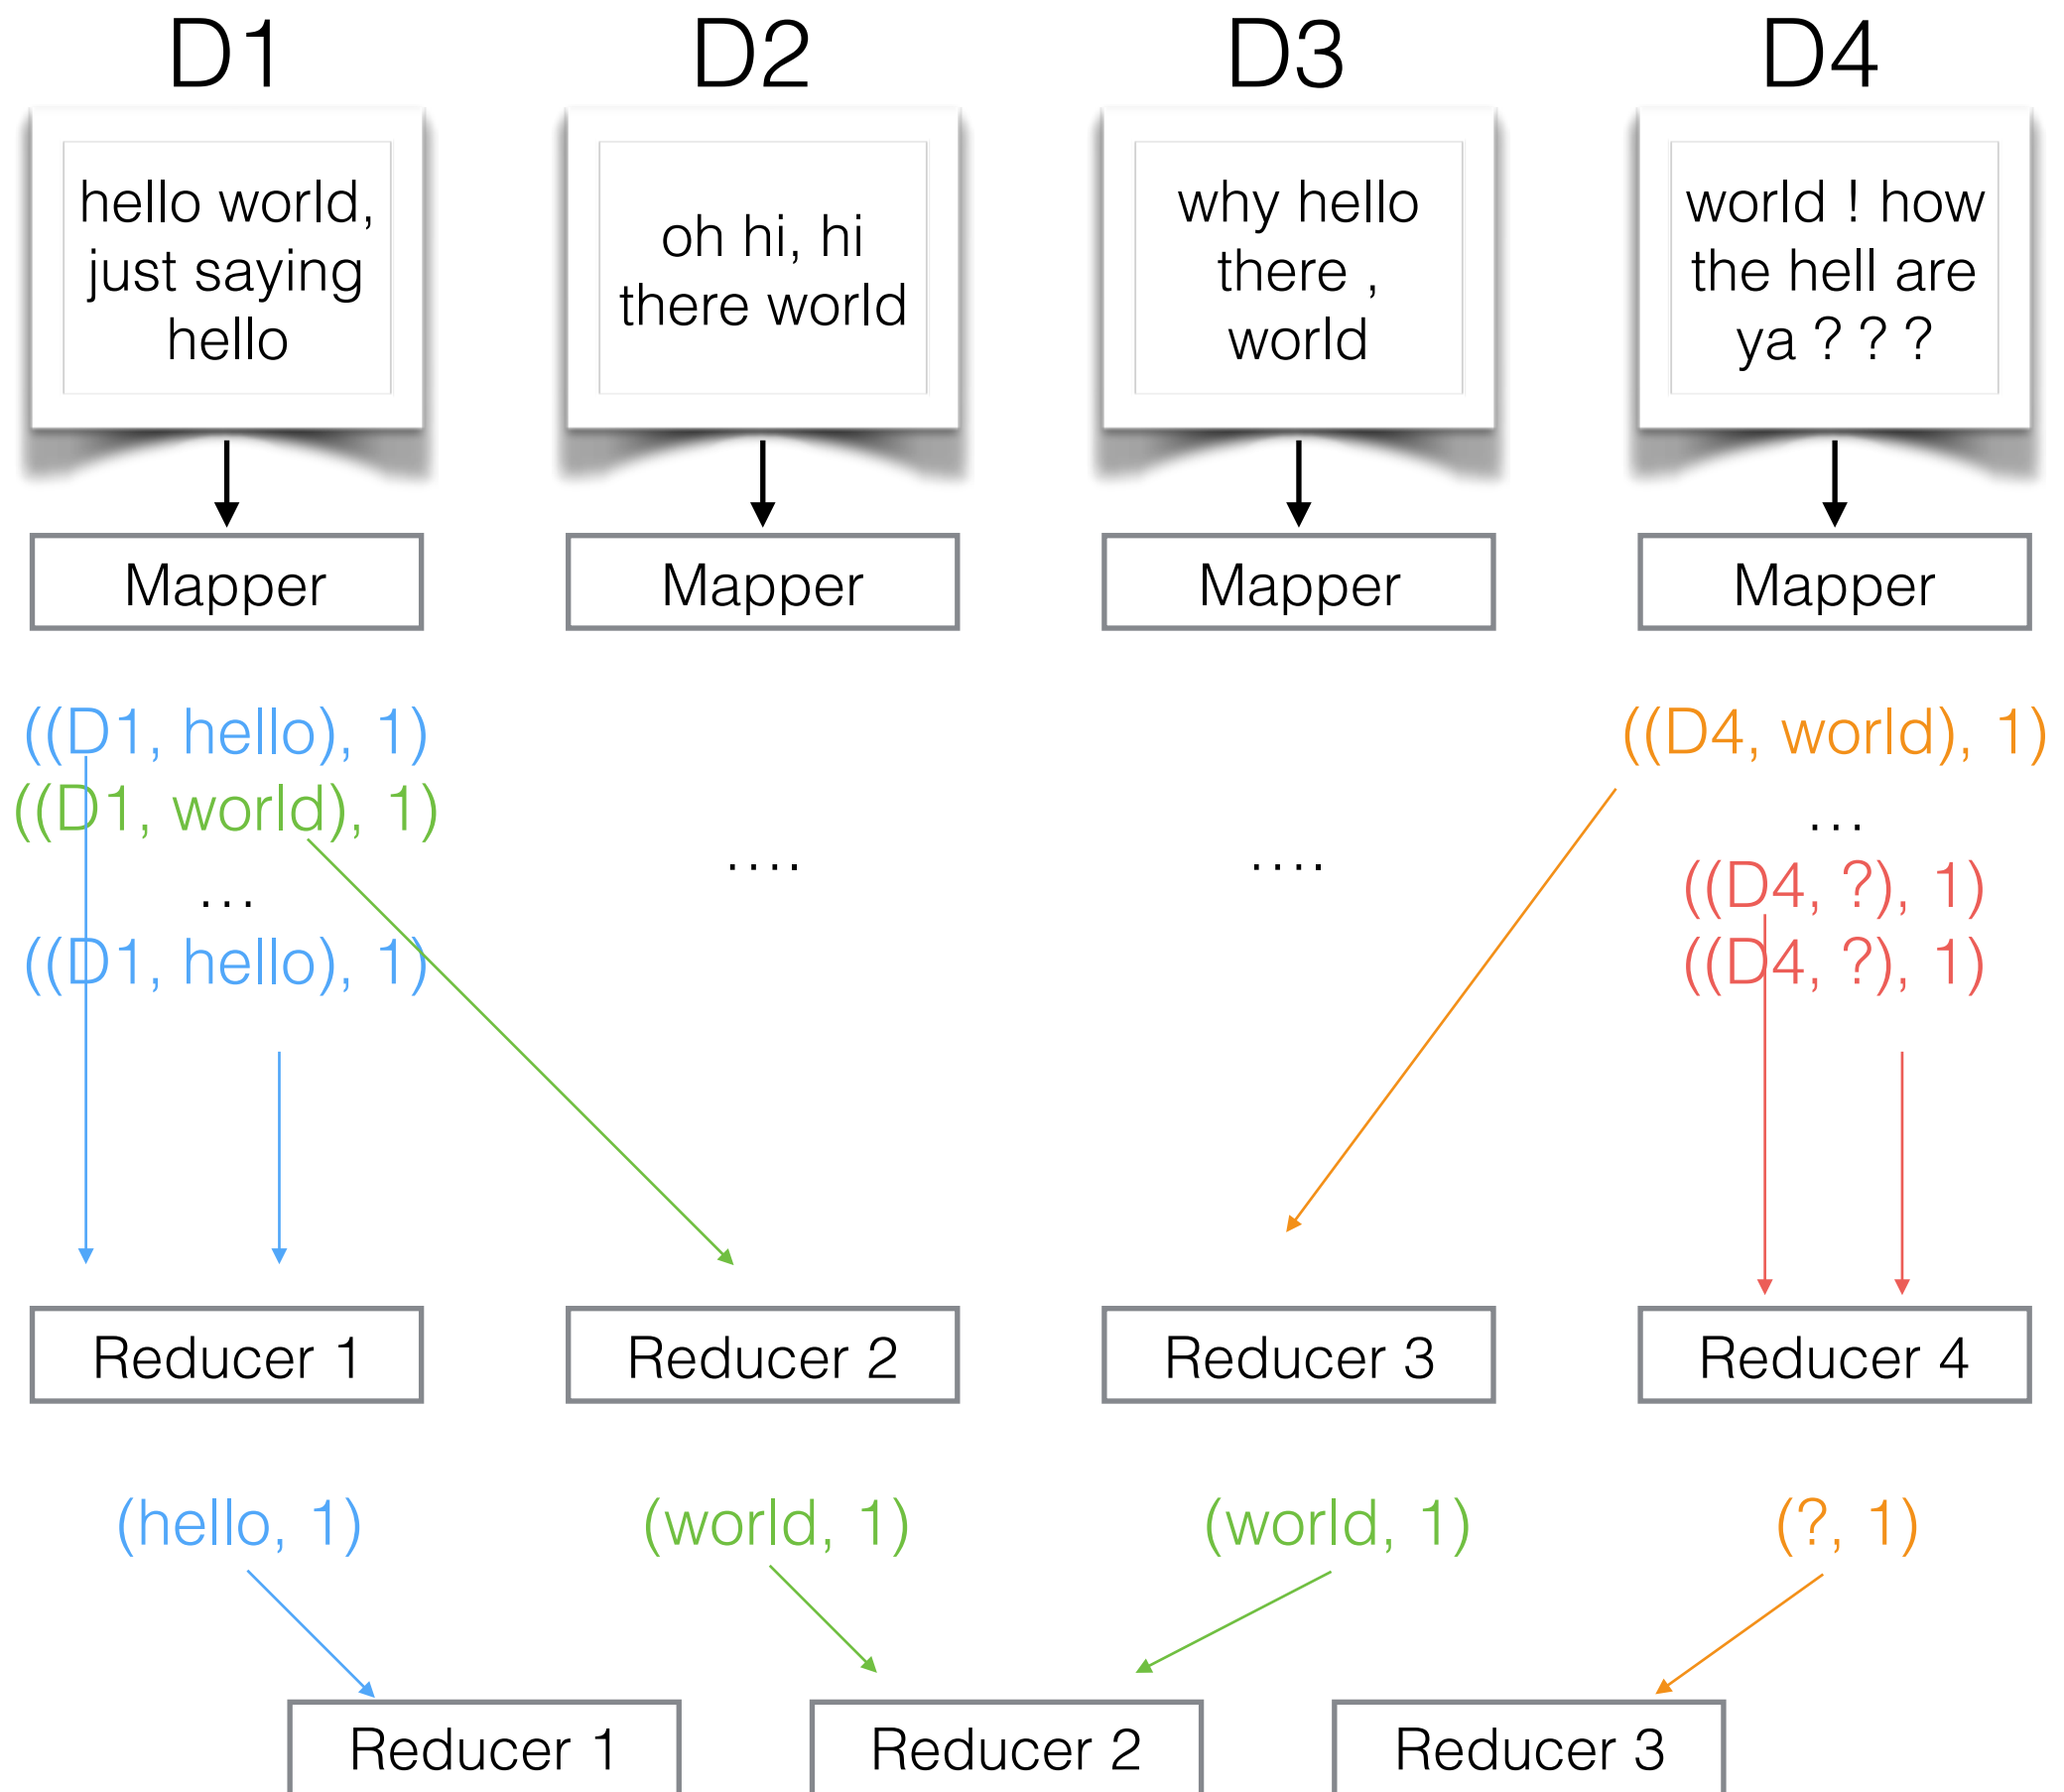
D4

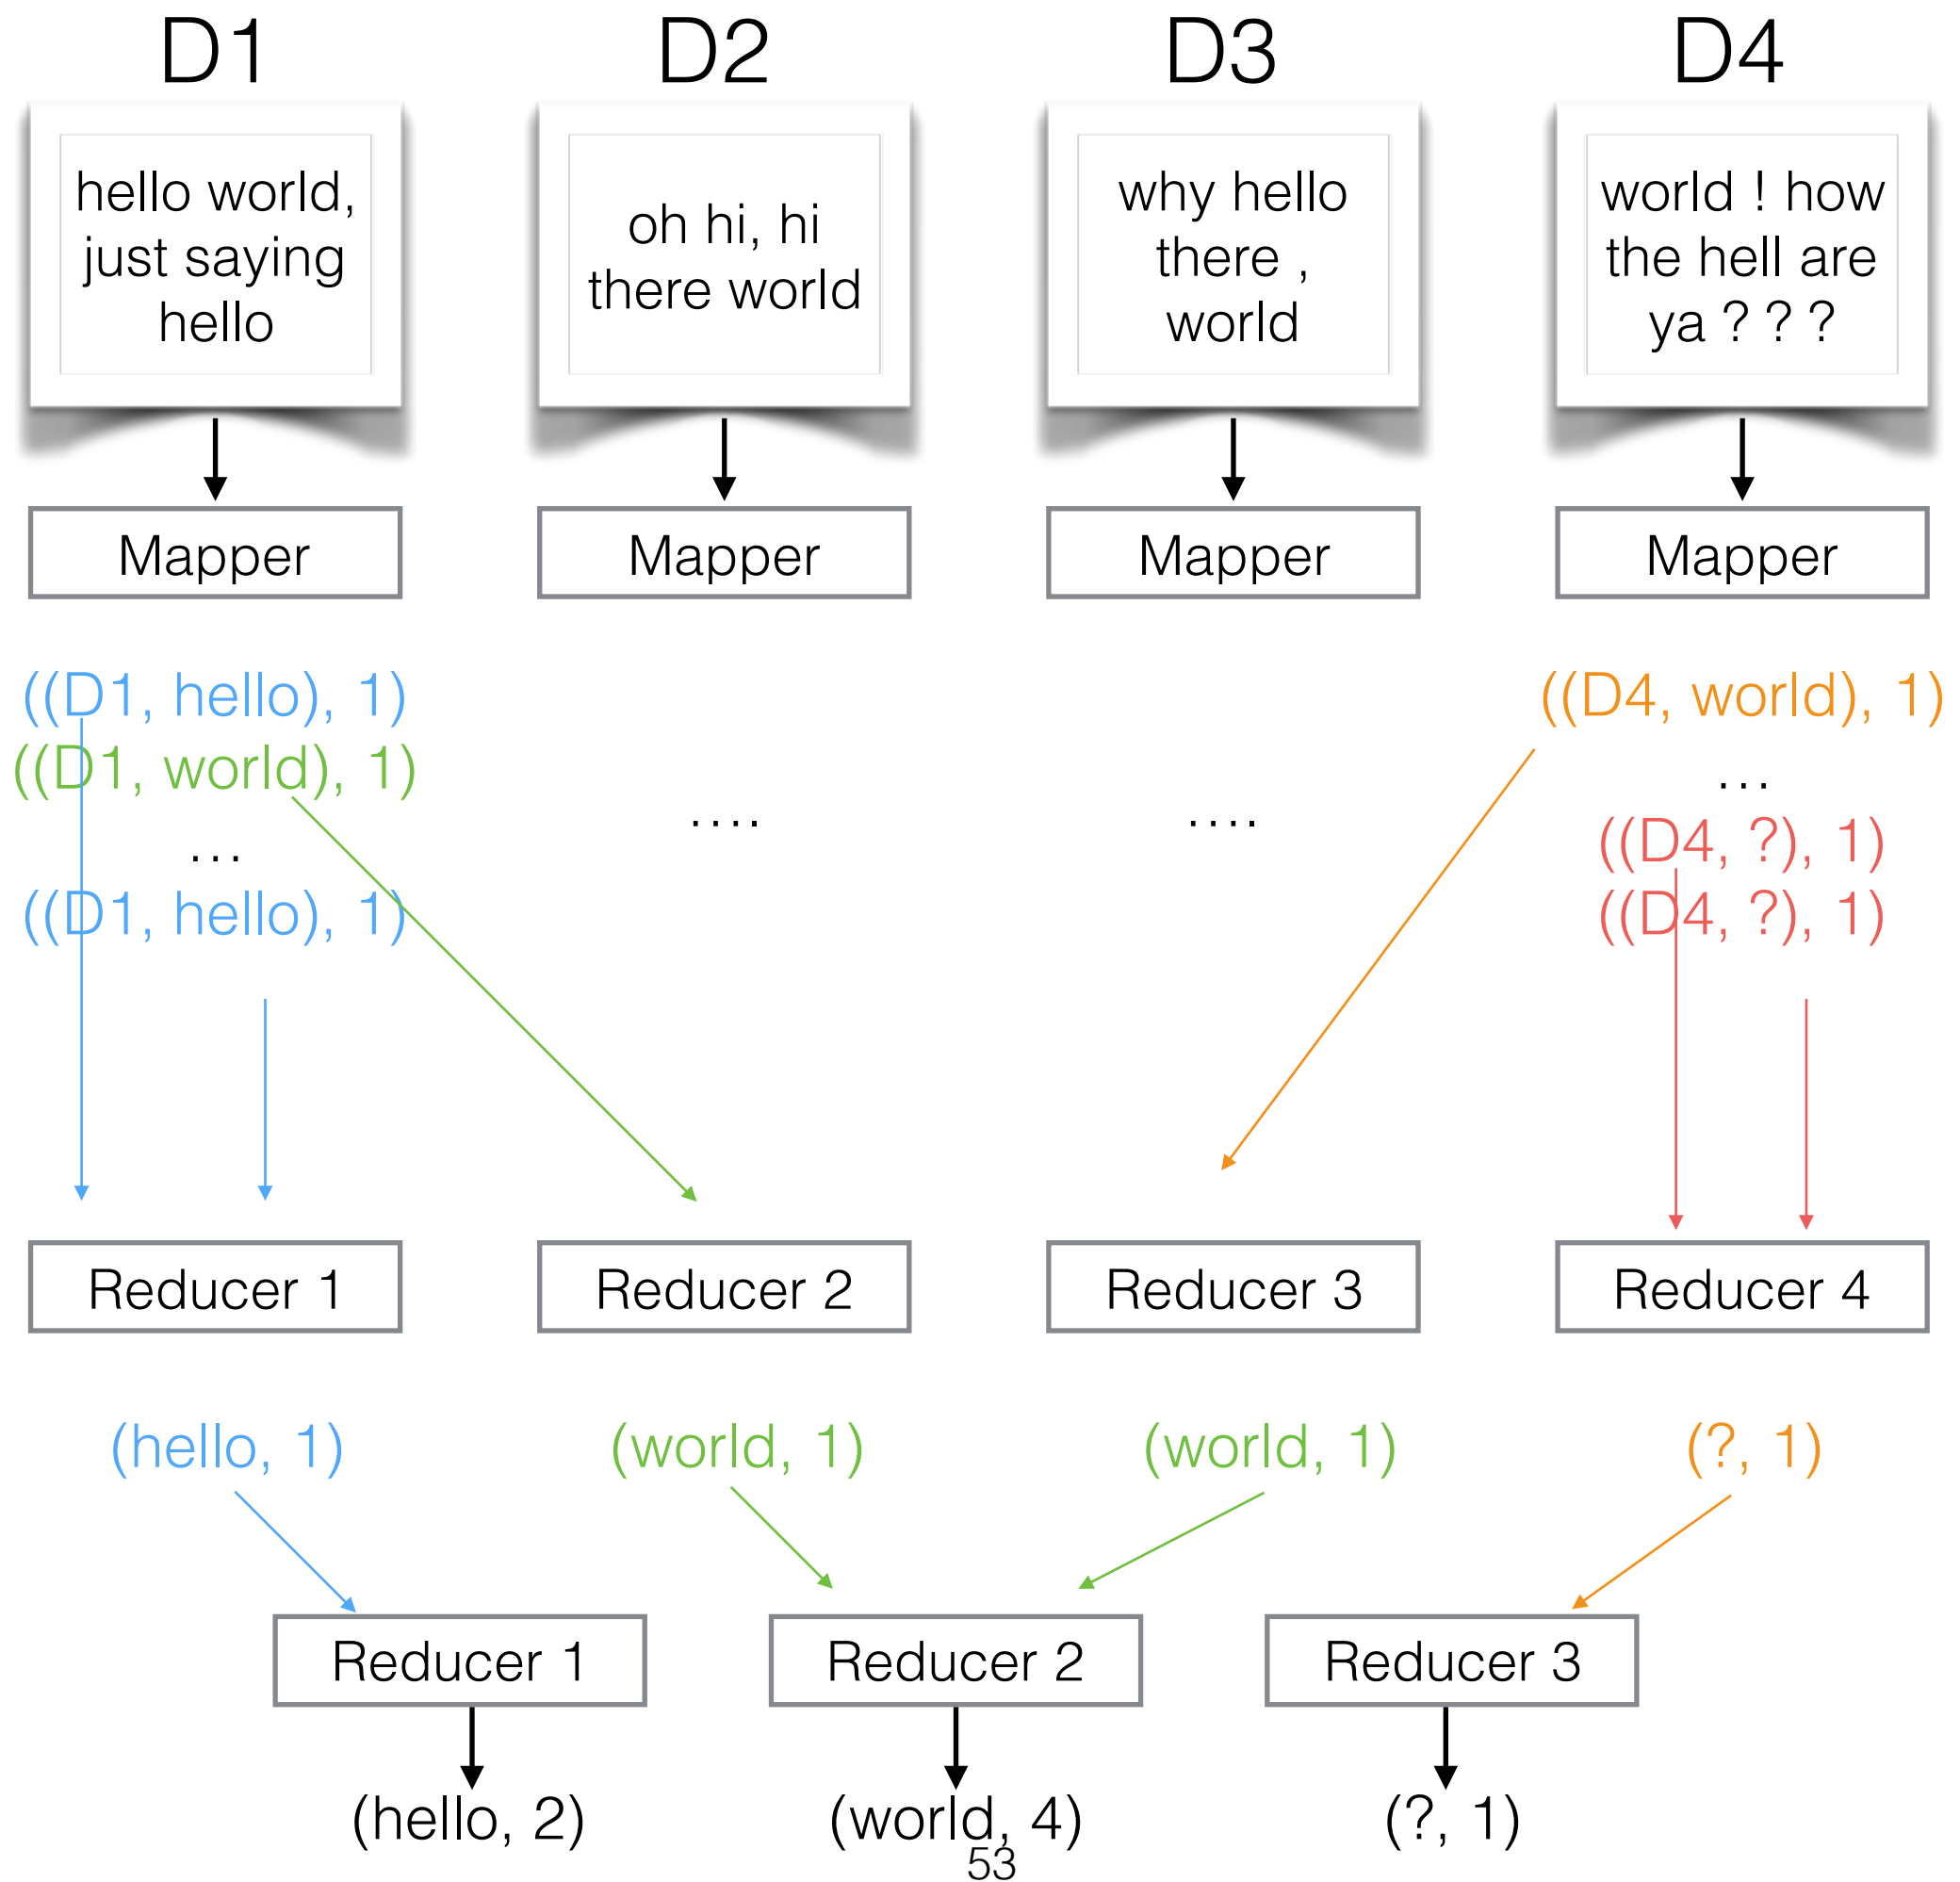
world ! how
the hell are
ya ? ? ?











(non)Clicker Question!

Find the number of unique documents that each word occurs in?

```
// enumerate occurrences of each word
// with count of 1
def MapFn1: String -> (String, Int) {
    ???
}
def ReduceFn1: String -> (String, Int) {
    ???
}
// sum the total counts of each word
def ReduceFn2: (String, Int) -> (String, Int) {
    ???
}
// define your pipeline
def main() {
    Table<String, String> table = read(table_path)
    Table<String, Int> output =
        table.MapFn1().ReduceFn1().ReduceFn2();
    write(output)
}
```

```

// enumerate occurrences of each word
// with count of 1
def MapFn1: (String, String) -> ((String, String), Int) {
    for w in input.value().split() {
        emit((input.key(), w), 1)
    }
}
def ReduceFn1: ((String, String), Int) -> (String, Int) {
    emit(input.key()[1], 1)
}
// sum the total counts of each word
def ReduceFn2: (String, Int) -> (String, Int) {
    sum = 0;
    for (w, c) in input { sum += c; }
    emit(w, sum);
}
// define your pipeline
def main() {
    Table<String, String> table = read(table_path)
    Table<String, Int> output =
        table.MapFn1().MapFn2().ReduceFn();
    write(output)
}

```

```

// enumerate occurrences of each word
// with count of 1
def MapFn1: (String, String) -> ((String, String), Int) {
    for w in input.value().split() {
        emit((input.key(), w), 1)
    }
}
// ignore the value list! ("unique")
def ReduceFn1: ((String, String), Int) -> (String, Int) {
    emit(input.key()[1], 1)
}
// sum the total counts of each word
def ReduceFn2: (String, Int) -> (String, Int) {
    sum = 0;
    for (w, c) in input { sum += c; }
    emit(w, sum);
}
// define your pipeline
def main() {
    Table<String, String> table = read(table_path)
    Table<String, Int> output =
        table.MapFn1().MapFn2().ReduceFn();
    write(output)
}

```


Find the number of unique documents that each word occurs in?

```
// enumerate occurrences
// of each word with count of 1
def MapFn1: {
  for w in input.value().split() {
    emit((input.key(), w), 1)
  }
}
def ReduceFn1: {
  emit(input.key()[1], 1)
}
// sum the total counts
// of each word
def ReduceFn2:{
  sum = 0;
  for (w, c) in input{ sum += c; }
  emit(w, sum);
}
```

Find the number of unique documents that each word occurs in?

```
// enumerate occurrences
// of each word with count of 1
def MapFn1: {
  for w in input.value().split() {
    emit((input.key(), w), 1)
  }
}
def ReduceFn1: {
  emit(input.key()[1], 1)
}
// sum the total counts
// of each word
def ReduceFn2:{
  sum = 0;
  for (w, c) in input{ sum += c; }
  emit(w, sum);
}
```

```
// enumerate occurrences
// of each word with count of 1
def MapFn1: {
  for w in input.value().split() {
    emit(input.key(), w)
  }
}
def ReduceFn1: {
  for w in input.value(){emit(w, 1)}
}
// sum the total counts
// of each word
def ReduceFn2:(S, I) -> (S, I){
  sum = 0;
  for (w, c) in input{ sum += c; }
  emit(w, sum);
}
```

(non)Clicker Question!

Find the number of unique documents that each word occurs in?

```
// enumerate occurrences
// of each word with count of 1
def MapFn1: {
  for w in input.value().split() {
    emit((input.key(), w), 1)
  }
}
def ReduceFn1: {
  emit(input.key()[1], 1)
}
// sum the total counts
// of each word
def ReduceFn2:{
  sum = 0;
  for (w, c) in input{ sum += c; }
  emit(w, sum);
}
```

```
// enumerate occurrences
// of each word with count of 1
def MapFn1: {
  for w in input.value().split() {
    emit(input.key(), w)
  }
}
def ReduceFn1: {
  for w in input.value(){emit(w, 1)}
}
// sum the total counts
// of each word
def ReduceFn2:(S, I) -> (S, I){
  sum = 0;
  for (w, c) in input{ sum += c; }
  emit(w, sum);
}
```

Do these produce the same output?

(a) Yes ₅₉

(b) No

Clicker Question!

```
Input K: V
Doc1 : here are some words
Doc2: words words words
Doc3: here are words
```

```
def MapFn1: (S, S) -> (S, S) {
  for w in input.value().split() {
    emit(input.key(), w)
  }
}
```

```
def ReduceFn1: (S, S) -> (S, I) {
  for w in input.value() {
    emit(w, 1)
  }
}

def ReduceFn2: (S, I) -> (S, I) {
  sum = 0;
  for (w, c) in input {
    sum += c;
  }
  emit(w, sum);
}
```

What will this produce?

- (a) here:2, are:2, some:1, words:3
- (b) here:2, are:2, some:1, words:5
- (c) here:1, are:1, some:1, words:1

Clicker Question!

```
Input K: V
Doc1 : here are some words
Doc2: words words words
Doc3: here are words
```

```
def MapFn1: (S, S) -> (S, S) {
  for w in input.value().split() {
    emit(input.key(), w)
  }
}
```

```
def ReduceFn1: (S, S) -> (S, I) {
  for w in input.value() {
    emit(w, 1)
  }
}
```

```
def ReduceFn2: (S, I) -> (S, I) {
  sum = 0;
  for (w, c) in input {
    sum += c;
  }
  emit(w, sum);
}
```

What will this produce?

- (a) here:2, are:2, some:1, words:3**
- (b) here:2, are:1, some:1, words:5**
- (c) here:1, are:1, some:1, words:1**

Clicker Question!

```
Input K: V
Doc1 : here are some words
Doc2: words words words
Doc3: here are words
```

```
def MapFn1: (S, S) -> (S, S) {
  for w in input.value().split() {
    emit(input.key(), w)
  }
}
```

Reducer is by
DocId only

```
def ReduceFn1: (S, S) -> (S, I) {
  for w in input.value() {
    emit(w, 1)
  }
}
```

```
def ReduceFn2: (S, I) -> (S, I) {
  sum = 0;
  for (w, c) in input {
    sum += c;
  }
  emit(w, sum);
}
```

What will this produce?

- (a) here:2, are:2, some:1, words:3
- (b) here:2, are:1, some:1, words:5**
- (c) here:1, are:1, some:1, words:1

Other MapReduce Functions

- Sort
- Unique
- Sample
- First
- Filter
- Join

Other MapReduce Functions

- Sort
- Unique
- Sample
- First
- Filter
- Join

Joins

Joins

- Joins are usually computed “under the hood” by most MR implementations (like in SQL)

Joins

- Joins are usually computed “under the hood” by most MR implementations (like in SQL)
- But you can imagine having to do them yourself...



Hacky Joins



- Joins are usually computed “under the hood” by most MR implementations (like in SQL)
- But you can imagine having to do them yourself...
- ...or, if you aren't that imaginative type, you can just look at the homework



Hacky Joins



- Joins are usually computed “under the hood” by most MR implementations (like in SQL)
- But you can imagine having to do them yourself...
- ...or, if you aren't that imaginative type, you can just look at the homework
- (sry)

Real Life Application

Real Life Application

Is Charles Mingus a **composer**?

Real Life Application

Is Charles Mingus a **composer**?

“Mingus is a **composer**”



Real Life Application

Is Charles Mingus a **composer**?

“Mingus is a **composer**”



Visions of Jazz: The First Century - Page 452 - Google Books Result

<https://books.google.com/books?isbn=0199879532>

Gary Giddins - 1998 - Music

If **Mingus is a composer** worthy of our attention, it must be because his melodies are one with his voicings and scaffolding. Set adrift among Harry Partch's globes ...

Jazz: There's a Mingus a-Monk us, in The Abstract Truth - Daily Kos

www.dailykos.com/story/.../Jazz-There-s-a-Mingus-a-Monk-us-in-The-Abstract-Trut... ▼

Mar 9, 2014 - **Mingus is a composer** and arranger. In fact a big band has been established which performs in Manhattan every week in NYC that just plays ...

Real Life Application

Is Charles Mingus a **1950s American jazz composer**?

“Mingus is a **1950s American jazz composer**”



No results found for "mingus is a 1950s american jazz composer".

Real Life Application

Is Charles Mingus a **1950s American jazz composer**?

Real Life Application

Is Charles Mingus a **1950s American jazz composer**?

... if **Mingus is a composer** worthy of our attention, it must be because...

Mingus dominated the scene back in the 1950s and 1960s.

Mingus was truly a product of America in all its historic complexities...

A virtuoso bassist and composer, **Mingus** irrevocably **changed the face of jazz**...

Real Life Application

ComposerX dominated the scene back in the 1950s and 1960s.



ComposerX is a **1950s composer**.

Real Life Application

Subject	Predicate	Object
Barack Obama	won	the electoral vote
Kamala Lopez	wrote	an op-ed for HuffPo
Charles Mingus	wrote	jazz
Barack Obama	opposed	the appropriations bill
Barack Obama	listens to	jazz

Category	Entity
Person	Barack Obama
Person	Kamala Lopez
Person	Charles Mingus
Huffington Post Columnists	Barack Obama
Huffington Post Columnists	Kamala Lopez
US Presidents	Barack Obama
Jazz Composers	Charles Mingus

Joins

Subject	Predicate	Object
Barack Obama	won	the electoral vote
Kamala Lopez	wrote	an op-ed for HuffPo
Charles Mingus	wrote	jazz
Barack Obama	opposed	the appropriations bill
Barack Obama	listens to	jazz

Category	Entity
Person	Barack Obama
Person	Kamala Lopez
Person	Charles Mingus
Huffington Post Columnists	Barack Obama
Huffington Post Columnists	Kamala Lopez
US Presidents	Barack Obama
Jazz Composers	Charles Mingus

Desired output:

Subject	Predicate	Object	Categories
Barack Obama	won	the electoral vote	Person, US_Presidents, Huffington_Post_Columnists
Kamala Lopez	wrote	an op-ed for HuffPo	Person, Huffington_Post_Columnists,

Joins

Subject	Predicate	Object
Barack Obama	won	the electoral vote
Kamala Lopez	wrote	an op-ed for HuffPo
Charles Mingus	wrote	jazz
Barack Obama	opposed	the appropriations bill
Barack Obama	listens to	jazz

Category	Entity
Person	Barack Obama
Person	Kamala Lopez
Person	Charles Mingus
Huffington Post Columnists	Barack Obama
Huffington Post Columnists	Kamala Lopez
US Presidents	Barack Obama
Jazz Composers	Charles Mingus

Desired output:

Subject	Predicate	Object	Categories
Barack Obama	won	the electoral vote	Person, US_Presidents, Huffington_Post_Columnists
Kamala Lopez	wrote	an op-ed for HuffPo	Person, Huffington_Post_Columnists,

Joins

Facts

Subject	Predicate	Object
Barack Obama	won	the electoral vote
Kamala Lopez	wrote	an op-ed for HuffPo
Charles Mingus	wrote	jazz
Barack Obama	opposed	the appropriations bill
Barack Obama	listens to	jazz

Categories

Category	Entity
Person	Barack Obama
Person	Kamala Lopez
Person	Charles Mingus
Huffington Post Columnists	Barack Obama
Huffington Post Columnists	Kamala Lopez
US Presidents	Barack Obama
Jazz Composers	Charles Mingus

```
Select * from Facts, Categories
Where Subject == Entity
```

Joins

Facts

Subject	Predicate	Object
Barack Obama	won	the electoral vote
Kamala Lopez	wrote	an op-ed for HuffPo
Charles Mingus	wrote	jazz
Barack Obama	opposed	the appropriations bill
Barack Obama	listens to	jazz

Categories

Category	Entity
Person	Barack Obama
Person	Kamala Lopez
Person	Charles Mingus
Huffington Post Columnists	Barack Obama
Huffington Post Columnists	Kamala Lopez
US Presidents	Barack Obama
Jazz Composers	Charles Mingus

```
Select * from Facts, Categories
Where Subject == Entity
GroupBy Subject
```

Joins

Facts

Subject	Predicate	Object
Barack Obama	won	the electoral vote
Kamala Lopez	wrote	an op-ed for HuffPo
Charles Mingus	wrote	jazz
Barack Obama	opposed	the appropriations bill
Barack Obama	listens to	jazz

Categories

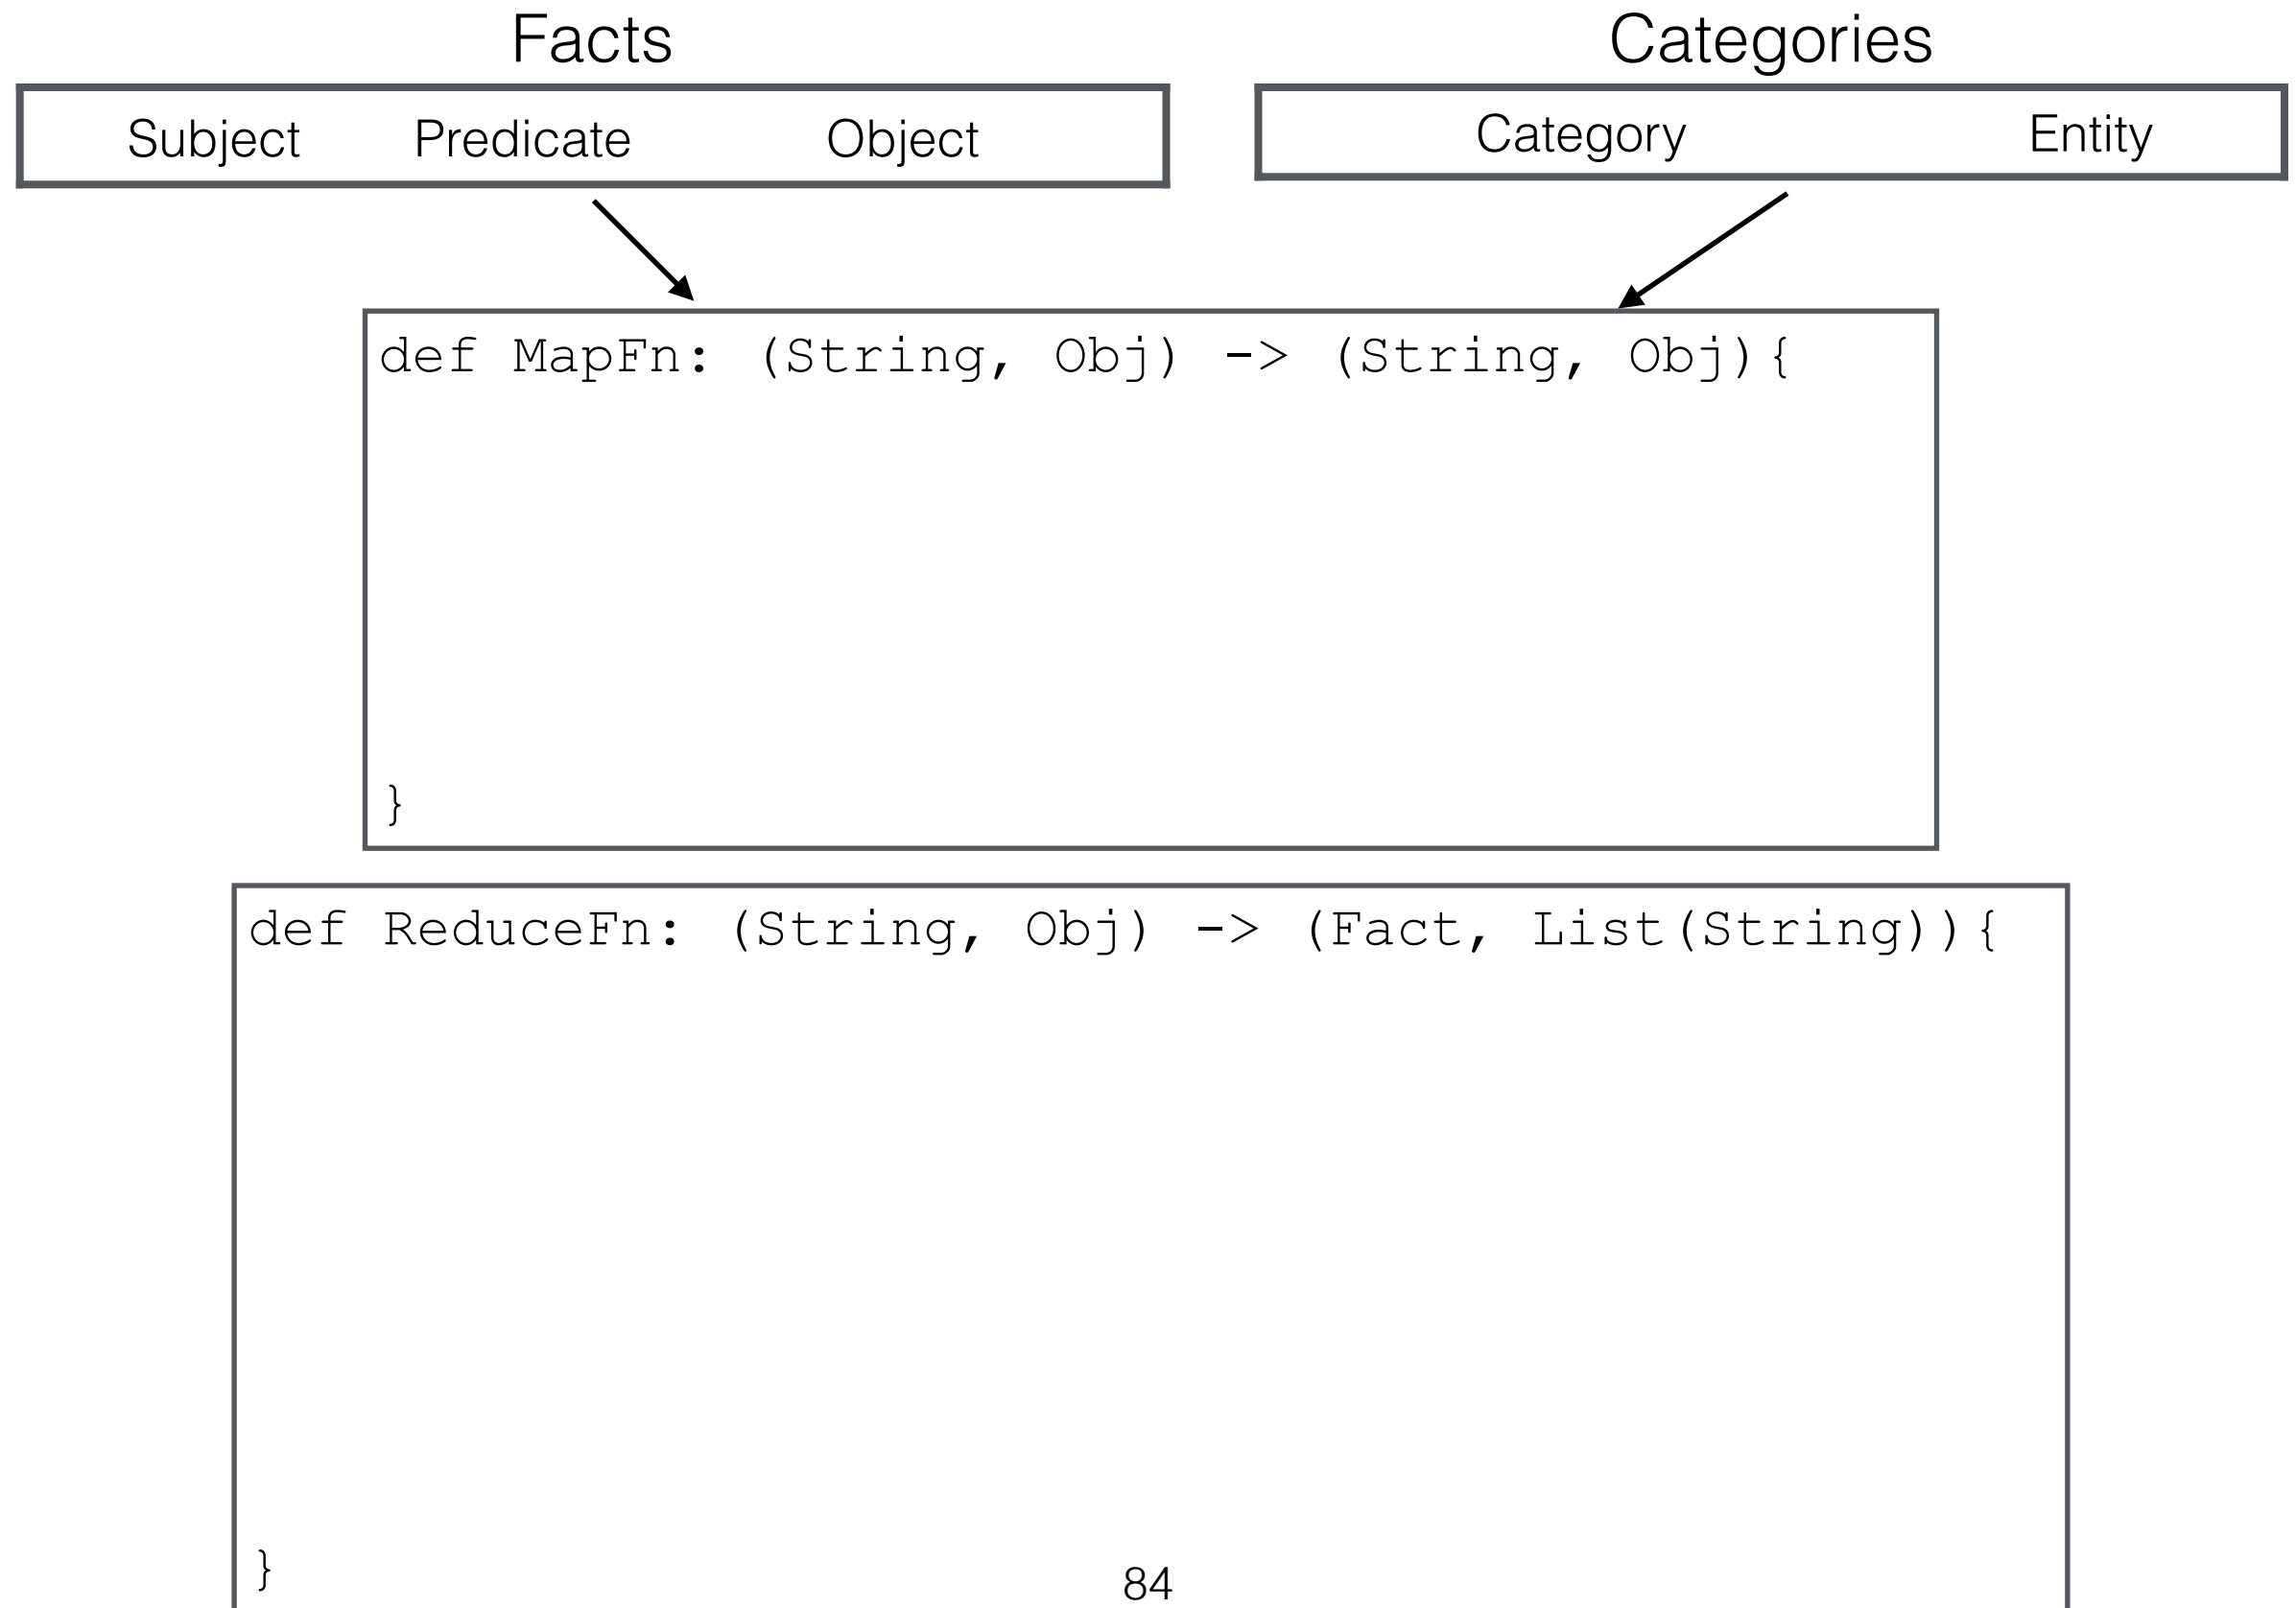
Category	Entity
Person	Barack Obama
Person	Kamala Lopez
Person	Charles Mingus
Huffington Post Columnists	Barack Obama
Huffington Post Columnists	Kamala Lopez
US Presidents	Barack Obama
Jazz Composers	Charles Mingus

```
Select * from Facts, Categories
Where Subject == Entity
GroupBy Subject
```

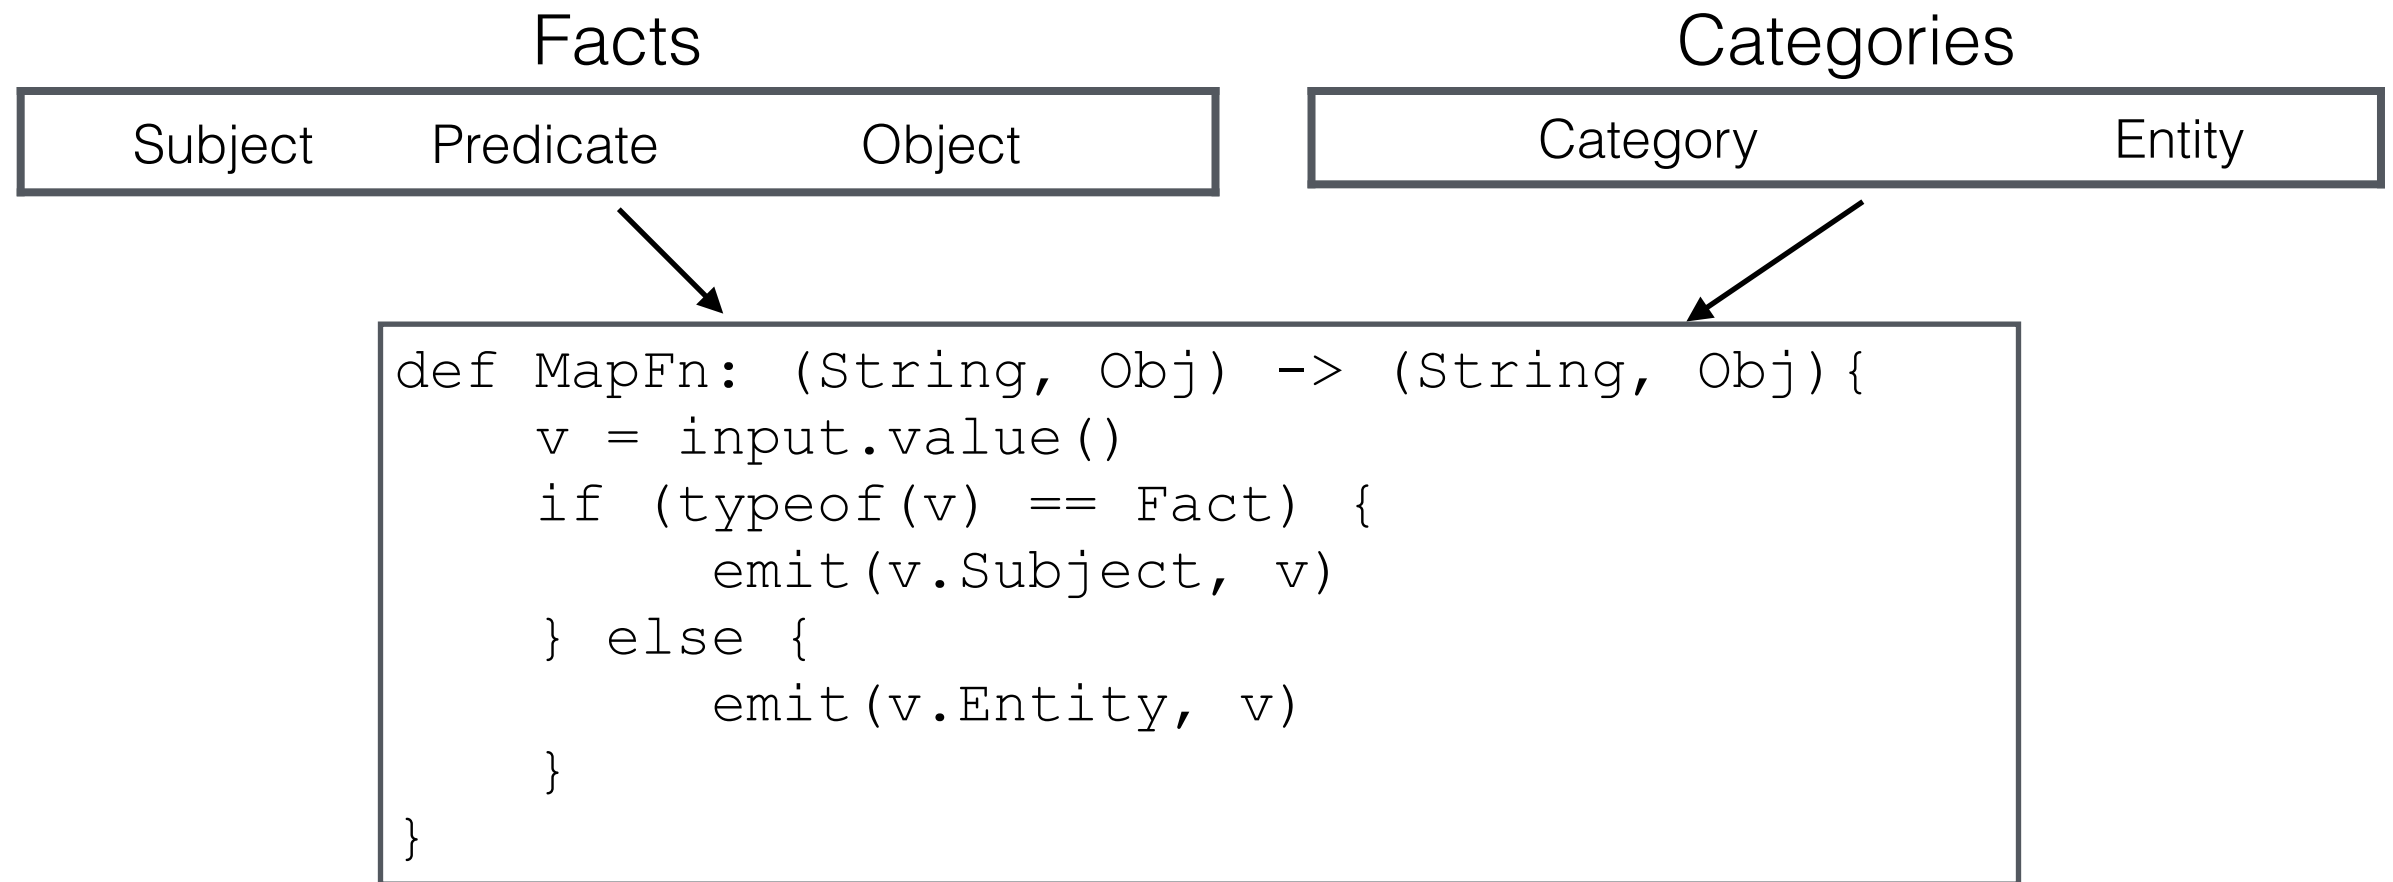


```
Key: String
Value: (list_of((String, String, String), list_of((String, String)))
```

DIY Joins



DIY Joins



```
def ReduceFn: (String, Obj) -> (Fact, List(String)) {
```

DIY Joins

Facts

Categories


Subject

Predicate

Object

Category

Entity



The diagram shows two input boxes at the top. The left box, labeled 'Facts', contains 'Subject', 'Predicate', and 'Object'. The right box, labeled 'Categories', contains 'Category' and 'Entity'. Arrows from both boxes point down to a central box containing a Scala MapFn function. The function takes a (String, Obj) pair and emits either (v.Subject, v) if the object is a Fact, or (v.Entity, v) otherwise.

```
def MapFn: (String, Obj) -> (String, Obj) {  
  v = input.value()  
  if (typeof(v) == Fact) {  
    emit(v.Subject, v)  
  } else {  
    emit(v.Entity, v)  
  }  
}
```

```
def ReduceFn: (String, Obj) -> (Fact, List(String)) {  
  all_cats = []; all_facts = []  
  for v in input.value() {  
  
  }  
}
```

DIY Joins

Facts

Categories


Subject

Predicate

Object

Category

Entity



```
def MapFn: (String, Obj) -> (String, Obj) {  
  v = input.value()  
  if (typeof(v) == Fact) {  
    emit(v.Subject, v)  
  } else {  
    emit(v.Entity, v)  
  }  
}
```

```
def ReduceFn: (String, Obj) -> (Fact, List(String)) {  
  all_cats = []; all_facts = []  
  for v in input.value() {  
    if (typeof(v) == Fact) {  
      all_facts.append(v)  
    } else {  
      all_cats.append(v.Category)  
    }  
  }  
}
```

DIY Joins

Facts

Categories


Subject

Predicate

Object

Category

Entity



```
def MapFn: (String, Obj) -> (String, Obj) {  
  v = input.value()  
  if (typeof(v) == Fact) {  
    emit(v.Subject, v)  
  } else {  
    emit(v.Entity, v)  
  }  
}
```

```
def ReduceFn: (String, Obj) -> (Fact, List(String)) {  
  all_cats = []; all_facts = []  
  for v in input.value() {  
    if (typeof(v) == Fact) {  
      all_facts.append(v)  
    } else {  
      all_cats.append(v.Category)  
    }  
  }  
  for f in all_facts { emit(f, all_cats); }  
}
```


DIY Joins

Facts

Categories


Subject

Predicate

Object

Category

Entity



```
def MapFn: (String, Obj) -> (String, Obj) {  
  v = input.value()  
  if (typeof(v) == Fact) {  
    emit(v.Subject, v)  
  } else {  
    emit(v.Entity, v)  
  }  
}
```

```
def ReduceFn: (String, Obj) -> (Fact, List(String)) {  
  all_cats = []; all_facts = []  
  for v in input.value() {  
    if (typeof(v) == Fact) {  
      all_facts.append(v)  
    } else {  
      all_cats.append(v.Category)  
    }  
  }  
  for f in all_facts { emit(f, all_cats); }  
}
```



Hacky Joins



Facts

Categories

Subject

Predicate

Object

Category


Entity



```
def MapFn: (String, Tuple) -> (String, Tuple) {  
  v = input.value()  
  if (len(v) == 3) {  
    emit(v[0], v)  
  } else {  
    emit(v[1], v)  
  }  
}
```



```
def ReduceFn: (String, Tuple) -> (Tuple, List(String)) {  
  all_cats = []; all_facts = []  
  for v in input.value() {  
    if (len(v) == 3) {  
      all_facts.append(v)  
    } else {  
      all_cats.append(v[0])  
    }  
  }  
  for f in all_facts { emit(f, all_cats); }  
}
```





Hacky Joins



Facts

Categories

Subject

Predicate

Object

Category

Entity

`<blink>`

This is not a thing you ever actually do! Please do not do this!

`</blink>`

```
def MapFn: (String, Tuple) => (String, Tuple) {  
  v = input.value()  
  if (len(v) == 3) {  
    emit(v[0], v)  
  } else {  
    emit(v[1], v)  
  }  
}  
  
def ReduceFn: (String, Tuple) => (Tuple, List[String]) {  
  all_cats = []; all_facts = []  
  for v in input.value() {  
    if (len(v) == 3) {  
      all_facts.append(v)  
    } else {  
      all_cats.append(v[0])  
    }  
  }  
  for f in all_facts { emit(f, all_cats); }  
}
```



Hacky Joins



Facts

Categories

Subject

Predicate

Object

Category

Entity

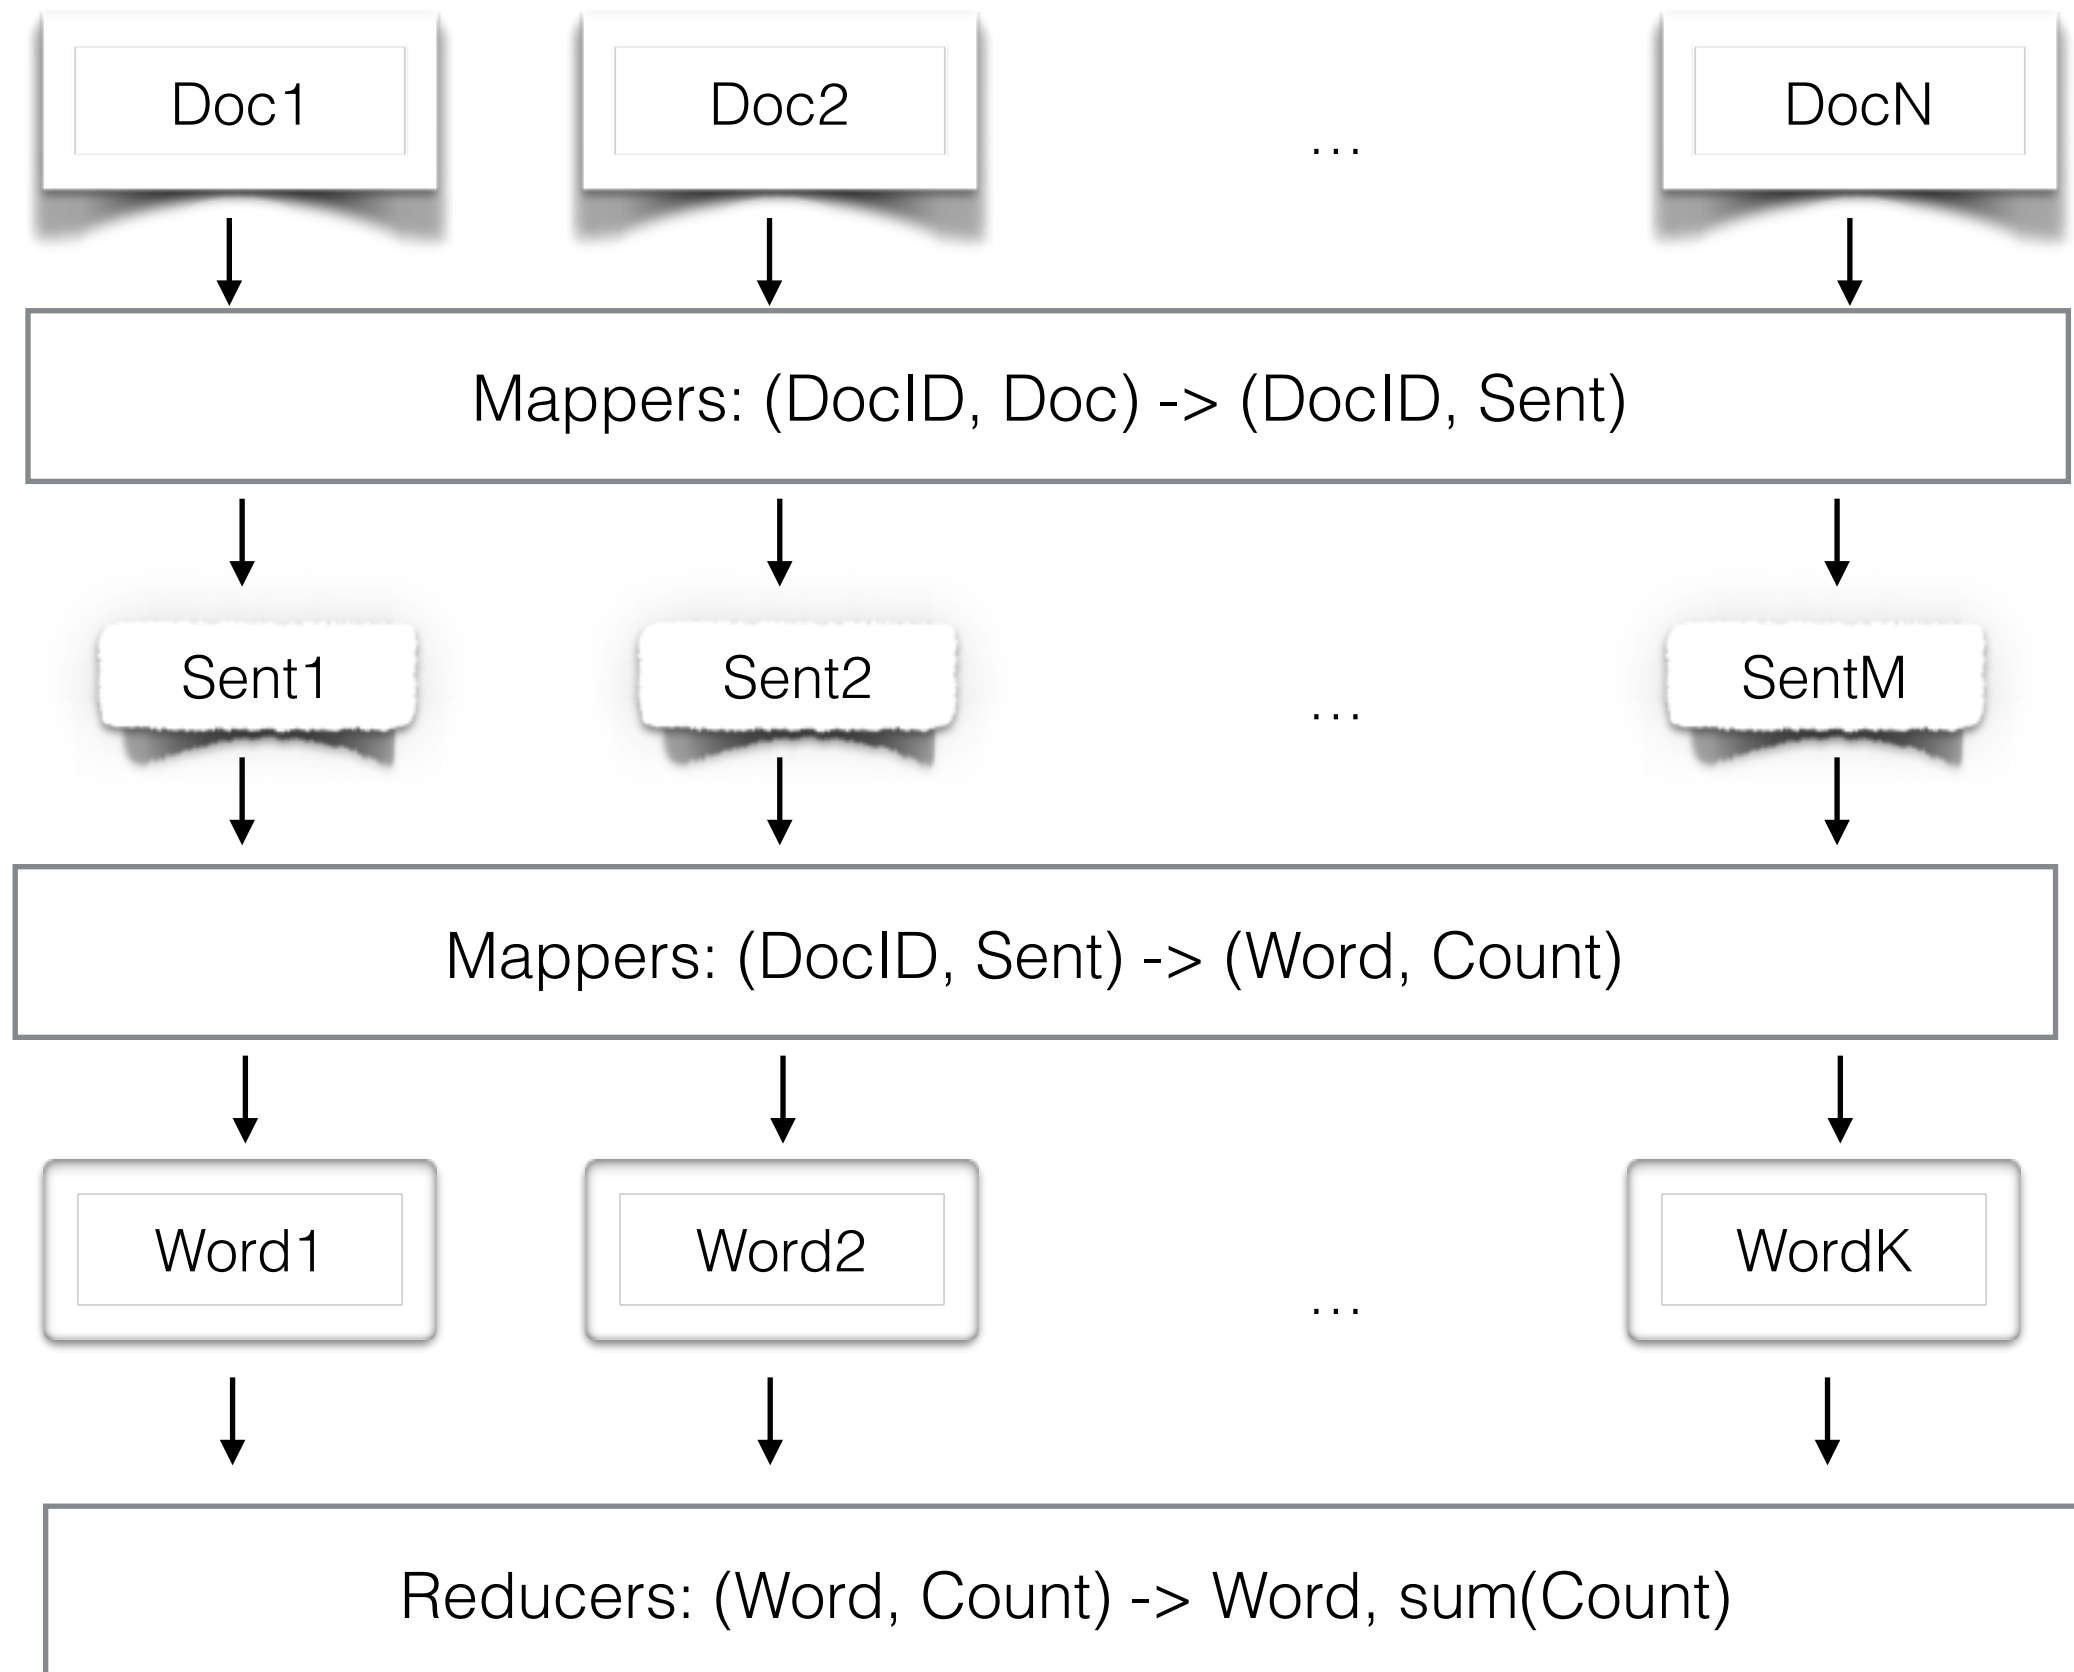
`<blink>`
This is not a thing you ever actually
do! Please do not do this!
`</blink>`

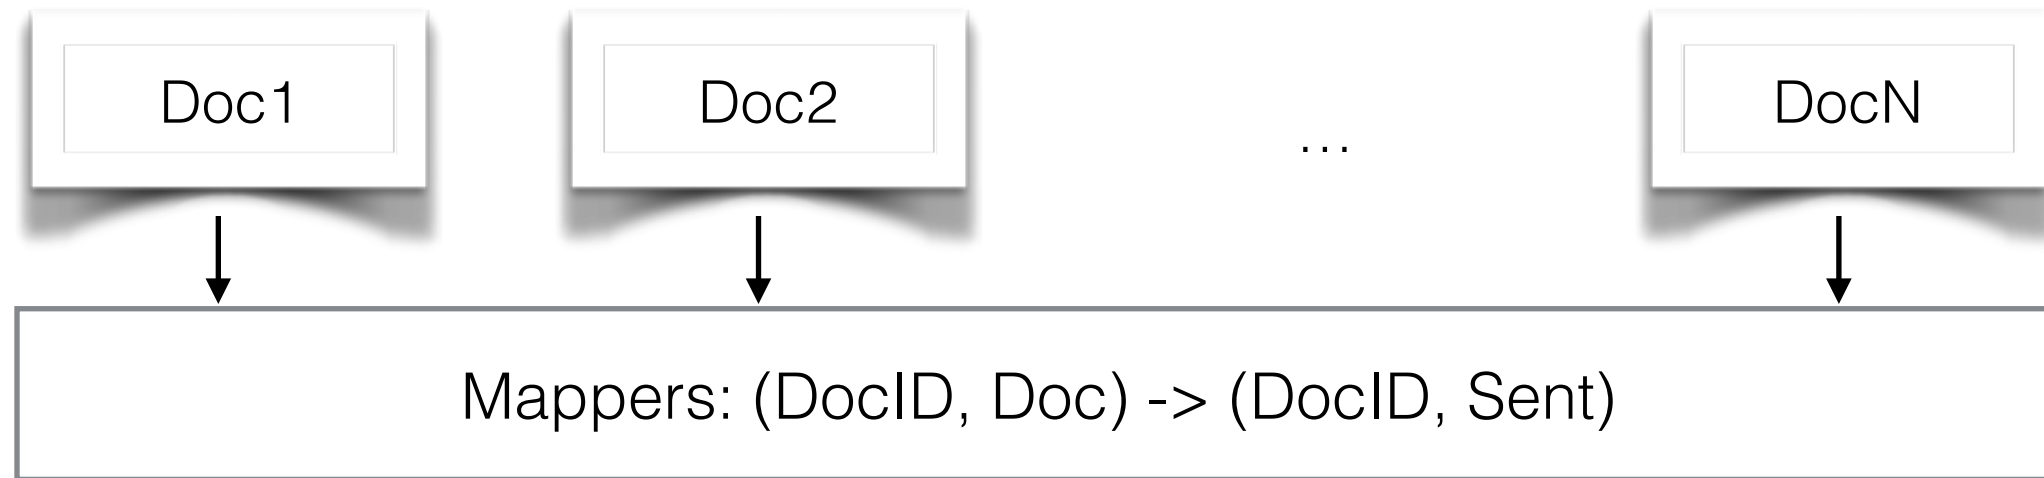
(But do it for the homework, and I will never tell,
and we will never speak of it, and if someone
asks you in an interview if you are the kind of
person who would do this in a map reduce, you
will deny deny deny. Agreed?)

```
}  
for f in all_facts { emit(f, all_cats); }  
}
```




Bottlenecks!






Clicker Question!

In the best-case scenario, how much parallelization could we get here (maximum number of mappers)?

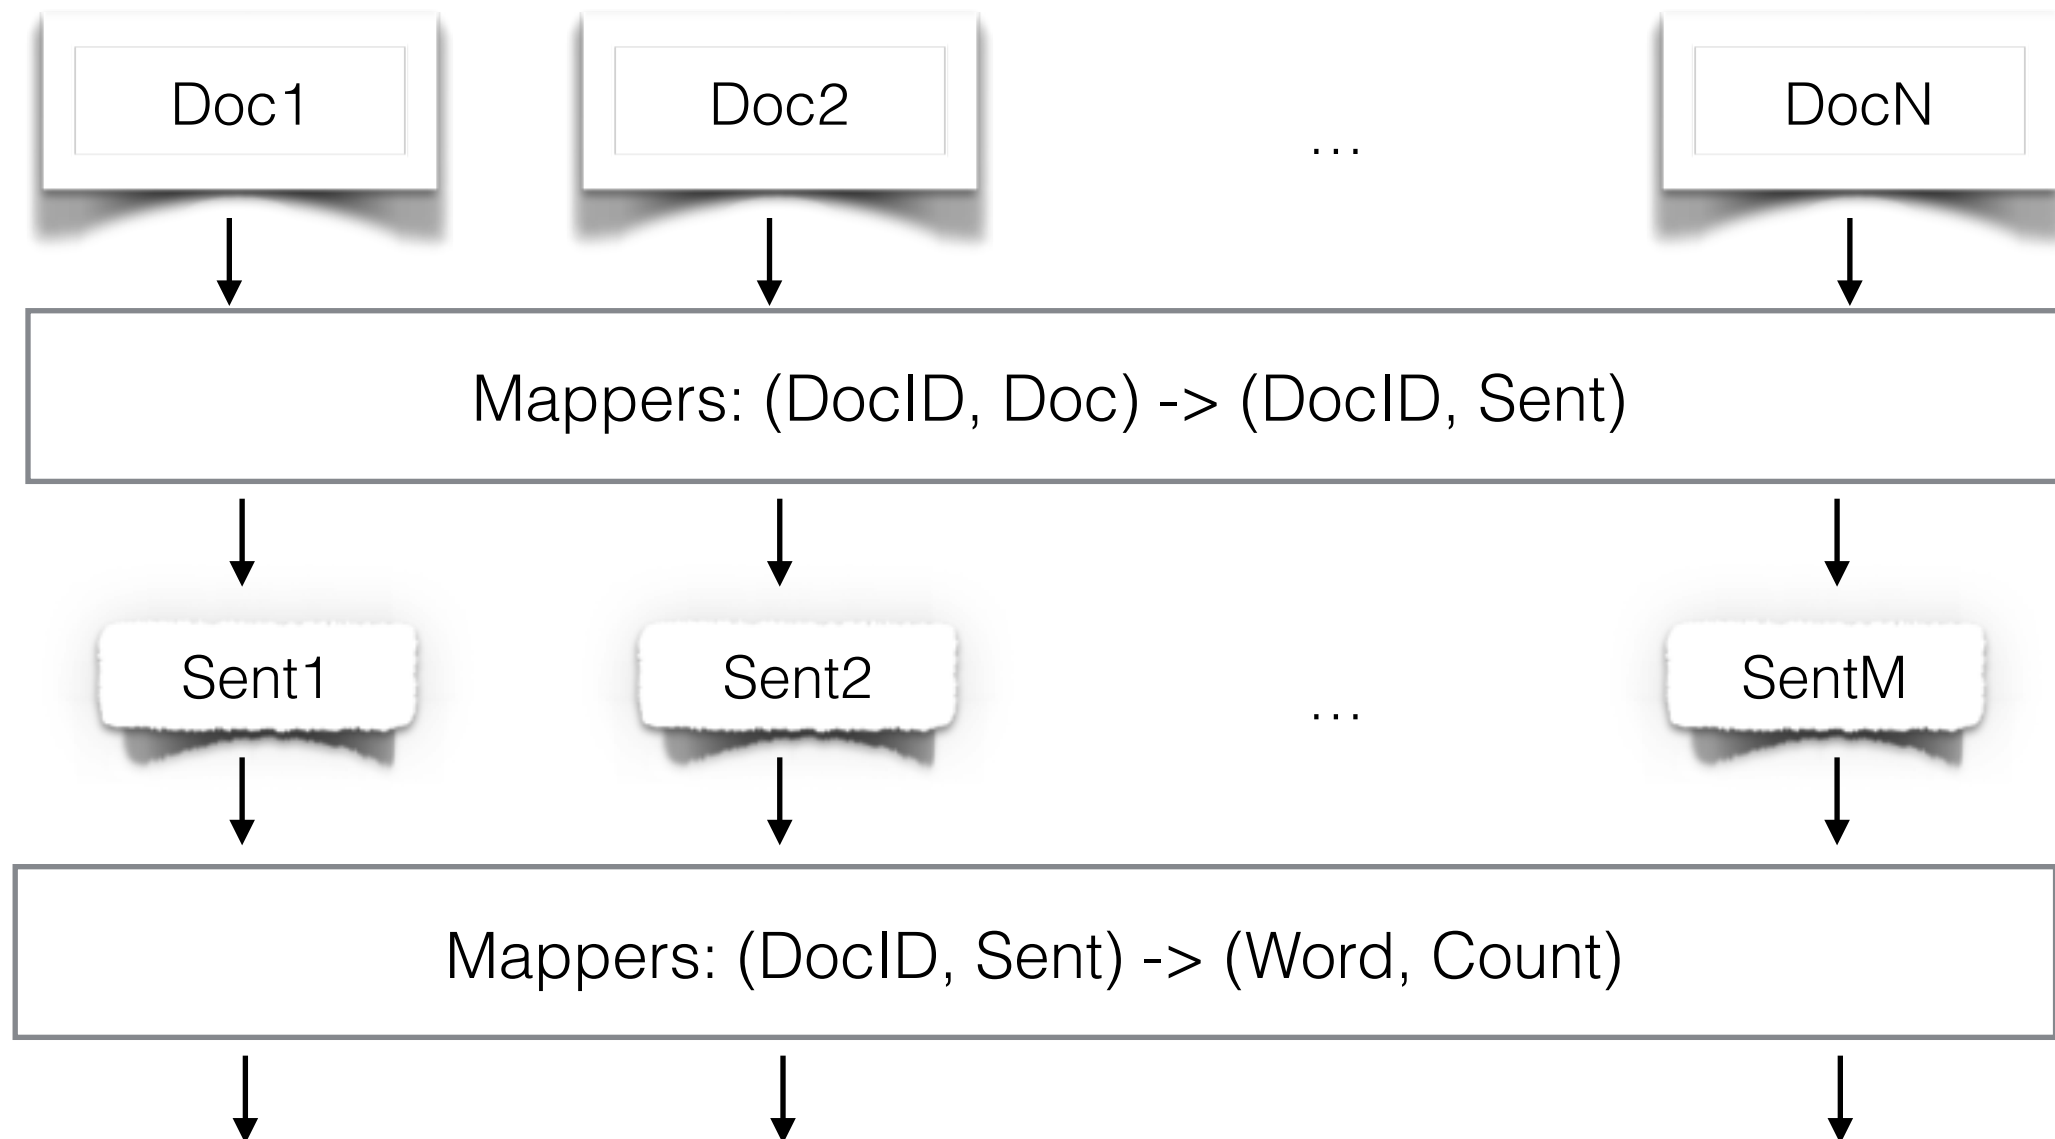
- (a) N
- (b) $\log(N)$
- (c) As many as we can afford.



Clicker Question!

In the best-case scenario, how much parallelization could we get here (maximum number of mappers)?



- (a) N
- (b) $\log(N)$
- (c) As many as we can afford.

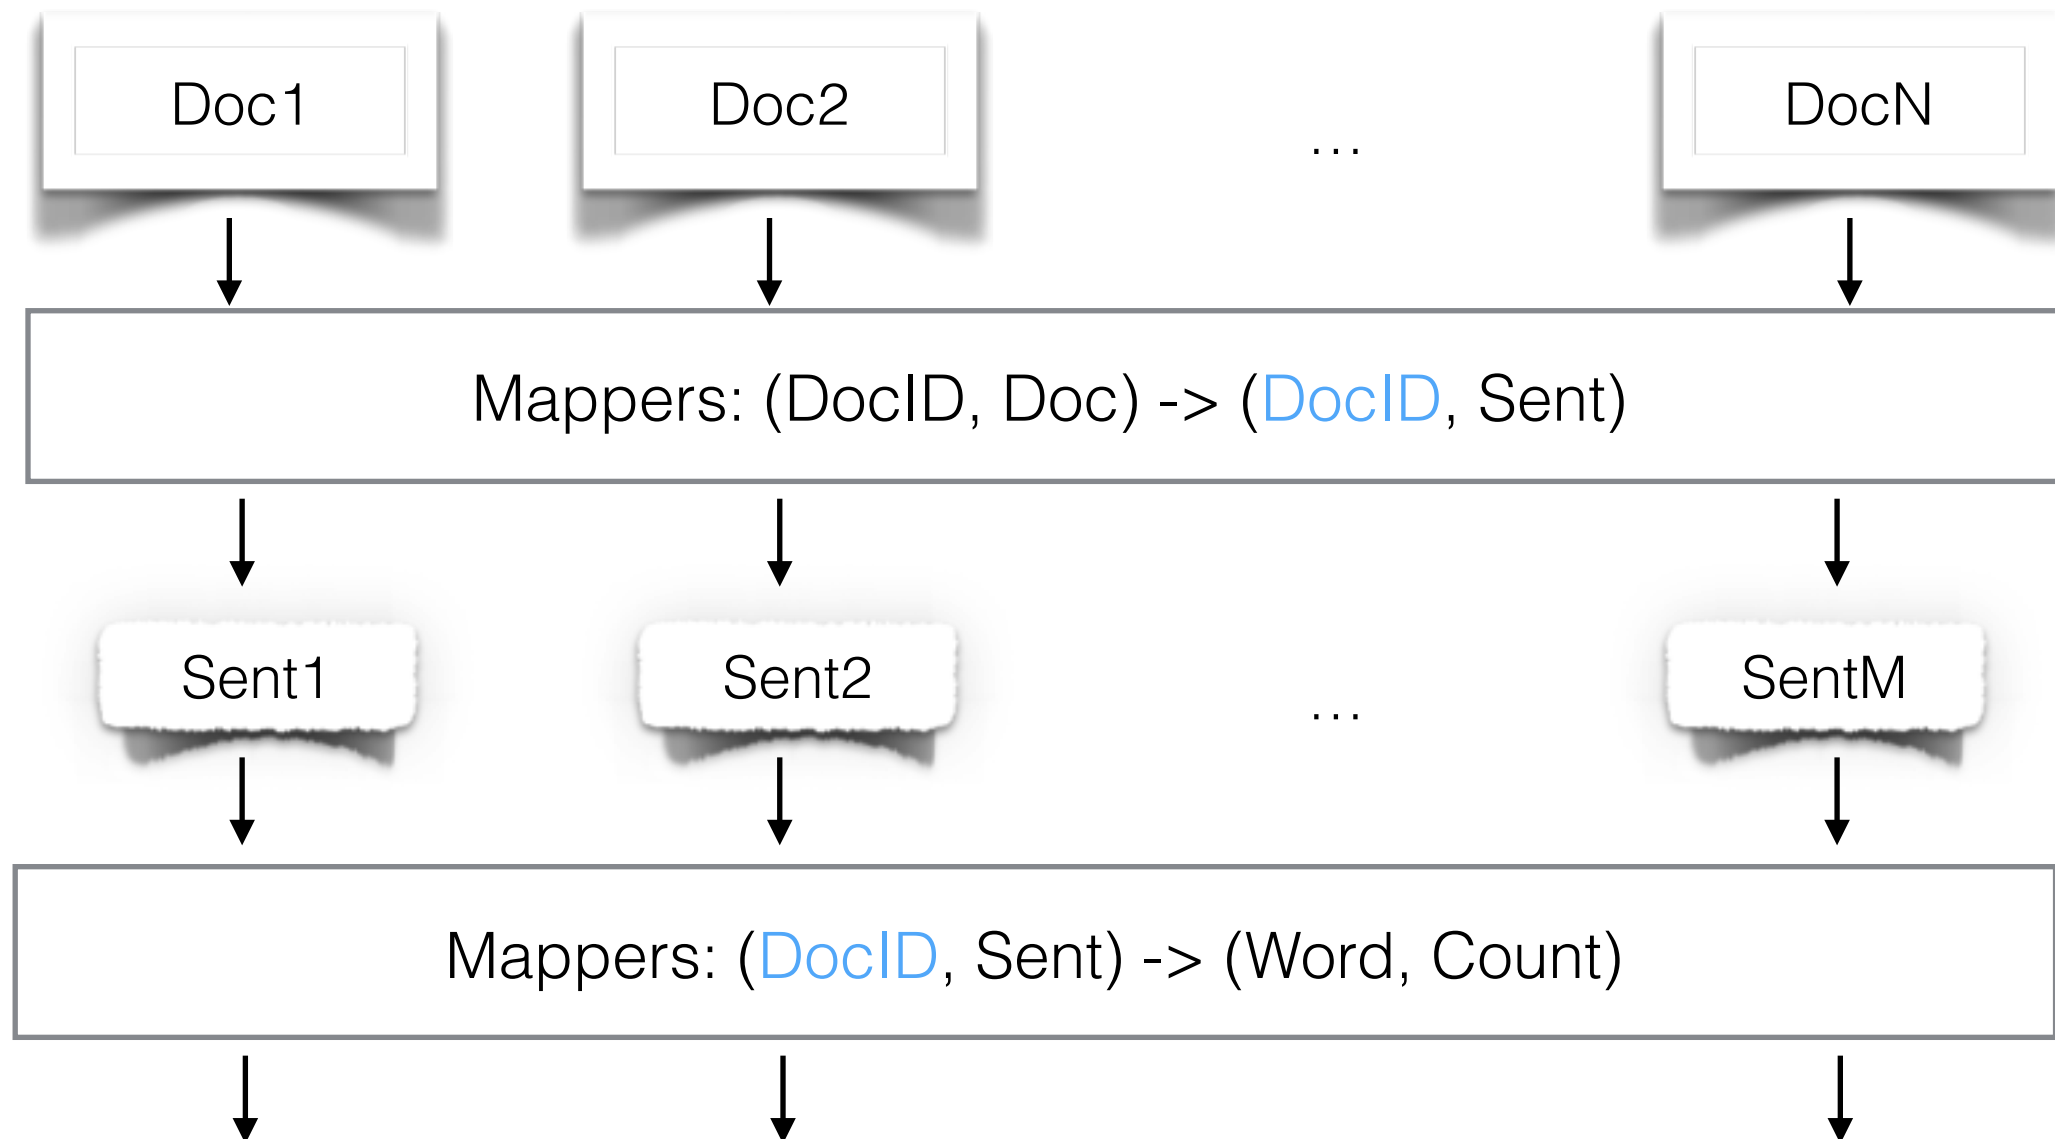


Clicker Question!
How about here?

- (a) N**
- (b) M**
- (c) N*M**







Clicker Question!
How about here?

- (a) N
- (b) M
- (c) $N * M$

Mapping doesn't require
the same keys to route
to the same machine.

Clicker Question!

Which is (likely to be) faster?

(a)

Mapper1:
(DocID, Doc) -> (DocID, Sent)



Mapper2:
(DocID, Sent) -> (Word, Count)



Reducer:
(Word, Count) -> Word,
sum(Count)

(b)

Mapper:
(DocID, Doc) -> (Word, Count)



Reducer:
(Word, Count) -> Word,
sum(Count)

(c) They are the same

Doc = list_of(Sentence)
Sentence = list_of(Word)

er Question!
Which is (likely to be) faster?

(a)

Mapper1:
(DocID, Doc) -> (DocID, Sent)



Mapper2:
(DocID, Sent) -> (Word, Count)



Reducer:
(Word, Count) -> Word,
sum(Count)

(b)

Mapper:
(DocID, Doc) -> (Word, Count)



Reducer:
(Word, Count) -> Word,
sum(Count)

(c) They are the same

Clicker Question!

Which is (likely to be) faster?

(a)

Mapper1:
(DocID, Doc) -> (DocID, Sent)



Mapper2:
(DocID, Sent) -> (Word, Count)



Reducer:
(Word, Count) -> Word,
sum(Count)

(b)

Mapper:
(DocID, Doc) -> (Word, Count)



Reducer:
(Word, Count) -> Word,
sum(Count)

(c) They are the same

Clicker Question!

Which is (likely to be) faster?

(a)

Mapper1:
(DocID, Doc) -> (DocID, Sent)

(b)

Mapper:
(DocID, Doc) -> (Word, Count)

Smaller jobs = more
dynamic load balancing
and faster recovery from
failure

Ma
(DocID, Sent)

Word,
t)



Reducer:
(Word, Count) -> Word,
sum(Count)

(c) They are the same

Clicker Question!

Which is (likely to be) faster?

(a)

Mapper1:
(DocID, Doc) -> (DocID, Sent)



In general, nested
loops should be
refactored into
multiple mappers
sum(Count)

(b)

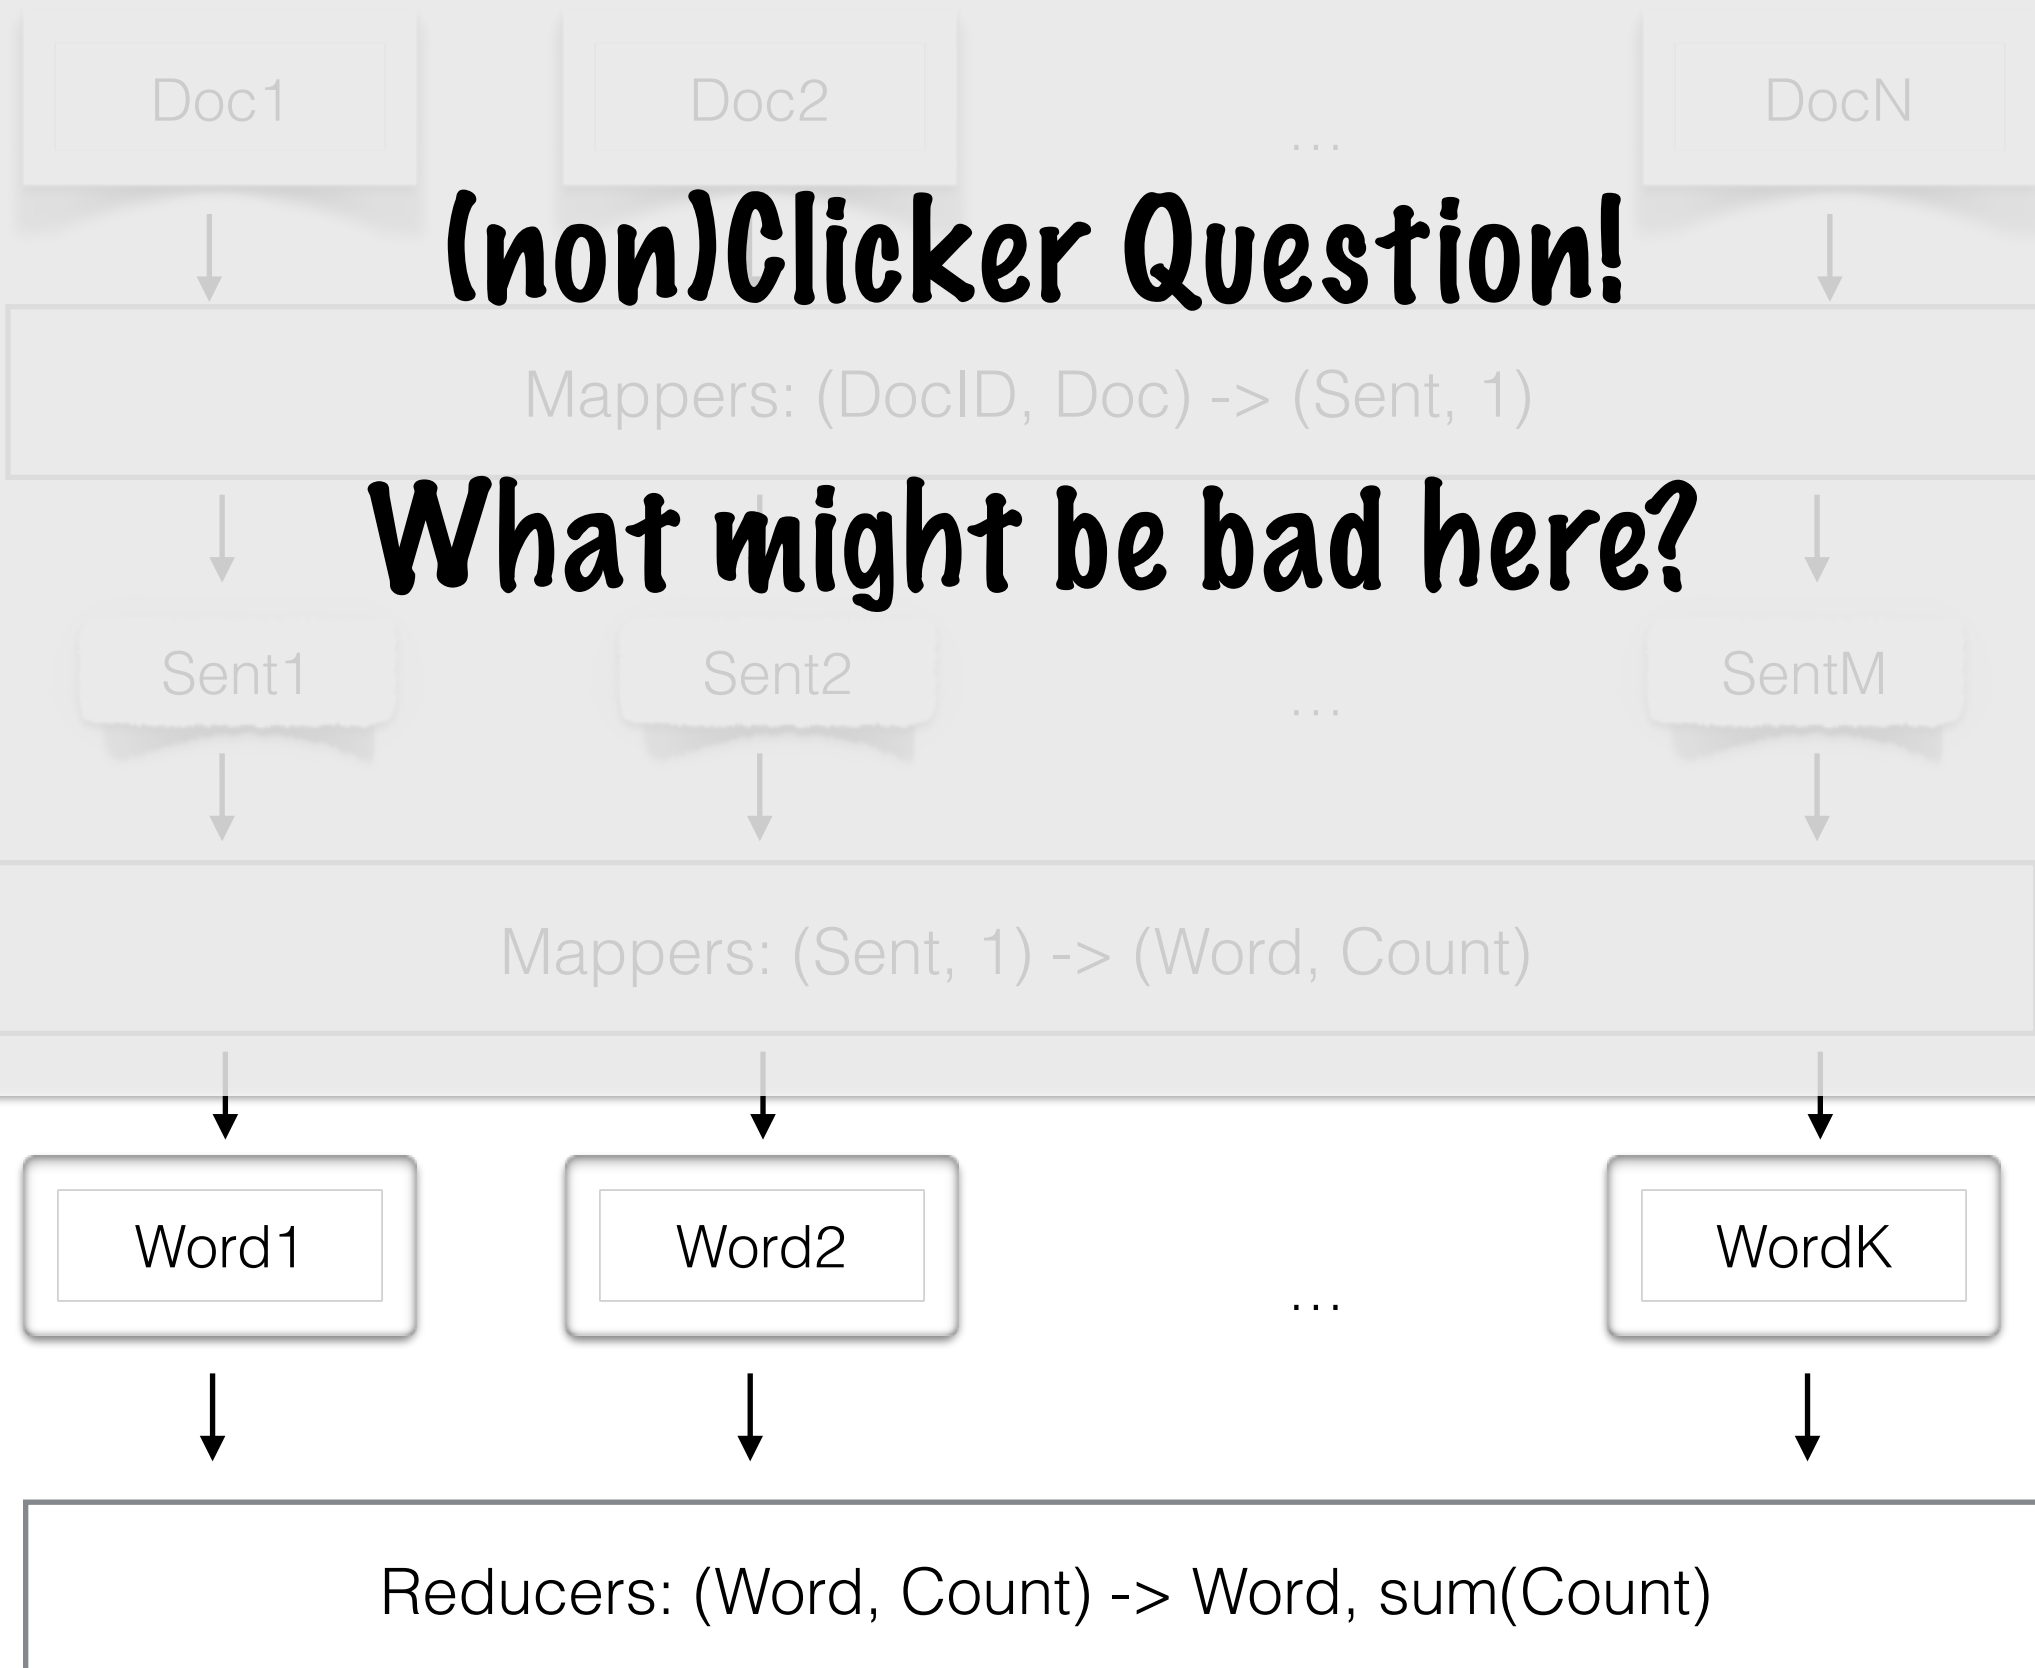
Mapper:
(DocID, Doc) -> (Word, Count)

for sentence in doc:
 for word in sentence:
 blah blah



Reducer:
(Word, Count) -> Word,
sum(Count)

(c) They are the same



Doc1

Doc2

...

DocN

(non)Clicker Question!

Mappers: (DocID, Doc) -> (Sent, 1)

What might be bad here?

Skewed Key Distributions!

**(Need all values with the same key to be together,
so can't automatically load balance)**

Sent1

Sent2

SentM

Word1


Word2

...


WordK

Reducers: (Word, Count) -> Word, sum(Count)


Zipf's Law




Zipf's Law




Zipf's Law



Zipf's Law



Zipf's Law



Real Life Application

Subject	Predicate	Object	Categories
Barack Obama	won	the electoral vote	Person, US_Presidents, Huffington_Post_Columnists
Kamala Lopez	wrote	an op-ed for HuffPo	Person, Huffington_Post_Columnists, Actor

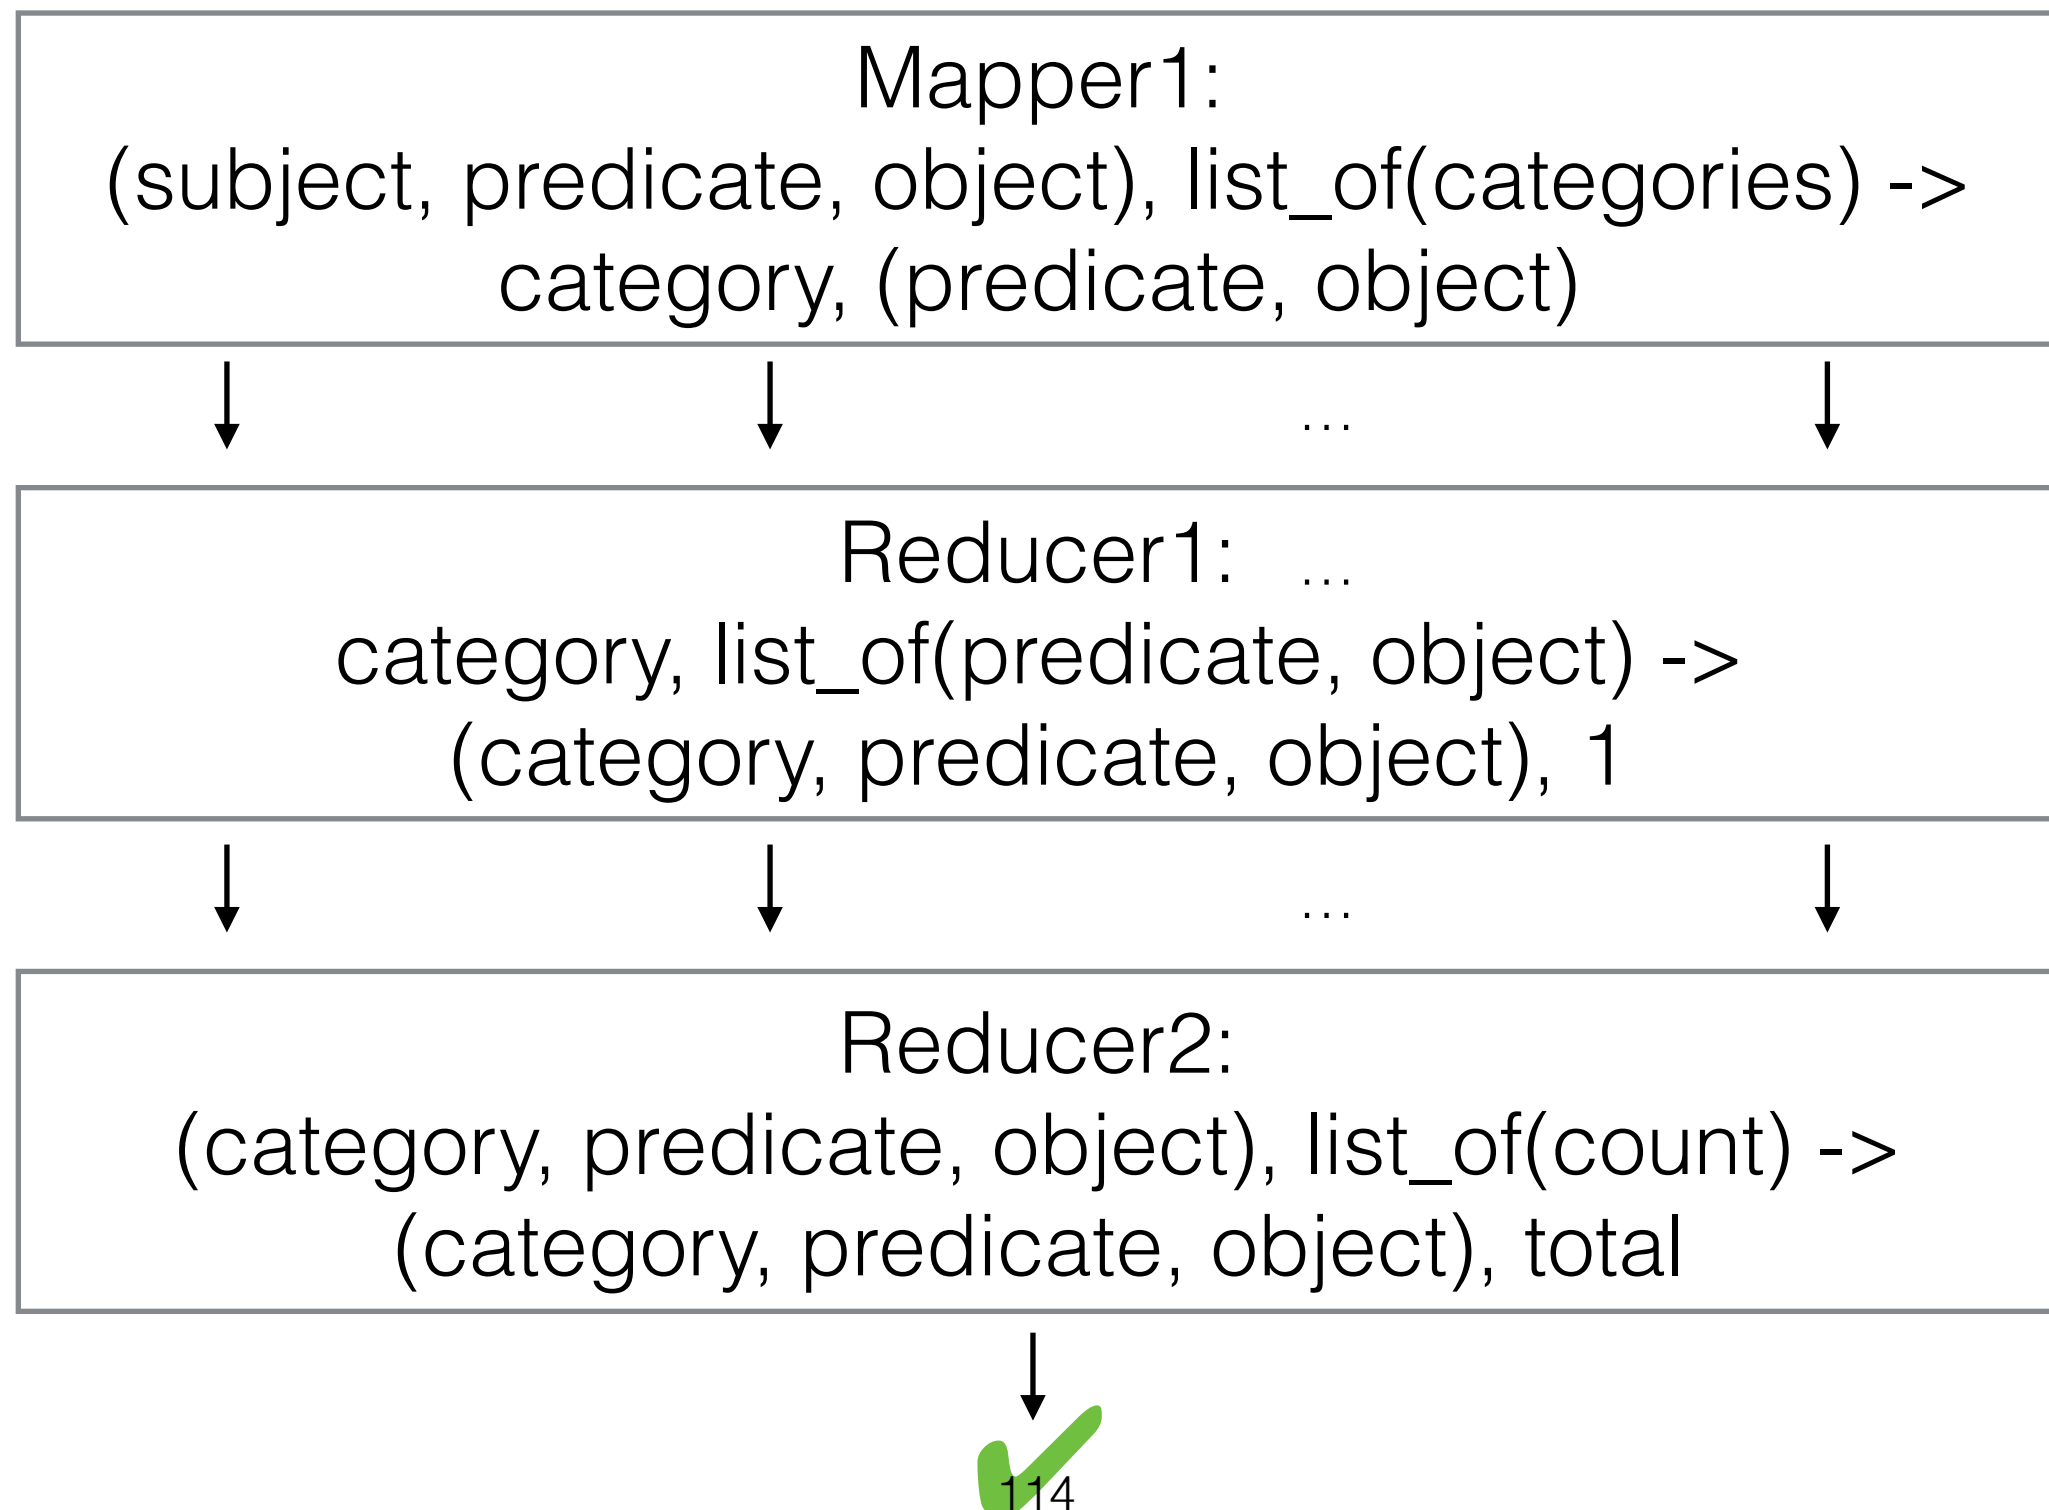
Predicate	Object	Category	Score
won	the electoral vote	US_Presidents	0.92
won	the electoral vote	Person	0.89
won	the electoral vote	Huffington Post Columnists	0.23
wrote	an op-ed for HuffPo	Huffington Post Columnists	0.99
wrote	an op-ed for HuffPo	Person	0.91

Real Life Application

Subject	Predicate	Object	Categories
Barack Obama	won	the electoral vote	Person, US_Presidents, Huffington_Post_Columnists
Kamala Lopez	wrote	an op-ed for HuffPo	Person, Huffington_Post_Columnists, Actor

Predicate	Object	Category	Score
won	the electoral vote	US_Presidents	702,345
won	the electoral vote	Person	812,485
won	the electoral vote	Huffington Post Columnists	24,571
wrote	an op-ed for HuffPo	Huffington Post Columnists	134,213
wrote	an op-ed for HuffPo	Person	136,091

First Attempt





First Attempt

Mapper1:

(subject, predicate, object), list_of(categories) ->
category, (predicate, object)



...



Reducer1: ...

category, list_of(predicate, object) ->
(category, predicate, object), 1



...



Reducer2:

(category, predicate, object), list_of(count) ->
(category, predicate, object), total





First Attempt

Mapper1:

(subject, predicate, object), list_of(categories) ->
category, (predicate, object)



...



Reducer1: ...

category, list_of(predicate, object) ->
(category, predicate, object), 1



(c

Every tuple involving a
single category (e.g.
"Person") has to go through
the same reducer...

>



First Attempt



Mapper1:

(subject, predicate, object), list_of(categories) ->
category, (predicate, object)



...



Reducer1: ...

category, list_of(predicate, object) ->
(category, predicate, object), 1



...



Reducer2:

(category, predicate, object), list_of(count) ->
(category, predicate, object), total





First Attempt



Mapper1:

(subject, predicate, object), list_of(categories) ->
category, (predicate, object)



...



Reducer1: ...

category, list_of(predicate, object) ->
(category, predicate, object), 1



...



Reducer2:

(category, predicate, object), list_of(count) ->
(category, predicate, object), total



So much better!

Mapper1:

(subject, predicate, object), list_of(categories) ->
(category, predicate, object), 1



...



Reducer2:

(category, predicate, object), list_of(count) ->
(category, predicate, object), total



Alright, scram.